

Custom Search

Search

# Re: crush: straw is dead, long live straw2

[\[Date Prev\]](#)[\[Date Next\]](#)[\[Thread Prev\]](#)[\[Thread Next\]](#)[\[Date Index\]](#)[\[Thread Index\]](#)

- *Subject:* Re: crush: straw is dead, long live straw2
- *From:* Sage Weil <[sweil@xxxxxxxxxx](mailto:sweil@xxxxxxxxxx)>
- *Date:* Fri, 12 Dec 2014 06:46:24 -0800 (PST)
- *Cc:* [ceph-devel@xxxxxxxxxxxxxxxxxx](mailto:ceph-devel@xxxxxxxxxxxxxxxxxx)
- *In-reply-to:* <[20141212091408.GN4150@thinkpad.thebehrens.net](mailto:20141212091408.GN4150@thinkpad.thebehrens.net)>
- *References:* <[alpine.DEB.2.00.1412081516010.18281@cobra.newdream.net](mailto:alpine.DEB.2.00.1412081516010.18281@cobra.newdream.net)>  
<[20141212091408.GN4150@thinkpad.thebehrens.net](mailto:20141212091408.GN4150@thinkpad.thebehrens.net)>
- *User-agent:* Alpine 2.00 (DEB 1167 2008-08-23)

0

On Fri, 12 Dec 2014, Thorsten Behrens wrote:

```
> Sage Weil wrote:
> > Calculating the ln() function is a bit annoying because it is a floating
> > point function and CRUSH is all fixed-point arithmetic (integer-based).
> > The current draft implementation uses a 128 KB lookup table (2 bytes per
> > entry for 16 bits of input precision). It seems to be about 25% slower
> > than the original straw code in my simple microbenchmark, but I'm not sure
> > that's a number we should trust since it loops through a zillion inputs
> > and will pull most of the lookup table into the processor caches.
> >
> > We could probably try a few others things:
> >
> [snip]
>
> My Hacker's Delight copy has a number of further variants (on
> pp. 215), unfortunately only covering log_2 (which is trivially the
> number of leading zeros in the binary representation) and log_10. But
> since the logarithms are simply related by a multiplicative constant,
> something clever might be doable here with a bit of thinking.
```

Any log will definitely do; the next steps are to do a division and then take the max. Simply counting 0's isn't precise enough, though.. we need a lot more precision than that.

The table lookup is (I think) on the order of 50-100 cycles with a cold cache; that's what we need to beat!

sage

--

To unsubscribe from this list: send the line "unsubscribe ceph-devel" in the body of a message to [majordomo@xxxxxxxxxxxxxxxxxx](mailto:majordomo@xxxxxxxxxxxxxxxxxx)

More majordomo info at <http://vger.kernel.org/majordomo-info.html>

- **Follow-Ups:**
  - [Re: crush: straw is dead, long live straw2](#)
    - *From:* Milosz Tanski
  - [Re: crush: straw is dead, long live straw2](#)
    - *From:* Joe Landman
- **References:**
  - [crush: straw is dead, long live straw2](#)
    - *From:* Sage Weil
  - [Re: crush: straw is dead, long live straw2](#)
    - *From:* Thorsten Behrens
- Prev by Date: [Re: the cpu optimize in ceph](#)
- Next by Date: [Reducing backfilling/recovery long tail](#)
- Previous by thread: [Re: crush: straw is dead, long live straw2](#)
- Next by thread: [Re: crush: straw is dead, long live straw2](#)
- Index(es):
  - [Date](#)
  - [Thread](#)

[\[Index of Archives\]](#)   [\[CEPH Users\]](#)   [\[Ceph Large\]](#)   [\[Information on CEPH\]](#)   [\[Linux BTRFS\]](#)   [\[Linux USB Devel\]](#)  
[\[Video for Linux\]](#)   [\[Linux Audio Users\]](#)   [\[Yosemite News\]](#)   [\[Linux Kernel\]](#)   [\[Linux SCSI\]](#)

---

