

UNIVERSITY OF CALIFORNIA
SANTA CRUZ

**CEPH: RELIABLE, SCALABLE, AND HIGH-PERFORMANCE
DISTRIBUTED STORAGE**

A dissertation submitted in partial satisfaction of the
requirements for the degree of

DOCTOR OF PHILOSOPHY

in

COMPUTER SCIENCE

by

Sage A. Weil

December 2007

The Dissertation of Sage A. Weil
is approved:

Professor Scott A. Brandt, Chair

Doctor Richard Golding

Professor Charlie McDowell

Professor Carlos Maltzahn

Professor Ethan L. Miller

Lisa C. Sloan
Vice Provost and Dean of Graduate Studies

Copyright © by

Sage A. Weil

2007

Table of Contents

List of Figures	vii
List of Tables	xiv
Abstract	xv
Dedication	xvii
Acknowledgments	xviii
1 Introduction	1
1.1 Contributions	3
1.2 Outline	5
2 Related Work	7
2.1 Local File Systems	7
2.2 Client-Server File Systems	10
2.3 Distributed File Systems	11
2.3.1 Wide-area File Systems	12
2.3.2 SAN File Systems	13
2.3.3 Object and Brick-based Storage	13
2.3.4 Non-POSIX Systems	15
3 Ceph Architecture	17
3.1 Ceph	17
3.2 Client Operation	20
3.2.1 File I/O and Capabilities	21
3.2.2 Client Synchronization	22
3.2.3 Namespace Operations	24
3.3 Communication Model	25

4	Distributed Metadata Management	27
4.1	Background and Related Work	30
4.1.1	Local File Systems	30
4.1.2	Distributed File Systems	32
4.2	Metadata Storage	37
4.2.1	Embedded Inodes	39
4.2.2	Remote Links and the Anchor Table	40
4.2.3	Large Journals	41
4.3	Adaptive Workload Distribution	42
4.3.1	Subtree Partitioning as Soft State	43
4.3.2	Metadata Replication	46
4.3.3	Locking	47
4.3.4	Load Balance	48
4.3.5	Subtree Migration	49
4.3.6	Directory Fragments	50
4.3.7	Traffic Control	52
4.4	Failure Recovery	53
4.4.1	Journal Structure	54
4.4.2	Failure Detection	55
4.4.3	Recovery	55
4.4.4	Discussion	59
4.5	Evaluation	59
4.5.1	Metadata Partitioning	60
4.5.2	Embedded Inodes	65
4.5.3	Journaling	69
4.5.4	Adaptive Distribution	71
4.5.5	Metadata Scaling	72
4.5.6	Failure Recovery	74
4.5.7	Availability	77
4.6	Future Work	78
4.7	Experiences	79
4.8	Conclusions	80
5	Data Distribution	82
5.1	Related Work	84
5.2	The CRUSH algorithm	86
5.2.1	Hierarchical Cluster Map	87
5.2.2	Replica Placement	88
5.2.3	Map Changes and Data Movement	95
5.2.4	Bucket Types	96
5.3	Evaluation	104
5.3.1	Data Distribution	104
5.3.2	Reorganization and Data Movement	107

5.3.3	Algorithm Performance	109
5.3.4	Reliability	112
5.4	Future Work	113
5.5	Conclusions	114
6	Distributed Object Storage	116
6.1	Overview	117
6.2	Distributed Object Storage	121
6.2.1	Data Placement	122
6.2.2	Cluster Maps	124
6.2.3	Communication and Failure Model	126
6.2.4	Monitors	127
6.2.5	Map Propagation	129
6.3	Reliable Autonomic Storage	130
6.3.1	Replication	130
6.3.2	Serialization versus Safety	132
6.3.3	Maps and Consistency	133
6.3.4	Versions and Logs	135
6.3.5	Failure Recovery	136
6.3.6	Client Locking and Caching	142
6.4	Performance and Scalability	144
6.4.1	Scalability	147
6.4.2	Failure Recovery	151
6.5	Future Work	153
6.6	Related Work	155
6.7	Conclusions	161
7	Local Object Storage	162
7.1	Object Storage Interface	162
7.2	Data Layout	164
7.3	Data Safety	164
7.4	Journaling	165
7.5	Evaluation	166
7.6	Related Work	167
7.7	Conclusion	168
8	Conclusion	169
8.1	Future Work	169
8.1.1	MDS Load Balancing	169
8.1.2	Client Interaction	170
8.1.3	Security	171
8.1.4	Quotas	171
8.1.5	Quality of Service	171

8.1.6	Snapshots	172
8.2	Summary	172
A	Communications Layer	175
A.1	Abstract Interface	175
A.2	SimpleMessenger	176
A.3	FakeMessenger	176
B	MDS Implementation Details	178
B.1	MDS Distributed Cache	178
B.1.1	Cache Structure	179
B.1.2	Replication and Authority	182
B.1.3	Subtree Partition	183
B.2	Metadata Storage	185
B.2.1	Directory Fragments and Versioning	185
B.2.2	Journal Entries	187
B.2.3	Anchor Table	188
B.3	Migration	189
B.3.1	Cache Infrastructure	189
B.3.2	Normal Migration	192
B.4	Failure Recovery	194
B.4.1	Journal Replay	195
B.4.2	Resolve Stage	195
B.4.3	Reconnect Stage	197
B.4.4	Rejoin Stage	197
B.5	Anchor Table	201
B.5.1	Table Updates	201
B.5.2	Failure Recovery	201
	Bibliography	203

List of Figures

3.1	System architecture. Clients perform file I/O by communicating directly with OSDs. Each process can either link directly to a client instance, or interact with a mounted file system.	18
4.1	All metadata for each directory, including file names (dentries) and the inodes they reference, are stored in a single object (identified by the directory's inode number) in the shared, distributed object store. Inode 1 is the root directory. . .	38
4.2	Structure of the metadata cache as seen by a single MDS. Metadata is partitioned into subtrees bounded by directory fragments. Each MDS replicates ancestor metadata for any locally managed subtrees. Large or busy directories are broken into multiple fragments, which can then form nested subtrees. . .	45
4.3	Directories are fragmented based on a tree in which each interior node has a 2^n children, and leaves correspond to individual fragments. Fragments are described by a bit pattern and mask, like an IP subnet, partitioning an integer namespace. Dentries are assigned to a specific fragment using a hash functions.	51
4.4	MDS performance as file system, cluster size, and client base are scaled. . . .	61

4.5	Percentage of cache devoted to ancestor inodes as the file system, client base and MDS cluster size scales. Hashed distributions devote large portions of their caches to ancestor directories. The dynamic subtree partition has slightly more ancestors than the static partition due to the re-delegation of subtrees nested within the hierarchy.	62
4.6	Cache hit rate as a function of cache size (as a fraction of total file system size). For smaller caches, inefficient cache utilization due to replicated ancestors results in lower hit rates.	63
4.7	No traffic control (top): nodes forward all requests to the authoritative MDS who slowly responds to them in sequence. Traffic control (bottom): the authoritative node quickly replicates the popular item and all nodes respond to requests.	64
4.8	Cumulative directory size for a variety of static file system snapshots. Many directories are very small (or even empty), and less than 5% contain more than 100 entries.	66
4.9	Study of hard link prevalence in file system snapshots. For each snapshot, the total file system size and the percentage of files with multiple links are shown. Stacked bars indicate the fraction of multilink files with all links in the same directory, with parallel cohorts linked by multiple directories, or with no obvious locality properties.	67

4.10	An MDS cluster adapts to a varying workload. At times 200 and 350 the workload shifts, as 400 clients begin creating files in private directories and then in a shared directory. In both cases load is redistributed, in the latter case after the large (and busy) directory is fragmented. At time 280 two additional MDSs join the cluster.	71
4.11	Per-MDS throughput under a variety of workloads and cluster sizes. As the cluster grows to 128 nodes, efficiency drops no more than 50% below perfect linear (horizontal) scaling for most workloads, allowing vastly improved performance over existing systems.	72
4.12	Average latency versus per-MDS throughput for different cluster sizes (<i>makedirs</i> workload).	74
4.13	Throughput in a small cluster before, during, and after MDS failures at time 400, 500, and two at 450. Unresponsive nodes are declared dead after 15 seconds, and in each case recovery for a 100 MB journal takes 4-6 seconds. . . .	75
4.14	Per-MDS throughput in a 4-node cluster under a compilation workload with heavy read sharing of <code>/lib</code> . Replication of popular metadata allows progress to continue even when the MDS managing the heavily shared directory suffers a failure.	77
5.1	A partial view of a four-level cluster map hierarchy consisting of rows, cabinets, and shelves of disks. Bold lines illustrate items selected by each <i>select</i> operation in the placement rule and fictitious mapping described by Table 5.1.	91

5.2	Reselection behavior of $select(6, disk)$ when device $r = 2$ (b) is rejected, where the boxes contain the CRUSH output \vec{R} of $n = 6$ devices numbered by rank. The left shows the “first n ” approach in which device ranks of existing devices (c, d, e, f) may shift. On the right, each rank has a probabilistically independent sequence of potential targets; here $f_r = 1$, and $r' = r + f_r n = 8$ (device h). . . .	94
5.3	Data movement in a binary hierarchy due to a node addition and the subsequent weight changes.	96
5.4	Node labeling strategy used for the binary tree comprising each tree bucket. . .	100
5.5	Efficiency of reorganization after adding or removing storage devices two levels deep into a four level, 7290 device CRUSH cluster hierarchy, versus $RUSH_P$ and $RUSH_T$. 1 is optimal.	108
5.6	Efficiency of reorganization after adding items to different bucket types. 1 is optimal. Straw and list buckets are normally optimal, although removing items from the tail of a list bucket induces worst case behavior. Tree bucket changes are bounded by the logarithm of the bucket size.	109
5.7	CRUSH and $RUSH_T$ computation times scale logarithmically relative to hierarchy size, while $RUSH_P$ scales linearly.	110
5.8	Low-level speed of mapping replicas into individual CRUSH buckets versus bucket size. Uniform buckets take constant time, tree buckets take logarithmic time, and list and straw buckets take linear time.	111

6.1	A cluster of many thousands of OSDs store all objects in the system. A small, tightly coupled cluster of monitors collectively manages the cluster map that describes the current cluster composition and the distribution of data. Each client instance exposes a simple storage interface to applications.	118
6.2	Objects are grouped into <i>placement groups</i> (PGs), and distributed to OSDs via CRUSH, a specialized replica placement function. Failed OSDs (<i>e. g.</i> <code>osd1</code>) are filtered out of the final mapping.	122
6.3	Replication strategies implemented by RADOS. Primary-copy processes both reads and writes on the first OSD and updates replicas in parallel, while chain forwards writes sequentially and processes reads at the tail. Splay replication combines parallel updates with reads at the tail to minimize update latency. . .	131
6.4	RADOS responds with an <i>ack</i> after the write has been applied to the buffer caches on all OSDs replicating the object (shown here with splay replication). Only after it has been safely committed to disk is a second <i>commit</i> notification sent to the client.	132
6.5	Each PG has a log summarizing recent object updates and deletions. The most recently applied operation is indicated by <i>last_update</i> . All updates above <i>last_complete</i> are known to have been applied, while any missing objects in the interval between <i>last_complete</i> and <i>last_update</i> are summarized in the missing list.	134

6.6	Reads in shared read/write workloads are usually unaffected by write operations. However, reads of uncommitted data can be delayed until the update commits. This increases read latency in certain cases, but maintains a fully consistent behavior for concurrent read and write operations in the event that all OSDs in the placement group simultaneously fail and the write <i>ack</i> is not delivered.	143
6.7	Per-OSD write performance. The horizontal line indicates the upper limit imposed by the physical disk. Replication has minimal impact on OSD throughput, although if the number of OSDs is fixed, <i>n</i> -way replication reduces total <i>effective</i> throughput by a factor of <i>n</i> because replicated data must be written to <i>n</i> OSDs.	144
6.8	Write latency for varying write sizes and primary-copy replication. More than two replicas incurs minimal additional cost for small writes because replicated updates occur concurrently. For large synchronous writes, transmission times dominate. Clients partially mask that latency for writes over 128 KB by acquiring exclusive locks and asynchronously flushing the data.	145
6.9	Latency for 4 KB writes under different replication strategies and levels of replication. Splay replication augments chain replication with parallel replica updates to achieve performance similar to primary-copy replication for high levels of replication.	146

6.10	Normalized write latency per replica for 1 MB writes under different replication strategies and levels of replication. Parallel replication does not improve latency for replications levels below six.	147
6.11	OSD write performance scales linearly with the size of the OSD cluster until the switch is saturated at 24 OSDs. CRUSH and hash performance improves when more PGs lower variance in OSD utilization.	148
6.12	Duplication of map updates received by individual OSDs as the size of the cluster grows. The number of placement groups on each OSD effects number of peers it has who may share map updates.	151
6.13	Write throughput over time as a saturated 20 OSD cluster recovers from two OSD failures at time 30. Data re-replication begins at time 50 and completes at time 80.	152
6.14	Write throughput over time as an unsaturated 20 OSD cluster recovers from one OSD failure at time 30 and two more at time 80.	152
7.1	Performance of EBOFS compared to general-purpose file systems. Although small writes suffer from coarse locking in the prototype, EBOFS nearly saturates the disk for writes larger than 32 KB. Since EBOFS lays out data in large extents when it is written in large increments, it has significantly better read performance.	166

List of Tables

5.1	A simple rule that distributes three replicas across three cabinets in the same row.	91
5.2	Summary of mapping speed and data reorganization efficiency of different bucket types when items are added to or removed from a bucket.	97
5.3	As the total utilization of available storage approaches capacity, the number of devices that would otherwise overflow and that require adjustment increases. CRUSH computation increases slightly while decreasing variance.	106
6.1	Data elements present in the OSD cluster map, which describes both cluster state and the distribution of data.	125
B.1	Sample anchor table with two anchored inodes. Note that the full path is shown only for illustration and is not part of the table. The hash mark (#) in the parent column is used to denote a particular fragment of the given directory, where 0 indicates the directory is not fragmented. The /usr entry has a ref count of two because it is referenced by both /usr/bin and /usr/lib.	189

Abstract

Ceph: Reliable, Scalable, and High-performance Distributed Storage

by

Sage A. Weil

As the size and performance requirements of storage systems have increased, file system designers have looked to new architectures to facilitate system scalability. The emerging object-based storage paradigm diverges from server-based (*e. g.* NFS) and SAN-based storage systems by coupling processors and memory with disk drives, allowing systems to delegate low-level file system operations (*e. g.* allocation and scheduling) to object storage devices (OSDs) and decouple I/O (read/write) from metadata (file open/close) operations. Even recent object-based systems inherit a variety of decades-old architectural choices going back to early UNIX file systems, however, limiting their ability to effectively scale.

This dissertation shows that device intelligence can be leveraged to provide reliable, scalable, and high-performance file service in a dynamic cluster environment. It presents a distributed metadata management architecture that provides excellent performance and scalability by adapting to highly variable system workloads while tolerating arbitrary node crashes. A flexible and robust data distribution function places data objects in a large, dynamic cluster of storage devices, simplifying metadata and facilitating system scalability, while providing a uniform distribution of data, protection from correlated device failure, and efficient data migration. This placement algorithm facilitates the creation of a reliable and scalable object storage service that distributes the complexity of consistent data replication, failure detection, and recovery

across a heterogeneous cluster of semi-autonomous devices.

These architectural components, which have been implemented in the Ceph distributed file system, are evaluated under a variety of workloads that show superior I/O performance, scalable metadata management, and failure recovery.

To Elise and Riley,
who make it all worthwhile,
and to my parents, to whom I owe so much.

Acknowledgments

I would like to thank my advisor, Scott Brandt, and Carlos Maltzahn for their invaluable guidance, support, and friendship over these past few years.

I thank Ethan Miller for his collaboration and guidance in my research, and for his help in editing this thesis. I gratefully acknowledge Darrell Long, Richard Golding, and Theodore Wong for their excellent suggestions for my work. I thank Charlie McDowell for his generous service on my committee.

I would like to thank the students of the SSRC for their support and friendship. Kristal Pollack, Feng Wang, Chris Xin, Andrew Leung, Chris Olson, Joel Wu, Tim Bisson, and Rosie Wacha have provided a stimulating and challenging environment, and have all been a pleasure to work with over the years.

Finally, I would like to thank Bill Loewe and Tyce McLarty for their continued support. In particular, their generous lobbying on my behalf helped secure access to computing resources at LLNL that were extremely helpful in my research.

This research was supported by Lawrence Livermore National Laboratory.

Chapter 1

Introduction

System designers have long sought to improve the performance of file systems, which have proved critical to the overall performance of an exceedingly broad class of applications. The scientific and high-performance computing communities in particular have driven advances in the performance and scalability of distributed storage systems, typically predicting more general purpose needs by a few years. Traditional solutions, exemplified by NFS [72], provide a straightforward model in which a server exports a file system hierarchy that clients can map into their local name space. Although widely used, the centralization inherent in the client/server model has proven a significant obstacle to scalable performance.

More recent distributed file systems have adopted architectures based on object-based storage, in which conventional hard disks are replaced with intelligent object storage devices (OSDs) which combine a CPU, network interface, and local cache with an underlying disk or RAID [14, 30, 31, 104, 107]. OSDs replace the traditional block-level interface with one in which clients can read or write byte ranges to much larger (and often variably sized) named

objects, distributing low-level block allocation decisions to the devices themselves. Clients typically interact with a metadata server (MDS) to perform metadata operations (*open, rename*), while communicating directly with OSDs to perform file I/O (reads and writes), significantly improving overall scalability.

Systems adopting this model continue to suffer from scalability limitations due to little or no distribution of the metadata workload. Continued reliance on traditional file system principles like allocation lists and inode tables and a reluctance to delegate intelligence to the OSDs have further limited scalability and performance, and increased the cost of reliability.

I have developed a prototype for Ceph [100], a distributed file system that provides excellent performance, reliability, and scalability. The architecture is based on the assumption that systems at the petabyte scale are inherently dynamic: large systems are inevitably built incrementally, node failures are the norm rather than the exception, and the quality and character of workloads are constantly shifting over time.

Ceph decouples data and metadata operations by eliminating file allocation tables and replacing them with a novel generating function. This allows Ceph to leverage the intelligence present in OSDs to distribute the complexity surrounding data access, update serialization, replication and reliability, failure detection, and recovery. Ceph utilizes an adaptive distributed metadata cluster architecture that streamlines common metadata operations while dramatically improving the scalability of metadata access, and with it, the scalability of the entire system. I discuss the goals and workload assumptions motivating choices in the design of the architecture, analyze their impact on system scalability, performance, and reliability, and relate experiences in implementing a functional system prototype.

1.1 Contributions

The thesis of this dissertation is that device intelligence can be leveraged to provide reliable, scalable, and high-performance file service in a dynamic cluster environment.

The primary contributions of this dissertation are threefold.

First, I present a distributed metadata management architecture that provides excellent performance and scalability while seamlessly tolerating arbitrary node crashes. Ceph’s MDS diverges from conventional metadata storage techniques, and in doing so facilitates adaptive file system and workload partitioning among servers, improved metadata availability, and failure recovery. Specifically, file system metadata updates are initially written to large, lazily-trimmed per-MDS journals that absorb temporary and repetitive updates. File (inode) metadata is then embedded in the file system namespace and stored inside per-directory objects for efficient read access and metadata prefetching. I present a comparative analysis of distributed metadata partitioning techniques, and describe the novel dynamic subtree partitioning approach used by Ceph. Notably, my MDS defines the namespace hierarchy in terms of directory fragments, facilitating fine-grained load distribution even for large or busy directories, and implements a traffic control mechanism for dispersing load generated by flash crowds—sudden concurrent access by thousands of client nodes—across multiple nodes in the MDS cluster.

The second contribution of this work is a robust hierarchical data distribution function that places data in a large distributed storage system. When used in place of conventional allocation tables, the algorithm efficiently addresses a range of critical storage-related placement issues, including statistically uniform data distribution, correlated failure and data safety, and

data migration in dynamically changing storage clusters. Specifically, the algorithm places sets of data objects (either object replicas, or parity or erasure coded fragments) in a hierarchy of storage devices using a flexible rule language. The flexible hierarchical cluster description facilitates an adjustable balance between extremely efficient placement calculations and mapping stability when devices are added or removed from the cluster. It further provides the ability to enforce the separation of replicas across user-defined failure domains, limiting exposure to data loss due to correlated failures.

My third contribution is a distributed object storage architecture that leverages device intelligence to provide a reliable and scalable storage abstraction with minimal oversight. Specifically, I describe an efficient, scalable, and low-overhead cluster management protocol that facilitates consistent and coherent data access through the propagation of small cluster maps that specify device membership and data distribution. This allows a dynamic cluster of semi-autonomous OSDs to self-manage consistent data replication, failure detection, and failure recovery while providing the illusion of a single logical object store with direct, high-performance client access to data.

These contributions have been implemented as part of the Ceph file system prototype. The prototype is written in roughly 40,000 semicolon-lines of C++ code and has been released as open source under the Lesser GNU Public License (LGPL) to serve as a reference implementation and research platform.

1.2 Outline

The remainder of this dissertation is structured as follows.

Chapter 2 introduces past and existing approaches to file system design. General architectural issues affecting performance, scalability, and reliability are introduced, and in certain cases related to Ceph’s contributions. In general, however, specifically related work is discussed in context in the chapters that follow.

Chapter 3 introduces the Ceph architecture and key design features. The overall operation of the system is described from the perspective of a client performing basic file system operations in order to provide an overview of how the various system components interact.

Chapter 4 describes the design, implementation, and performance characteristics of Ceph’s metadata server (MDS). I focus on the design implications of Ceph’s unconventional approach to file (inode) storage and update journaling on metadata storage, dynamic workload distribution, and failure recovery. A variety of static file system snapshots and workload traces are analyzed to motivate design decisions and performance analysis, and performance is evaluated under a range of micro-benchmarks and workloads under both normal-use and failure scenarios.

Chapter 5 describes CRUSH, the special-purpose data placement algorithm used by Ceph to distribute data among object storage devices. Specifically, I address the problem of managing data layout in large, dynamic clusters of devices by eliminating conventional allocation maps and replacing them with a low-cost *function*. A variety of issues related to scalability, data safety, and performance are considered.

Chapter 6 introduces RADOS, Ceph's Reliable Autonomic Distributed Object Store. RADOS provides an extremely scalable storage cluster management platform that exposes a simple object storage abstraction. System components utilizing the object store can expect scalable, reliable, and high-performance access to objects without concerning themselves with the details of replication, data redistribution (when cluster membership expands or contracts), failure detection, or failure recovery.

Chapter 7 describes EBOFS, the low-level object storage library utilized by RADOS to efficiently manage collections of objects on locally attached disks.

Finally, Chapter 8 provides some concluding remarks and describes avenues of continued research.

Chapter 2

Related Work

This research draws upon a large body of experience and research with file and storage systems. This includes experience with local file systems (which interact with locally attached disks), network file systems that interact with a remote server, and a range of distributed file system architectures in which file system data is spread across multiple hosts.

2.1 Local File Systems

File system design over the past three decades has been heavily influenced by the original Unix file system, and the Fast File System (FFS) [65] found in BSD Unix. The interface and behavior—semantics—of these early file systems formed the basis of the POSIX SUS (Single Unix Specification) standard [12] to which most modern systems and applications conform.

BSD’s FFS, like most “local” file systems, is designed to operate on a locally attached hard disk storing fixed-size sectors or blocks. (In contrast, Ceph is a “cluster” file system that

operates on a large number of hosts connected by a network.) Hard disks at the time were slow to “seek” or position themselves over a given block, but once positioned could read data relatively quickly. To avoid positioning delays as much as possible, FFS featured *cylinder groups*, representing localized regions of disk falling within the same cylindrical regions of the spinning hard disk platter. Related data and metadata were stored in the same cylinder group, such that most seeks involved minimal repositioning of the disk arm.

Over the past few decades, disk performance has increased significantly. However, positioning latencies have dropped much more slowly than disk transfer rates have increased, making careful layout of data on disk even more important than before. The desire to reduce position latencies motivated the design of the Log-structured File System (LFS) in the Sprite network operating system [26, 78]. LFS avoided positioning delays for write workloads by laying out new data on disk in a sequential log. Metadata structures were periodically flushed to disk (in the log) to allow previously written data to be located and re-read. However, in order to reclaim disk space from deleted files, LFS required a “cleaner” to relocate partially deallocated regions of the log, which degraded performance under many workloads [87, 88].

DualFS [73] separates data and metadata storage—not unlike object-based cluster file systems—by writing metadata to a separate device using an LFS-style log, mitigating cleaner effects by eliminating bulky file data. Although—as with LFS—a special inode map file is required to translate inode numbers to locations on disk, DualFS improves metadata performance by storing recently modified dentries and inodes close together in the log.

Although the cylinder group concept has been refined over the years, the basic metadata structures and allocation strategies introduced by FFS still persist in typical local file sys-

tems, such as ext3 [95]. As file system sizes have scaled more quickly than hard disk architectures, metadata structures have adapted. For example, file sizes have increased while disk block sizes have remained small—usually 4 KB—for efficient storage of small files. For large files, this gives rise to extremely long block allocation lists enumerating the location of file data. Modern local file systems are exemplified by XFS [92], which replaces allocation lists with *extents*—start and length pairs—to compactly describe large regions of disk.

Modern file systems have also introduced reliability features to streamline fast recovery after a failure. As with most commonly used alternatives, XFS maintains an on-disk *journal* file that it uses to log metadata updates that it is about to perform. In the event of a failure, the journal can be read to ensure that updates were successfully and correctly applied, or clean up any partially applied updates. Although journals come with a performance penalty—disks have to frequently reposition themselves to append to the journal, and each metadata update is written to disk twice—they avoid the need for more costly consistency checks during failure recovery (which are increasingly expensive as file systems scale).

Many systems also implement a form of *soft updates*, in which modifications are written only to unallocated regions of disk and in careful order to ensure that the on-disk image is always consistent (or easily repaired) [64]. The WAFL file system [40], for example, uses a copy-on-write approach when update the hierarchical metadata structures, writing all new data to unallocated regions of disk, such that changes are committed by simply updating a pointer to the root tree node. WAFL also maintains a journal-like structure, but does so only to preserve update durability between commits, not to maintain file system consistency. A similar technique is used by EBOFS, the object-storage component of Ceph (see Chapter 7).

2.2 Client-Server File Systems

The desire to share storage resources and data in networked computing environments has given rise to a number of client-server file systems. The most commonly used are NFS [72] and CIFS [38], which allow a central server to “export” a local file system to remote systems, who then map it into their own namespace. Centralizing storage has facilitated the creation of specialized, high-performance storage systems and popularized so-called NAS—network attached storage. However, the centralization inherent the client-server architecture has proved a significant impediment to scalability, because all file system operations must be processed by a single server. In large installations, administrators typically assign subsets of their data to specific servers, and map them into a single namespace by mounting multiple servers on each client. While reasonably effective and widespread, this approach can cause administrative headaches when certain data sets grow or contract and require manual migration to other servers.

Networked file systems typically relax consistency semantics in order to preserve cache performance in a distributed environment. For example, NFS clients write file data back to the server asynchronously, such that concurrent file access from other client hosts may not always return the most recent copy of data. Similarly, clients typically cache file metadata (*e. g.* as seen by *stat*) for a fixed interval to limit client interaction with and load on the file server. This relaxation of file system consistency can cause problems for many applications, precluding their use in NFS-based environments.

2.3 Distributed File Systems

Distributed file systems attempt to address the fundamental load balancing and scaling challenges inherent in client-server systems. Early distributed file systems include AFS [43], Coda [85], and that found in Sprite [71]. These utilize a central server to coordinate file system access, and issue *leases* to explicitly promise the validity of data (or metadata) provided to the client cache for a specified period, while subsequent *callbacks* can revoke the promise in the event of a conflicting access. However, while Sprite simply disabled caching in the event of (relatively rare) write sharing, AFS adopts an alternative consistency model constrained by file open and close instead of individual read and write events. In contrast, NFS (prior to version 4, which adds limited lease-like support for file “delegations”) is stateless by design, sacrificing consistency in the presence of data sharing.

Sprite partitions the file system among different servers by statically separating the namespace into “domains,” and dynamically mapping each to server. AFS utilizes a hybrid scheme that partitions files based on both their name and identifier among volumes, each of which is then assigned to a server. In both systems, partition boundaries are visible to the user through the lack of support for *link* or atomic *rename*. In contrast to these systems, Ceph aims to provide a unified file system namespace and fully consistent POSIX semantics.

Distributed architectures frequently replicate data across multiple servers for reliability and availability in the presence of failure. In the Harp [61] file system, servers employ primary-copy replication and write-ahead logs to allow fast recovery in the event of a node failure. A range of storage subsystems use similar techniques to provide reliable block-based

access to storage. Petal [56] provides a sparse virtual disk abstraction that is managed by a dynamic cluster of servers. FAB [82] employs a replica consistency protocol based on majority voting. Ceph similarly provides a distributed and consistent storage service with (unlike Petal) an object-based interface, a comparatively simple (relative to FAB) replica consistency protocol, and greater scalability (see Chapter 6).

Much like local file systems leverage RAID for reliability, systems like Frangipani [94] exploit a reliable distributed storage-layer abstraction like Petal to build a distributed file system. Boxwood [63] similarly provides a file service based on reliable, replicated block devices, although it does so by first constructing a rich distributed B-tree data structure. Ceph takes a somewhat analogous approach with its use of the RADOS abstraction to provide reliable and scalable storage, although it uses an object-based (rather than block-based) interface (see Section 2.3.3).

2.3.1 Wide-area File Systems

Many systems consider distribution of file system components over a wide area or even (in some cases) in a global environment. The xFS [99] file system used an invalidation-based cache-consistency protocol to facilitate aggressive client caching, in certain cases allowing reads from peer's caches [23, 22], to minimize costly communication over wide-area networks, though with a significant cost in complexity. The OceanStore [52] file system aims to build a globally available file storage service with erasure codes and a location-independent routing abstraction. Pangaea [83] targets a similar wide-area environment, aggressively replicating file contents where they are read, while relaxing replica consistency where strong guaran-

tees are not explicitly required. All of these systems focus largely on minimizing costly network communication (in some cases at the expense of consistency) and scalability over a wide area; high-performance access to in a clustered environment—where network communication is relatively cheap—is not a design goal as it is in Ceph.

2.3.2 SAN File Systems

Recognizing that most file I/O is inherently parallel—that is, I/O to different files is unrelated semantically—most recent “cluster” file systems are based on the basic idea of shared access to underlying storage devices. So-called SAN—storage area network—file systems are based on hard disks or RAID controllers communicating via a fibre channel (or similar) network, allowing any connected host to issue commands to any connected disk.

Most SAN file systems utilize a distributed lock manager (DLM) [51] to coordinate access to shared block storage. Additional steps are taken to limit the probability of lock contention, exemplified by the partial distribution of block allocation in GPFS [86]. Other systems, such as StorageTank [66], recognizing that contention for locks protecting file system metadata can limit scalability, utilize an auxiliary cluster of servers for managing file system metadata—the hierarchical namespace, access control, and block allocation.

2.3.3 Object and Brick-based Storage

Recently, many file systems and platforms, including Federated Array of Bricks (FAB) [82], PVFS [55], and pNFS [39] have been designed around clusters of network attached storage servers [31]. Like StorageTank, metadata is managed by an independent server (or servers),

but file I/O takes place by interacting directly with a cluster of storage servers (often termed “bricks”) instead of fibre channel disks. This allows certain functionality—such as low-level block allocation decisions—to be distributed to the devices themselves.

Lustre [14], the Panasas file system [104], zFS [76], Sorrento [93], Ursa Minor [1], and Kybos [107] are based on the closely-related object-based storage paradigm [7] popularized by NASD [33]. Instead of performing I/O in terms of small, fixed-size blocks, data is instead stored in *objects* that have a name (typically taken from a flat 128-bit namespace), a variable size (anywhere from bytes to gigabytes), and other per-object metadata, such as named attributes describing access permissions. Object-based storage allows easily distributed functionality like low-level block allocation or security capability enforcement to be performed by semi-intelligent devices, reducing contention on servers managing metadata and improving overall system scalability.

Although these object-based systems most closely resemble Ceph, none of them has the combination of scalable and adaptable metadata management, reliability, and fault tolerance that Ceph provides. Lustre and Panasas in particular delegate responsibility for low-level block allocation to OSDs, reducing the amount of file metadata, but otherwise do little to exploit device intelligence. With the exception of Sorrento, all of these systems use explicit allocation maps to specify where objects are stored, an approach that forces the involvement of the object directory server(s) in any migration of data between storage devices, limiting efficiency. Like Ceph, Sorrento uses a function to describe data placement on storage nodes instead of relying on traditional allocation tables. However, Sorrento’s hashed distribution [47] lacks Ceph’s support for efficient data migration, device weighting, and separation of replicas across failure domains

(see Chapter 5). Most importantly, Ceph’s functional specification of data layout allows OSDs to manage migration and failure recovery in a distributed fashion (see Chapter 6).

Kybos leverages device intelligence to manage quality-of-service reservations (something that Ceph does not do), and stores data using a network RAID protocol with two-phase updates (Ceph uses a simpler replication protocol). Ursa Minor provides a distributed object storage service that supports a range of redundancy strategies—including replication and parity-based encoding (RAID)—and consistency models, depending on application requirements. All of these systems, however, have limited support for efficient distributed metadata management, limiting their scalability and performance (see Chapter 4).

2.3.4 Non-POSIX Systems

A number of distributed file systems adopt alternative (non-POSIX) file system interfaces. Farsite [2, 25] federates a large number of unreliable workstations into a distributed file system providing Windows file service with Byzantine fault-tolerance. Notably, this does not include support for files linked to multiple names in the hierarchy (so-called “hard links”). However, like Ceph, Farsite dynamically distributes management of the file system namespace among available servers; dynamic subtree partitioning is discussed in greater detail in Chapter 4.

The Sorrento [93] file system is POSIX-like, but—much like AFS—adopts a relaxed consistency model that simplifies metadata management and update consistency in the presence of write sharing. Other systems simplify consistency more drastically: Cedar [32] translates the consistency problem into one of versioning by making shared files immutable, while in Venti [75], all data is considered immutable.

Although architecturally the Google File System [30] resembles object-based systems, it eschews standard file system interfaces entirely, and is optimized for very large files and a workload consisting largely of reads and file appends.

Chapter 3

Ceph Architecture

This chapter introduces the basic architectural features of Ceph. The overall operation of the system is presented by describing the Ceph client, and its interaction with various system components while performing basic file operations.

3.1 Ceph

The Ceph file system has three main components: the client, each instance of which exposes a near-POSIX file system interface to a host or process; a cluster of OSDs, which collectively stores all data and metadata; and a metadata server cluster, which manages the namespace (file names and directories) while coordinating security, consistency and coherence (see Figure 3.1). I say the Ceph interface is near-POSIX because I find it appropriate to extend the interface and selectively relax consistency semantics in order to better align both with the needs of applications and improve system performance (discussed in Section 3.2.2).

The primary goals of the architecture are scalability (to hundreds of petabytes and




Figure 3.1: System architecture. Clients perform file I/O by communicating directly with OSDs. Each process can either link directly to a client instance, or interact with a mounted file system.

beyond), performance, and reliability. Scalability is considered in a variety of dimensions, including the overall storage capacity and throughput of the system, and performance in terms of individual clients, directories, or files. Our target workload may include such extreme cases as tens or hundreds of thousands of hosts concurrently reading from or writing to the same file or creating files in the same directory. Such scenarios, common in scientific applications running on supercomputing clusters, are increasingly indicative of tomorrow’s general purpose workloads. More importantly, distributed file system workloads are inherently dynamic, with significant variation in data and metadata access as active applications and data sets change over time. Ceph directly addresses the issue of scalability while simultaneously achieving high performance, reliability and availability through three fundamental design features: decoupled data and metadata, dynamic distributed metadata management, and reliable autonomic distributed object storage.

- **Decoupled Data and Metadata**—Ceph maximizes the separation of file system metadata management from the storage of file data. Metadata operations (*open*, *rename*, etc.) are collectively managed by a metadata server cluster, while clients interact di-

rectly with OSDs to perform file I/O (reads and writes). Object-based storage has long promised to improve the scalability of file systems by delegating low-level block allocation decisions to individual devices. However, in contrast to existing object-based file systems [14, 104, 31, 30] which replace long per-file block lists with shorter object lists, Ceph eliminates allocation lists entirely. Instead, a simple function is used to name the objects containing file data based on inode number, byte range, and striping strategy, while a special-purpose data distribution function assigns objects to specific storage devices. This allows any party to calculate (rather than look up) the name and location of objects comprising a file's contents, eliminating the need to maintain and distribute object lists, simplifying the design of the system, and reducing the metadata cluster workload.

- **Dynamic Distributed Metadata Management**—Because file system metadata operations make up as much as half of typical file system workloads [77], effective metadata management is critical to overall system performance. Ceph utilizes a novel metadata cluster architecture based on dynamic subtree partitioning [102] that adaptively and intelligently distributes responsibility for managing the file system directory hierarchy among tens or even hundreds of MDSs. A (dynamic) hierarchical partition preserves locality in each MDS's workload, facilitating efficient updates and aggressive prefetching to improve performance for common workloads. Significantly, the workload distribution among metadata servers is based entirely on current access patterns, allowing Ceph to effectively utilize available MDS resources under any workload and achieve near-linear metadata performance scaling in the number of MDSs.

- **Reliable Autonomic Distributed Object Storage**—Large systems composed of many thousands of devices are inherently dynamic: they are built incrementally; they grow and contract as new storage is deployed and old devices are decommissioned; device failures are frequent and expected; and large volumes of data are created, moved, and deleted. All of these factors require that the distribution of data evolve to effectively utilize available resources and maintain the desired level of data replication. Ceph delegates responsibility for data migration, replication, failure detection, and failure recovery to the cluster of OSDs that is storing the data, while at a high level, OSDs collectively provide a single distributed and reliable object store to clients and metadata servers. This approach allows Ceph to more effectively leverage the intelligence (CPU and memory) present on each OSD to achieve reliable, highly available object storage with linear scaling.

3.2 Client Operation

I introduce the overall operation and interaction of Ceph's components and its interaction with applications by describing Ceph's client operation. The Ceph client runs on each host executing application code and exposes a file system interface to applications. In the Ceph prototype, the client code runs entirely in user space and can be accessed either by linking to it directly or as a mounted file system via FUSE (a user-space file system interface). Each client maintains its own file data cache, independent of the kernel page or buffer caches, making it accessible to applications that link to the client directly.

3.2.1 File I/O and Capabilities

When a process opens a file, the client sends a request to a node in the MDS cluster (see Section 4.3.1). An MDS traverses the file system hierarchy to translate the file name into the *file inode*, which includes a unique inode number, the file owner, mode, size, and other per-file metadata. If the file exists and access is granted, the MDS returns the inode number, file size, and information about the striping strategy used to map file data into objects. The MDS may also issue the client a *capability* (if it does not already have one) specifying which read or write operations are permitted. Capabilities currently include four bits controlling the client's ability to read, cache reads, write, and buffer writes. In the future, capabilities will include security keys allowing clients to prove to OSDs that they are authorized to read or write data [58, 69] (the prototype currently trusts all clients). Involvement in file I/O is limited to managing capabilities to preserve file consistency and achieve proper semantics.

Ceph generalizes a range of striping strategies to map file data onto a sequence of objects. Successive *stripe_unit* byte blocks are assigned to the first *stripe_count* objects, until objects reach a maximum *object_size* and we move to the next set of *stripe_count* objects. Whatever the layout (by default Ceph simply breaks files into 8 MB chunks), an additional field specifies how many replicas are stored of each object.

To avoid any need for file allocation metadata, object names are constructed by concatenating the file inode number and the object number. Object replicas are then assigned to OSDs using CRUSH, a globally known mapping function (described in detail in Chapter 5). For example, if one or more clients open a file for read-only access, an MDS grants them the

capability to read and cache file content. Armed with the inode number, layout, and file size, the clients can name and locate all objects containing file data and read directly from the OSD cluster. Any objects or byte ranges that don't exist are defined to be file "holes", or zeros. Similarly, if a client opens a file for writing, it is granted the capability to write with buffering, and any data it generates at any offset in the file is simply written to the appropriate object on the appropriate OSD. The client relinquishes the capability on file close and provides the MDS with the new file size (the largest offset written), which redefines the set of objects that (may) exist and contain file data.

3.2.2 Client Synchronization

POSIX semantics sensibly require that reads reflect any data previously written, and that writes are atomic (*i. e.* the result of overlapping, concurrent writes will reflect a particular order of occurrence). When a file is opened by multiple clients with either multiple writers or a mix of readers and writers, the MDS will revoke any previously issued read caching and write buffering capabilities, forcing all client I/O to be synchronous. That is, each application read or write operation will block until it is acknowledged by the OSD, effectively placing the burden of update serialization and synchronization with the OSD storing each object. Achieving atomicity is more complicated when writes span object boundaries. The prototype currently uses a simple locking mechanism to achieve correct serialization, although an alternative method that implements something closer to a true transaction is under consideration.

Not surprisingly, synchronous I/O can be a performance killer for applications, particularly those doing small reads or writes, due to the latency penalty—at least one round-trip

to the OSD. Although read-write sharing is relatively rare in general-purpose workloads [77], it is more common in scientific computing applications [98], where performance is often critical. For this reason, it is often desirable to relax consistency at the expense of strict standards conformance in situations where applications do not rely on it. Although Ceph supports such relaxation via a global switch, and many other distributed file systems punt on this issue [55, 72, 93], this is an imprecise and unsatisfying solution: either performance suffers, or consistency is lost system-wide.

For precisely this reason, a set of high performance computing extensions to the POSIX I/O interface have been proposed by the high-performance computing (HPC) community [103], a subset of which are implemented by Ceph. Most notably, these include an `O_LAZY` flag for *open* that allows applications to explicitly relax the usual coherency requirements for a shared-write file. Performance-conscious applications who manage their own consistency (e. g. by writing to different parts of the same file, a common pattern in HPC workloads [98]) are then allowed to buffer writes or cache reads when I/O would otherwise be performed synchronously. If desired, applications can then explicitly synchronize with two additional calls: *lazyio-propagate* will flush a given byte range to the object store, while *lazyio-synchronize* will ensure that the effects of previous propagations are reflected in any subsequent reads. The latter is implemented efficiently by provisionally invalidating cached data, such that subsequent read requests will be sent to the OSD but only return data if it is newer. The Ceph synchronization model thus retains its simplicity by providing correct read-write and shared-write semantics between clients via synchronous I/O, and extending the application interface to relax consistency for performance conscious distributed applications.

3.2.3 Namespace Operations

Client interaction with the file system namespace is managed by the metadata server cluster. Both read operations (*e. g. readdir, stat*) and updates (*e. g. unlink, chmod*) are synchronously applied by an MDS to ensure serialization, consistency, correct security, and safety. For simplicity, no metadata locks or leases are issued to clients. For HPC workloads in particular, callbacks offer minimal upside at a high potential cost in complexity.

Instead, Ceph optimizes for the most common metadata access scenarios. A *readdir* followed by a *stat* of each file (*e. g. ls -l*) is an extremely common access pattern and notorious performance killer in large directories. A *readdir* in Ceph requires only a single MDS request, which fetches the entire directory, including inode contents. By default, if a *readdir* is immediately followed by one or more *stats*, the briefly cached information is returned; otherwise it is discarded. Although this relaxes coherence slightly in that an intervening inode modification may go unnoticed, Ceph can optionally make this trade for vastly improved performance. This behavior is explicitly captured by the *readdirplus* [103] extension, which returns *lstat* results with directory entries (as some OS-specific implementations of *getdir* already do).

Ceph can allow consistency to be further relaxed by caching metadata longer, much like earlier versions of NFS, which typically cache for 30 seconds. However, this approach breaks coherency in a way that is often critical to applications, such as those using *stat* to determine if a file has been updated—they either behave incorrectly, or end up waiting for old cached values to time out.

Ceph opts instead to again provide correct behavior and extend the interface in in-

stances where it adversely affects performance. This choice is most clearly illustrated by a *stat* operation on a file currently opened by multiple clients for writing. In order to return a correct file size and modification time, the MDS revokes any write capabilities to momentarily stop updates and collect up-to-date size and mtime values from all writers. The highest values are returned with the *stat* reply, and capabilities are reissued to allow further progress. Although stopping multiple writers may seem drastic, it is necessary to ensure proper serializability. (For a single writer, a correct value can be retrieved from the writing client without interrupting progress.) Applications who find coherent behavior unnecessary—victims of a POSIX interface that doesn't align with their needs—can opt to use the new *statlite* operation [103], which takes a bit mask specifying which inode fields are not required to be coherent.

3.3 Communication Model

Ceph adopts an asynchronous messaging model for inter-node communication. In contrast to systems based on RPC, outgoing messages are queued for later delivery without blocking, and exchanges need not (and generally do not) consist of request and response pairs. This allows for greater flexibility and efficiency in data and information flow by facilitating asymmetric exchanges.

The communications model assumes ordered fail-stop delivery of messages between any given pair of nodes. That is, all sent messages will be delivered in their entirety in the order they were sent. The implementation currently utilizes TCP, although the interface is constructed to easily support alternative transports such as RDS [70]. In the event of a communications

failure (*e. g.* a transport error, such as a break in an underlying TCP connection), the sender is asynchronously notified of the error.

Chapter 4

Distributed Metadata Management

As demand for storage has increased, the centralization inherent in client-server storage architectures has proven a significant obstacle to scalable performance. Recent distributed file systems have adopted architectures based on object-based or brick-based storage, distributing low-level block allocation decisions to the storage devices and simplifying file system metadata. This lends itself well to architectures in which clients interact with a separate metadata server (MDS) to perform metadata operations (*e. g. open, rename*) while communicating directly with storage devices (OSDs) for file I/O.

In systems that decouple data and metadata access, efficient metadata performance becomes critical to overall system performance. Prior workload study has shown that metadata operations account for as much as 30-70% of all file systems operations [77]. While conventional write buffering and data prefetching techniques reduce interaction with storage devices for data I/O, metadata operations are frequently synchronous. Further, because data is managed independently from metadata, the metadata server workload consists almost entirely of small

reads and writes to relatively small data structures, making it important for the MDS to optimize its own I/O with underlying storage devices.

As distributed file systems scale to petabytes of data and beyond, the distribution of workload becomes difficult. System capacity and workloads usually grow incrementally over time, and can vary independently, making static partitions of the file system to individual servers inefficient: individual nodes may become overloaded while others sit idle. Typical server workloads vary on a daily basis, as demand shifts throughout the day [27], and in many systems—compute clusters in particular—workloads may vary wildly as different jobs start and finish [98]. Such variance in file system workloads demands *adaptive* approaches to metadata management in order to effectively utilize resources at scale [25, 79, 100, 102].

The combined demand for metadata performance and dynamic load balancing poses a significant challenge to system reliability. Storage systems, by their nature, demand strong data safety, while distributed systems introduce more system components that may fail. Most conventional file systems employ some form of *journal*, which provides a sequential log of operations that allows fast recovery after a failure. In a dedicated MDS, however, the large number of small updates leads to a pathological I/O profile when combined with conventional metadata storage and journaling approaches. Alternative techniques based on *soft updates* require careful ordering of updates [64], limiting parallelism and placing an even greater burden on the storage subsystem.

I present a metadata management approach addressing the unique performance requirements of a clustered metadata server that is capable of tolerating arbitrary node crashes¹.

¹The MDS cluster is tolerant of any number or combination of MDS node crashes, provided there is sufficient opportunity for new MDS processes to recover in their place. The cluster does not tolerate Byzantine failures,

In contrast to previous work in this area, my architecture and implementation maintain metadata performance before, during, and after failure by simultaneously addressing the efficiency of metadata I/O, cluster adaptability, and failure recovery.

Ceph’s metadata storage differs from that in conventional file systems in two key ways. First, the relatively small per-file *inode* metadata structures in our environment (due in part to our use of an object-based storage layer that minimizes allocation metadata) make it practical to embed them in directories with the directory entries (*dentries*) that refer to them, while support for multiple pathnames referencing a single file (hard links) is preserved through an auxiliary table. Based on prior study and my own analysis of file system snapshots and workloads, Ceph stores each directory’s contents (*dentries* and embedded *inodes*) together in the shared object storage pool. This facilitates *inode* prefetching by exploiting the high degree of directory locality present in most workloads—with negligible additional I/O cost—and facilitates an efficient subtree-based load partitioning strategy by embedding all metadata in the hierarchical namespace.

Second, Ceph utilizes per-metadata server update journals that are allowed to grow very large, to hundreds of megabytes or more. This allows it to distill the effects of multiple updates and short-lived files into a much smaller set of updates to the primary metadata structures, minimizing inefficient random-access I/O, while facilitating streamlined recovery and subsequent MDS performance after a node failure.

Finally, I adopt a subtree-based partitioning strategy [25, 102] to dynamically distribute workload. A hierarchical approach preserves locality within each MDS workload, while

wherein processes behave incorrectly or maliciously.

selective metadata replication preserves cluster throughput and availability when a subset of nodes fail. Ceph’s approach redefines the directory hierarchy in terms of directory *fragments*, facilitating fine-grained load distribution and simple, efficient storage.

In this chapter I focus on the **failure recovery**, **metadata I/O efficiency**, and **adaptability** implications of the combined approach to metadata storage, journaling, and workload distribution and relate my experiences constructing a working implementation. I analyze a variety of static file system snapshots and workload traces to motivate MDS design and performance analysis, present a simulation based-analysis of metadata partitioning approaches to demonstrate the architectural advantages of a dynamic subtree-based approach, and evaluate my implementation under a range of micro-benchmarks, workload traces, and failure scenarios.

4.1 Background and Related Work

Metadata in file systems presenting a POSIX file system interface is normally represented by three basic structures: inodes, dentries, and directories. Per-file or directory metadata (*e. g.* modification time, size, data location) is stored in *inodes*. Each *directory* has some number of file names or directory entries (*dentries*), each referencing an inode by number.

4.1.1 Local File Systems

In the original Unix file system, as well as BSD’s FFS [65]—whose basic design most later file systems have largely preserved—inodes were stored in tables in reserved regions on disk, their precise location indicated by their inode number. In contrast, C-FFS [29] embeds

most inodes inside directories with the dentries that refer to them, as well as taking steps to keep each directory’s file data together on disk. Despite significant performance benefits (10-300%), few file systems since have adopted similar embedding techniques. This may be due to C-FFS’s need to maintain a map of all directories to support *multilink files* (usually termed *hard links*)—inodes which are referenced by multiple filenames. In contrast, our metadata server architecture embeds *all* inodes and maintains a much smaller table that indexes only those directories with nested multilink files (we consider table utilization in detail in Section 4.5.2.3).

The Log-Structured File System (LFS) [26, 78] stores all data and metadata in a sequential log for improved write performance. However, the “cleaner” used to reclaim space—by rewriting partially deallocated log segments to new regions of disk—can incur a significant overhead for many workloads [87, 88]. DualFS [73] separates data and metadata management—much like object-based cluster file systems—by writing metadata to a separate device using an LFS-style log, mitigating cleaner effects by eliminating bulky file data. Although—as with LFS—a special inode map file is required to translate inode numbers to locations on disk, DualFS improves metadata locality by storing recently modified dentries and inodes close together in the log. hFS [113] combines elements of FFS and LFS by storing metadata and small file data in an LFS-like log and large file data in FFS-style allocation groups, allowing the effects of many operations to be combined into an atomic segment write (as in LFS) while also avoiding the need for a cleaner entirely.

Most modern file systems [92, 95] employ *journaling* to maintain file system consistency after a crash. Metadata (and, in certain cases, data) updates are first written to a log-structured journal before being applied to the regular on-disk metadata structures. Although

journals come with a performance penalty—disks have to frequently reposition themselves to append to the journal, and each metadata update is written to disk twice—they avoid the need for more costly consistency checks during failure recovery (which are increasingly expensive as file systems scale). Many systems also implement a form of *soft updates*, in which modifications are written only to unallocated regions of disk and in careful order to ensure that the on-disk image is always consistent (or easily repaired) [40, 64]. Ceph’s MDS combines elements of LFS and journaling, adapted to a distributed environment.

4.1.2 Distributed File Systems

Distributed file systems must partition management of the overall file system namespace across multiple servers in order to scale. I discuss prior approaches in terms of the partitioning strategy to introduce the efficiency issues involved.

4.1.2.1 Static Subtree Partitioning

The most common approach is to statically partition the directory hierarchy and assign each subtree to a particular server, in some cases migrating subtrees when it becomes necessary to correct load imbalance—this approach is taken by Sprite [71], StorageTank [66], PanFS [68], and others. However, because partition boundaries are typically visible through the lack of *link* and atomic *rename* support, these systems resemble *ad hoc* collections of NFS or CIFS [38] servers. Static partitions fail to account for the growth or contraction of individual subtrees over time, often requiring intervention of system administrators to repartition or (in some cases) manually rebalance data across servers. Although StorageTank and PanFS volumes or file sets

are smaller than the physical server capacity and can be migrated between servers, the unit of load distribution remains coarse and fixed.

4.1.2.2 Hash-based Partitioning

To limit workload imbalance associated with static subtree partitioning, Lustre’s clustered MDS distributes directories randomly among “inode groups,” each of which is assigned to a server [14]. This provides a finer distribution of data and workload by decorrelating the partition from the hierarchy. A number of other distributed file systems use a similarly random distribution: Vesta [20], Intermezzo [13], RAMA [67], and zFS [76] all hash the file pathname and/or some other unique identifier to determine the location of metadata and/or data. As long as such a mapping is well defined, this simple strategy has a number of advantages. Clients can locate and contact the responsible MDS directly and, for average workloads and well-behaved hash functions, requests are evenly distributed across the cluster. Further, hot-spots of activity in the hierarchical directory structure, such as heavy create activity in a single directory, do not correlate to individual metadata servers because metadata location has no relation to the directory hierarchy. However, hot-spots consisting of individual files can still overwhelm a single responsible MDS.

More significantly, distributing metadata by hashing eliminates all hierarchical locality, and with it many of the locality benefits typical of local file systems. Some systems distribute metadata based on a hash of the *directory* portion of a path only to allow directory contents to be grouped on MDS nodes and on disk. This approach facilitates prefetching and other methods of exploiting locality within the metadata workload. Even so, to satisfy POSIX

directory access semantics, the MDS cluster must traverse prefix (ancestor) directories containing a requested piece of metadata to ensure that the directory permissions allow the current user access to the metadata and data in question. Because the files and directories located on each MDS are scattered throughout the directory hierarchy, a hashed metadata distribution results in high overhead, either from a traversal of metadata scattered on multiple servers, or from the cache of prefixes replicated locally. Prefix caches between nodes will exhibit a high degree of overlap because parent directory inodes must be replicated for each MDS serving one or more of their children, consuming memory resources that could cache other data.

Approaches like ANU (adaptive nonuniform randomization) [108] seek to distribute fixed file sets (*e. g.* Lustre’s inode groups) among servers using a hash-like function that is dynamically adjusted based on measured server load. As with static subtree partitions, ANU relies on an existing fixed partition of metadata into file sets; in contrast, Ceph’s MDS balances load using a flexible dynamic partition and heuristics similar to those in ANU to prevent load thrashing.

4.1.2.3 Lazy Hybrid

Lazy Hybrid (LH) metadata management [15] seeks to capitalize on the benefits of a hashed distribution while avoiding the problems associated with path traversal by merging the net effect of the permission check into each file metadata record. Like other hashing approaches, LH uses a hash of the file’s full path name to distribute metadata. To alleviate the potentially high cost of traversing paths scattered across the cluster, LH uses a dual-entry access control list that stores the effective access information for the entire path traversal with the metadata

for each file. It has been shown that this information can usually be represented very compactly even for large general-purpose file systems [74]. LH need only traverse the path when access controls need to be updated because an ancestor directory's access permissions are changed, affecting the effective permissions of all files nested beneath it. Similarly, renaming or moving a directory affects the path name hash output and hence metadata location of all files nested beneath it, requiring metadata to be migrated between MDSs. Previous trace analysis has shown that changes like this happen very infrequently [77] and it is likely that they will affect small numbers of files when they do occur. Moreover, it is possible to perform this update at a later time to avoid a sudden burst of network activity between metadata servers, by having each MDS maintain a log of recent updates that have not fully propagated and then lazily update nested items as they are requested.

LH avoids path traversal in most cases, provided certain metadata operations are sufficiently infrequent in the workload. Analysis has shown that update cost can be amortized to one network trip per affected file; as long as updates are eventually applied more quickly than they are created (changes to directories containing lots of items could trigger potentially millions of updates with a single update), LH delivers a net savings and good scalability. Like other file hashing approaches, it avoids overloading a single MDS in the presence of directory hot-spots by scattering directories. However, in doing so the locality benefits are lost while the system remains vulnerable to individually popular files. More importantly, the low update overhead essential to LH performance is predicated on the low prevalence of specific metadata operations, which may not hold for all workloads.

4.1.2.4 Dynamic Subtree Partitioning

A few research file systems have recently adopted *dynamic* subtree-based metadata partitions, in which subtrees of the hierarchy are redistributed across a cluster of servers in response to the current workload. The directory service in Farsite [25] dynamically redistributes subtrees described by a hierarchical file identifier, but unlike Ceph, implements Windows (non-POSIX) semantics, with no support for hard links. More significantly, Farsite also relies on a hierarchical file identifier in order to minimize growth of its description of the current metadata distribution; in contrast, Ceph requires no global state to specify the namespace partition, improving scalability. Envoy [79] allows client hosts to claim management of subtrees based on their current workload through a distributed locking strategy. Neither, however, evaluates the comparative benefits of a subtree-based approach quantitatively, nor do either present or evaluate an implementation with working failure recovery. The Farsite MDS utilizes an independent atomic-action state machine substrate, which the authors suggest could be integrated with Farsite's existing Byzantine-fault tolerant substrate; the performance implications of doing so are not discussed. A recovery procedure for Envoy is described, but not evaluated nor considered in the context of system performance.

4.1.2.5 Metadata Storage

Distributed approaches to metadata management commonly utilize a shared storage subsystem, facilitating load balancing and failure recovery. However, existing approaches take minimal steps to optimize the I/O generated by the metadata server. Because Lustre metadata in each inode group is stored in a modified ext3 file system, the large volume of updates imposes a

significant random I/O load due to the use of conventional journaling and inode storage. PanFS and Envoy, two other object-based file systems, both store each file’s metadata as attributes on the file data objects, subjecting the distributed object store to a similar workload.

I describe a combined metadata management strategy involving embedded inodes, large MDS journals, and dynamic subtree partitioning specifically in the context of metadata I/O efficiency, metadata partition efficiency and adaptability, and sustained performance in the presence of arbitrary node failure. Ceph’s evaluation is based on the study of general purpose file system usage—content snapshots and workload traces—and evaluation of both normal operation and failure recovery of a working implementation.

4.2 Metadata Storage

Metadata, like data, is stored by a cluster of object storage devices (OSDs) that form a replicated, reliable object storage pool, described in Chapter 6. A variety of other reliable, distributed storage layers have been discussed in the literature [1, 30, 82, 100, 107]; in this chapter I ignore the details of object storage, assuming only that performance is based on the behavior of modern hard disks—characterized by expensive seeks and high sequential throughput—and focus instead on minimizing the I/O workload generated by the MDS by favoring sequential over small, random I/Os.

The full metadata contents of each directory (file names and inodes) are stored in a single object whose name corresponds to the directory’s inode number, as shown in Figure 4.1. Each MDS also maintains a journal that contains recently created or modified metadata not yet




Figure 4.1: All metadata for each directory, including file names (dentries) and the inodes they reference, are stored in a single object (identified by the directory's inode number) in the shared, distributed object store. Inode 1 is the root directory.

committed to the directory storage objects. MDS journals are striped over many fixed-size (but large) objects with sequential identifiers. Journals grow by spilling onto new objects, and are trimmed by deleting older objects.

Normally, the MDS submits a separate I/O operation for each journal entry in order to minimize the latency associated with any individual update operation. Under heavy loads, however, journal throughput can still be limited by the storage layer due to the large number of small (though sequential) write operations. To avoid this, the MDS will write journal data to the object store in larger increments only when it is necessary to keep the overall I/O rate below a fixed (tunable) maximum, improving overall throughput (and, correspondingly, average latency) under heavy load.

4.2.1 Embedded Inodes

File path resolution is performed by loading the directory inode and then dentry for each path component to determine the file’s inode number. Although most file systems try to keep inodes near the directory entries that reference them, multiple read operations are typically required. This can be problematic for many common usage patterns, such as a *readdir* followed by a *stat* on every file (as with `ls -l` or `find`), as each *stat* needs data at a different location on disk. Moreover, *stat* operations typically occur in the order dentries appear in the directory, which rarely corresponds to the order inodes occur on disk.

The key reason that inodes and dentries exist as independent structures is that POSIX file systems allow multiple dentries (and thus, multiple file names) to refer to a single inode (file). In practice, however, multiple “hard links” to individual files are relatively rare, and are used primarily for temporary files. That is, inodes rarely have more than one link for very long (see Section 4.5.2.3).

The efficiency problems associated by scattered inodes are avoided by storing dentries and inodes together. I call one (usually the only) dentry referring to each inode the *primary dentry*, and store the inode adjacent to it, inside the directory object. When a directory is loaded off disk to perform a single lookup or a *readdir*, the dentries and inodes for the entire directory are fetched into the MDS cache in a single read operation. Given that analysis of file system workloads has shown a high degree of directory locality [77, 89], our streamlined inode prefetching results in an improved I/O profile with minimal cost; the increased directory size does not significantly effect overall access time (see Section 4.5.2.2), and prefetched metadata

are placed low in the cache LRU list to avoid displacing other metadata in the working set.

4.2.2 Remote Links and the Anchor Table

Any additional dentries that reference an inode are called *remote dentries*, and are stored with the inode number for the inode they refer to, as with dentries in conventional file systems. However, such links may exist anywhere in the file system hierarchy relative to the primary dentry, whose location may also change. I introduce an auxiliary *anchor table* which logically takes the place of a conventional inode table by allowing inodes to be located by their inode number. However, we populate the table by only *anchoring* those inodes with more than one link (*i. e.* only those inodes we may need to locate by number).

To anchor an inode, all ancestor inodes must either already exist or be added to the anchor table, which consists of small, fixed size records (*anchors*) of the form $\langle ino, parent, nref \rangle$ —backpointers that eventually link each anchored inode back to the root. *Ino* identifies the inode in question, *parent* identifies its immediate parent directory by inode number and fragment id (see Section 4.3.6), and *nref* is a reference count that includes the inode anchor and any parent pointers from other anchor table records. Thus, given the inode number of an anchored inode we can retrieve the path by successively looking up parents until a known inode (*i. e.* one that an MDS already has in its in-memory cache) or the root directory is reached.

A key property of an anchor table consisting of directory backpointers is that a directory rename that may effect the path for an unbounded number of anchored inodes (*e. g.* near the root of the hierarchy) requires only a single anchor table update transaction. That transaction will include the backpointer change for the renamed directory, a reference count decrement

on the old parent, and count increment on the new parent. It may also include the subsequent removal of ancestors whose counts have reached zero, or the insertion of the new parent and its ancestors if they were not already present in the table.

4.2.3 Large Journals

In most file systems, the journal acts as a temporary staging area to facilitate safe updates of often complicated metadata structures without fear of corruption. Journal contents are typically discarded as soon as those updates complete. In contrast, log-based file systems such as LFS [26, 78] and hFS [113] use the log as the primary storage structure, citing its superior write performance. To manage deletion and deallocation, log-based systems—with notable exception of hFS—typically introduce a “cleaner” process to rewrite partially deallocated segments so that disk space can be reused.

Ceph’s metadata manager takes a hybrid approach. Updates are first written to an MDS journal, and the affected metadata is marked “dirty” and pinned in the MDS cache. Although we eventually commit the change to the primary per-directory metadata objects, we delay this until the relevant entry must be trimmed from the tail of the journal, and allow the journal to become very large (hundreds of megabytes or more). Most significantly, each journal’s contents are reflected by the MDS’s in-core state; changes are later committed to the primary metadata structures without requiring the journal to be re-read (as with LFS’s cleaner). MDS journals can be viewed as a means of recovering the contents of each node’s in-memory metadata cache—and indirectly the file system state—thereby allowing the MDS to adapt conventional techniques like delayed writeback and group commit to a clustered setting for im-

proved performance without compromising safety.

Large journals reduce the number of directory updates in two key ways. First, in most workloads, individual metadata objects are updated multiple times before becoming idle. Similarly, most new files are temporary and are removed shortly after creation [27, 77]. As the journal is trimmed and changes are committed to the directory objects, all but the most recent update for any given object can be ignored; for deleted files, all journal entries are moot. Second, all updates to a given directory over the lifetime of the journal are effectively committed—and the corresponding in-memory copies marked “clean”—in a single update transaction.

This strategy allows the MDS to optimize its I/O profile by streaming most updates to disk in an efficient, sequential manner while also limiting the number of random updates to per-directory metadata objects. At the same time, unlike logging file systems like LFS and hFS, we also optimize future read access by grouping dentry and inode metadata by directory.

Aside from enabling failure recovery, large journals allow a recovering MDS to prime its cache with a large quantity of warm metadata. This avoids an inefficient start from a cold cache and the large number of read I/Os that would be required to equivalently re-populate it, preserving MDS performance after recovery.

4.3 Adaptive Workload Distribution

As file systems scale it becomes necessary to distribute workload over multiple machines in order to achieve acceptable performance. Most clustered file systems use a static subtree partition, in which fixed portions of the file system hierarchy are assigned to different

servers. The problem with static partitions is that neither file system contents nor workloads are static: both the storage needs for and client access to different parts of the hierarchy will grow or contract over time, often requiring the intervention of system administrators to repartition or re-balance file systems across devices or servers. Further, transient variations in workload can lead to wide variations in server load, preventing the effective utilization of available hardware resources.

Ceph capitalizes on the benefits of a *dynamic* subtree-based partition to achieve scalable performance. Subtrees of the hierarchy are adaptively migrated between nodes to correct load imbalance, while favoring a coarse partition in order to preserve locality within the workload managed by each MDS. This typically allows client nodes to interact with only a small set of servers, while also limiting the amount of ancestor metadata that must be replicated to support consistent path resolution and access control.

Farsite and Envoy also utilize a subtree-based partition and dynamic load distribution [25, 79], although neither provides a comparative evaluation of dynamic partitioning versus other approaches. More importantly, none of the prior research in this area focuses on the implications of MDS node failure on metadata partitioning (or vice-versa). In this section, I focus on the construction of a distributed metadata cache that preserves implementation simplicity while facilitating scalability *and* tolerating arbitrary MDS node crashes.

4.3.1 Subtree Partitioning as Soft State

Ceph’s architecture allows an arbitrary and adaptive subtree-based partition of the file system across a dynamic cluster of metadata servers, while providing a single semantic

namespace (*i. e.* full support for atomic rename and hard links, even when spanning subtrees currently managed by different MDS nodes). At the same time, Ceph allows extremely fine-grained load distribution when necessary. Instead of defining subtrees in terms of directories (which makes directories the indivisible unit of load distribution), Ceph allows large or busy directories to fragment, and defines the partition in terms of individual directory fragments (see Section 4.3.6). Each subtree is defined by the directory fragment at its root, and zero or more *bounds*, where each bound is the root directory fragment for a nested subtree. At any given subtree boundary, the directory inode belongs to the parent subtree (it is stored in the containing directory), while the directory fragment—and the metadata it contains—belongs to the child subtree.

A distinguished MDS (*mds0*) is always responsible for the root inode. Beyond that, metadata is partitioned in terms of metadata that currently exists in the cluster’s collective in-memory cache. Because the cached subset is a connected subtree of the overall hierarchy, this simultaneously partitions the total file system. Any given piece of metadata (directory fragment, dentry, or inode) occupies a specific position within the file system hierarchy, by virtue of Ceph’s embedded inodes (see Section 4.2.1). This allows a subtree partition to simultaneously partition both namespace (directory fragment and dentry) and file (inode) metadata. Every metadata item present in an MDS’s in-memory cache is either *authoritative*—if it falls within a subtree of the namespace that is locally managed—or a *replica* of an item managed by another MDS.

In contrast to many other distributed architectures [14, 25, 44, 68, 71], there is no central management or description of the overall partition. The partition of cached metadata




Figure 4.2: Structure of the metadata cache as seen by a single MDS. Metadata is partitioned into subtrees bounded by directory fragments. Each MDS replicates ancestor metadata for any locally managed subtrees. Large or busy directories are broken into multiple fragments, which can then form nested subtrees.

is “soft state” in that it is described only by the cluster’s collective in-memory cache (and, by extension, journal contents). Each MDS is aware only of the authority for metadata present in its own cache. Clients cache subtree boundaries as they encounter them to efficiently direct metadata requests.

4.3.2 Metadata Replication

Metadata replication serves two key purposes. First, it allows each node to cache a connected hierarchy, which is important for maintaining cache consistency. Metadata is also replicated for availability—both when the cluster is under heavy load and when a subset of nodes fail.

Ceph imposes four *distributed cache constraints* that are critical for failure recovery and simplify normal cache operations: (1) all cached metadata must be attached to the root of the hierarchy (*i. e.* all ancestors must also be cached); (2) any directory fragment bounds for an authoritative subtree, *i. e.* the root directory fragments of any child sub-trees, though not necessary they dendries they contain, must have local replicas; (3) the authority for any replicated metadata must be known; and (4) all replicas must be known by the authority.

Constraints (1) and (2) ensure that path resolution can start at any node, and that an MDS resolving a path will know when traversal reaches a subtree boundary. Replicating ancestor metadata further ensures that, for example, the failure of *mds0*—which manages the root directory—will not prevent path traversal on other nodes managing nested subtrees. Similarly, popular directories are flagged for replication, allowing most read-only requests to be serviced even if the authoritative MDS has failed.

Constraints (2) and (3) ensure that it will be clear to which MDS we should delegate if path traversal reaches the edge of an authoritative subtree: the bounding directory fragment replica will indicate the authority for the nested subtree. Note that replicating a directory fragment does not imply replication of its contents (dentries and inodes)—only that the replica will remain informed of the fragment’s authority. Finally, constraint (4) facilitates locking when operations affect replicated metadata, and ensures authoritative metadata can be expired from the cache after replicas are destroyed.

4.3.3 Locking

Each piece of replicated metadata is protected by a lock—a simple distributed state machine that controls whether an MDS can read or modify a given set of fields. Locking is fine-grained: although each dentry is protected by a single lock (controlling whether that element of the namespace can be read), each inode has five locks, each controlling a different set of related fields (*e. g.* link count and anchor status; file ownership and mode; file size and *mtime*; directory fragmentation). Lock acquisition is ordered by $\langle locktype, object \rangle$ to avoid deadlock.

Each lock’s state machine is constructed to minimize MDS interaction for the expected usage of the protected fields. Most fields are protected by a simple lock that keeps replicas consistent and readable by default, but allows an exclusive write lock on the authority for updates. File size and *mtime*, on the other hand, are protected by a lock with states corresponding to modes of client file access: single client, shared read, or mixed read/write or shared write. A “scatter” lock regulates the *mtime* field for directory inodes when fragment(s) are managed by different MDSs than the inode: concurrent updates are allowed unless the lock

is moved to a state that combines values and allows *mtime* to be read.

Ordinarily all update operations are forwarded to the metadata object authority for serialization and journaling. A few operations (*link*, *unlink*, and *rename*) affect multiple metadata objects that may be managed by different MDS nodes through *slave updates*. For example, if *link* is creating a hard link in a directory on the current MDS that refers to an inode on a different MDS, it will issue a slave request to increment the link count. Slaved updates are applied using a two-phase commit protocol: once all slaves have journaled a “prepare” event, the coordinating MDS journals the update (effectively committing the transaction), and slaves journal matching “committed” events to close the transaction state.

4.3.4 Load Balance

Each MDS monitors the popularity of cached metadata through counters associated with each inode and directory fragment. Each inode popularity vector includes a read and a write counter, while directory fragments additionally monitor *readdir* operations and the frequency that metadata is fetched from or committed to the object store. In addition to its own popularity, each directory maintains three additional load vectors that form summations over metadata nested deeper within the hierarchy: one for all nested metadata, one for all nested metadata for which the current node is authoritative, and one for authoritative metadata within the current subtree only.

When an MDS services client requests, the appropriate counters are incremented on the affected metadata and its ancestors in order to provide a hierarchical view of metadata popularity, which is in turn used to inform replication and migration decisions. An exponential

decay factor is applied when counters are read, providing a smoothed approximation of recent popularity.

Nodes in the MDS cluster periodically share their overall load level, as well as the cumulative popularity of locally managed metadata grouped by the authoritative MDS of the immediate ancestor for each subtree. This allows each overloaded node to determine the fraction of its overall workload to shed to underloaded nodes, while favoring migrations that reunify child subtrees with their parents. The hierarchical popularity accounting is then used to select appropriately sized subtrees to migrate.

While the popularity counters measure the frequency of metadata access, a final counter for each directory measures popularity *spread*. Each directory has a short list of the clients to last access its contents. The spread counter is incremented only when a new client's request is processed. Spread helps inform metadata replication decisions for availability, as demonstrated in Section 4.5.7.

4.3.5 Subtree Migration

Each subtree migration is a transfer of all cached metadata for the subtree. This exchange includes non-dirty metadata (that, strictly speaking, is not needed for correctness) because cache-worthiness is presumably unrelated to which MDS manages it, and the one-time cost during migration is significantly less expensive than re-fetching even a small subset of the data from the metadata store.

Migration involves a few initial message exchanges to set up the transfer, and then a two-phase commit: the “importing” node journals a copy of all imported metadata (Import-

Start), the “exporting” node commits the migration by journaling an Export event, and the importer journals an ImportFinish to close the transaction in its journal. Any “bystander” nodes who also replicate the root of the subtree are notified before the migration begins and after it completes, so that any cache expiration messages can be directed at both the old and new subtree authorities. This is necessary to ensure expirations are reliably delivered in the face of importer or exporter node failure, to preserve cache invariant (3). (A failure of both nodes is handled during the *resolve* phase; see Section 4.4.3.4.)

The relative ease with which subtrees are migrated is made possible by embedded inodes, which place all metadata within a single hierarchical namespace with a well-defined partition. Similarly, the use of a shared object store for metadata storage facilitates migration for arbitrarily large subtrees of the file system, as each migration involves only the transfer of *cached* metadata.

4.3.6 Directory Fragments

Metadata replication facilitates the distribution of read operations only, while subtrees are normally partitioned at the granularity of the directory hierarchy. Neither mechanism addresses load imbalance associated with individual directories that are extremely large or heavily updated (as in many high performance computing applications [98]). Moreover, our basic approach to directory metadata storage (each directory’s dentries and inodes written to a single object) does not scale well to large directories, nor does prefetching the entire directory always make sense in such a scenario.

In order to address both issues, Ceph extends the directory hierarchy to allow di-




Figure 4.3: Directories are fragmented based on a tree in which each interior node has a 2^n children, and leaves correspond to individual fragments. Fragments are described by a bit pattern and mask, like an IP subnet, partitioning an integer namespace. Dentries are assigned to a specific fragment using a hash functions.

rectory contents to be broken into multiple fragments. The one-to-many relationship between directory inodes and directory fragments is specified by a *frag tree* structure stored in the inode. An integer namespace is partitioned into one or more fragments based on a tree in which interior vertices split by powers of two, and the leaves are individual fragments, as illustrated in Figure 4.3. Each fragment is described by a bit mask (indicating which bits are significant) and a value for those bits, like an IP network and netmask. Directory entries are mapped into a particular directory fragment by hashing the file name and looking up the resulting value in the frag tree.

Each directory fragment's metadata are stored in a separate object, allowing extremely large directories to be stored efficiently. *Readdir* proceeds in fragment order and returns metadata in fragment-sized chunks, preserving prefetching performance for applications that touch every file in the directory. Because our subtree-based metadata partition is defined in terms of directory fragments, load can be finely delegated by migrating fragments between MDS nodes. Any fragment can be split into 2^n sub-fragments if it becomes large or busy. Conversely, fragments can be rejoined if load decreases or a directory shrinks. There is an I/O cost as-

sociated with fragment splits and merges, as both operations involve writing out the resulting per-directory fragment metadata objects.

4.3.7 Traffic Control

To effectively adapt to a changing workload, the MDS cluster must also cope with situations where a large number of clients access the same file or directory in the hierarchy at the same time, even suddenly and without warning. Extremely popular files and directories and sudden “flash crowds” are common in scientific computing workloads [98] (where large numbers of nodes may be acting in unison) and general purpose workloads where large numbers of users access similar files due to external events. If tens of thousands of clients access a single MDS simultaneously, that node will not be able to handle the request workload efficiently.

The fundamental problem is client knowledge of the metadata partition: if all clients know where to access any given piece of metadata at any time (based on a well-defined hashing strategy, for instance) then there is nothing to prevent them from simultaneously accessing the same item. Similarly, if clients are ignorant of the metadata distribution, then their requests must be directed randomly and forwarded within the MDS cluster, or pass through some sort of proxy, in either case requiring an extra network hop for all requests. Ideally, one would like a combination of the two situations: access to unpopular items to be directed at the authoritative MDS nodes, and access to popular items to be directed at many or all nodes (each replicating the popular metadata) to distribute traffic.

Ceph controls how client requests are directed by using clients’ initial ignorance of the metadata distribution to achieve near-ideal traffic flow for both popular and unpopular metadata.

All MDS responses sent to clients include current replication information—that is, which MDS nodes the client should contact in the future—for the metadata requested and their ancestors, which are then cached on the client. For unpopular items, the MDS cluster tells clients to direct future requests only at the authoritative node, while for popular items the client is told the item is replicated on many or all nodes. Because the popularity metric approximates the prevalence of an item in all client caches, the MDS cluster can effectively bound the number of nodes believing any particular file or subtree of the file hierarchy is located in any one place at all times, thus avoiding potential flash crowds before they can occur while still allowing most requests for unpopular data to be directed efficiently.

This strategy works for both explored and unexplored portions of the hierarchy. Because client requests are directed based on the deepest known prefix, any potential flood of requests will initiate from a set of mutually known (and thus popular) directories—in extreme cases, the root directory, which is known to all clients and consequently highly replicated.

4.4 Failure Recovery

Robust failure recovery is critical in distributed systems, where a large number of distinct pieces of hardware increases the likelihood of failure. Although our metadata architecture is based on journaling, a relatively standard mechanism for facilitating failure recovery, a number of factors complicate recovery. First, journals are large: entries near the tail may be significantly out of date. Second, the migration of metadata between servers introduces intervals of ambiguous authority that must be resolved during recovery. Third, our distributed cache

relies on a large amount of “soft” (non-journaled) state in order to function properly, including the identity of replicas and lock states. Because journaling such metadata would be exorbitantly expensive, it must be reliably reconstructed during failure recovery. Finally, open file state is associated with inodes—independent of the namespace—and shared only with clients, who are not always aware of metadata updates that may have relocated an open file’s position in the hierarchy.

4.4.1 Journal Structure

Each MDS maintains an independent journal containing a chronological sequence of atomic *events*, each of which typically includes some modified or contextual metadata. Ceph logically divides the event sequence into *segments*, and begins each segment with a special *subtree map* event that describes which subtrees of the hierarchy the MDS was authoritative for at that point in time.

Because inodes are embedded within the hierarchy, each journaled update must be accompanied by the modified metadata’s ancestors in order to locate the item within the hierarchy. To avoid duplication of contextual metadata in the journal, ancestors are journaled only up to the root of the containing subtree, and each ancestor is journaled only once per segment (unless it is subsequently modified). Between the subtree map and events that follow, each segment provides the necessary context to correctly interpret all metadata updates it contains.

Each piece of dirty metadata in the MDS cache is placed on a linked list for the segment it was most recently journaled in. Before old segments are trimmed from the tail of the journal, any directory fragments still referenced by the segment dirty list are committed.

New segments are created shortly after the journal extends onto a new storage object, such that as segments are trimmed the entire backing objects can be deleted to reclaim disk space. This allows each MDS to trim its own journal with minimal overhead by simply traversing the appropriate dirty list, in contrast to the cleaner in log-based file systems.

4.4.2 Failure Detection

Each active MDS sends regular beacon messages to a central monitor responsible for coordinating cluster membership. If an MDS does not check in for a sufficiently long interval, it is declared down and the status change is broadcast to surviving nodes. MDS processes who do not receive timely positive acknowledgement in response to their own beacons commit suicide.

4.4.3 Recovery

Each MDS consists entirely of a running process with a large in-memory cache, utilizing a distributed, shared object storage substrate. MDS failure does not result in data unavailability in the way that hardware failure in systems utilizing a local hard disk or NVRAM might. Once an MDS is declared down, a process running on any hardware node can be chosen to recover in its place.

Recovery is broken into four stages. The contents of the journal are first read into memory, ambiguous subtree authority and the fate of distributed transactions are resolved, client sessions are reestablished to restore open file state, and finally the node rejoins the cluster's distributed cache.

4.4.3.1 Replay

Journal *replay* begins with the first complete segment in the journal. The MDS simply reads all journal events in chronological order and adds any modified or contextual metadata to its in-memory cache. This serves to recover any dirty metadata that may not have been committed before the failure, as well as to prime the new MDS's cache.

Special note is made of events that indicate distributed transactions. For example, the MDS maintains a list of ImportStart events that are not followed by a matching ImportFinish, as indicates an ambiguous subtree import that may or may not have been committed by the exporting MDS. Similar message pairs reflect anchor table updates or slave update transactions made on behalf of *link*, *unlink*, or *rename* operations coordinated by other nodes.

4.4.3.2 Resolve

During the *resolve* phase, the fate of ambiguous transactions in the journal is determined. Each recovering node broadcasts a *resolve* message to all MDSs that includes a list of locally managed subtrees (described by the root and bound directory fragment identifiers), ambiguous subtrees that were mid-import at the time of failure, and any slave updates (see Section 4.3.3) initiated by the target MDS whose fate is unknown. Any surviving (non-failed) nodes in the cluster send similar resolve messages to each recovering MDS.

As each resolve is processed, the recovering MDS updates its own cache to reflect the authority of subtrees explicitly claimed by other nodes. Ambiguous slave updates are cross-checked against the local list of recently committed transactions, and a reply is generated to inform the recovering node. Once all resolve messages have been received, ambiguous imports

are resolved by simply checking if the subtree was unambiguously claimed by another node. (In the case of failure during subtree migration, only the importer will be unsure of the result; the exporter either did or did not write a commit to its journal.) Surviving bystanders also examine resolve messages to learn interrupted migration outcomes.

Finally, each recovering node trims all non-authoritative metadata from its cache that is not an ancestor of authoritative metadata (and thus required) or subtree bound. This is necessary because it is impossible to know whether any non-authoritative metadata was updated, moved, or even deleted at some point after it was mentioned in the local journal. To ensure the remaining replicated ancestor replicas are correct, all nodes replicating a renamed file or directory mention the event in their journal. This restores all distributed cache constraints except (4)—replica identity (see Section 4.3.2).

4.4.3.3 Reconnect and Open File State

Open file state, unlike other soft state, is shared not with other MDSs but with clients mounting the file system. Each recovering MDS reestablishes prior client sessions and queries clients for previously issued open file handles, which are necessary to recreate the corresponding capability and lock state in the MDS cache.

However, file opens are not synchronously journaled by the MDS, both to avoid the additional latency, and because most files are opened read-only (reliable *atime* updates are not deemed a priority). Instead, the MDS periodically writes recently opened inodes to the journal, and clients describe their capabilities by both inode number and last known filename. If an open file's inode was not recovered from the cache, the filename will allow it to be located within the

hierarchy. (Again, mentioning renamed replicas provides key information needed to reliably interpret paths for non-authoritative inodes.)

4.4.3.4 Rejoin

The final stage of recovery restores distributed cache and lock state. Recovering nodes send a *weak rejoin* message to each MDS, declaring any metadata replicas that the recipient is authoritative for, and any recovered file capabilities the sender is not authoritative for.

Surviving nodes (who have lost no state) send *strong rejoin* messages to recovering nodes, declaring any replicas that they hold, as well as asserting lock state. Any declared replicas that the recovering node does not already hold in its cache are obtained from the survivor: an item's absence in the journal implies it has not been modified and the replica is therefore up to date.

Recovering nodes use rejoin messages to initialize the lists of known replicas associated with each piece of metadata, and to choose initial lock states that are compatible with any surviving nodes. Recovered capabilities listed in weak rejoins that fall within locally managed subtrees are claimed and managed locally, with new client sessions established as necessary. Although migrating capability management is not strictly necessary (the other MDS clearly managed the capability just before the failure), reusing the capability migration mechanism in place for subtree migration simplified our implementation. Finally, a *rejoin ack* message is sent to initialize replica locks, and the recovered node becomes active.

4.4.4 Discussion

Any “dirty” metadata recovered from the journal will be restored to the in-memory cache. However, at least some of that metadata will likely have been committed by the prior MDS instance before it failed. Although version numbers are attached to all metadata objects to ensure correctness, some unnecessary I/O may result after recovery.

On the other hand, the use of a large journal provides a recovered node with the warmest subset of the node’s in-memory cache prior to failure. Consequently, simply killing an MDS process and restarting it is actually faster than if it were cleanly shut down. This durability resembles that in the Google File System’s MDS [30].

A limitation of this design (and, partly, my implementation) is that the resolve and rejoin recovery phases require the participation of all failed nodes. Although in certain cases this requirement can be relaxed, in general a multi-node failure requires parallel recovery.

4.5 Evaluation

The design of Ceph’s MDS architecture is evaluated in three stages. First, I evaluate the relative efficiency of static subtree, dynamic subtree, and hash-based approaches in a simulation environment to demonstrate the advantages of an dynamic, adaptive approach. I next evaluate the performance of Ceph’s MDS, by analyzing static file system snapshots and captured file system traces in order to better understand metadata workloads and their impact on performance. Finally, I measure the performance of the implementation with a range of micro-benchmarks and capture workloads under normal-use and failure scenarios.

4.5.1 Metadata Partitioning

I evaluate the relative performance and scalability of a dynamic subtree-based partition compared to alternative approaches: namely, a static subtree partition (*e. g.* a collection of NFS servers), directory-based hashing (*e. g.* Lustre), file-based hashing, and Lazy Hybrid [15], a hybrid hash-based approach proposed by Brandt *et al.* . This comparative evaluation is performed in an event-driven simulation environment.

Initially, I fix the amount of MDS memory per node and scale the entire system: file system size, number of MDS servers, and client base. Figure 4.4 shows the performance degradation of individual MDS nodes for different system sizes under a predominately static (file system and client) workload. Dynamic and static subtree partitioning show the best performance, the only difference between the two being that the static strategy does not employ load balancing to adjust the initial partition. In a real workload environment, a static partition is unlikely to be practical as file systems and workloads evolve over time and hierarchies are not typically as easily partitioned as our workload (a large collection of home directories). The apparent performance penalty for load balancing is due to an unfair distribution of metadata: some MDS nodes manage a small amount of metadata with extremely high efficiency while others have poor cache performance, resulting in a higher (though unfair) overall cluster throughput. More significantly, the performance of file and directory hashed distributions degrades more quickly than subtree based partitions due to inefficiencies analyzed in Section 4.5.1.1.

File hashing and lazy hybrid distributions show significantly lower performance due to inefficient metadata I/O operations, which involve disk requests to load individual inodes




Figure 4.4: MDS performance as file system, cluster size, and client base are scaled.

into cache. In contrast, the subtree and directory hashing partitioning strategies exploit the presence of locality in the workload by embedding inodes and storing entire directories together on disk to allow efficient lookups and prefetching. The benefits of this approach are best seen by contrasting the performance of the directory and file hashing strategies, which are otherwise identical.

Lazy Hybrid performance is interesting because it scales almost linearly due to its ability to avoid performing most path traversals under the evaluated workload. However, this ability is predicated on the rarity of modifications to the directory permissions and hierarchy which must be (lazily, but eventually) propagated to potentially large quantities of metadata.

4.5.1.1 Prefix Caching

The performance of metadata partitioning strategies is tightly linked to metadata cache efficiency. One of the primary factors affecting cache utilization is the need to cache prefix inodes of ancestor directories for the purposes of path traversal. The overhead associated




Figure 4.5: Percentage of cache devoted to ancestor inodes as the file system, client base and MDS cluster size scales. Hashed distributions devote large portions of their caches to ancestor directories. The dynamic subtree partition has slightly more ancestors than the static partition due to the re-delegation of subtrees nested within the hierarchy.

with caching ancestors for hashed partitions is particularly high because directories are scattered throughout the hierarchy and the prefix directory inodes to locate them must be replicated widely throughout the cluster. Figure 4.5 shows the percentage of MDS cache associated with ancestors as file system, client base and cluster size scale (as in Figure 4.4). The utilization for the static subtree partitioning represents a baseline for the file system simulated and is related to the ratio of directories to files, the average branching factor and average file depth. The dynamic subtree partition devotes slightly more cache to prefixes to anchor subtrees nested within the hierarchy that have been re-delegated to other MDS nodes to balance load.

The consumption of cache memory by ancestor metadata has the effect of decreasing the cache hit rate and thus overall MDS performance. The extent to which this affects performance is related to the average depth of directories in the hierarchy; obviously, a mostly flat namespace is more easily distributed—Lazy Hybrid tries to artificially flatten the namespace to achieve this effect. Ancestor caching overhead is also greater for smaller cache sizes both




Figure 4.6: Cache hit rate as a function of cache size (as a fraction of total file system size). For smaller caches, inefficient cache utilization due to replicated ancestors results in lower hit rates.

because memory is more scarce and because the demand for prefixes for path traversal is related to the distribution of requests throughout the file system, not just factors proportional to the size of the cache. Figure 4.6 shows how cache performance varies with the cache size, expressed as a fraction of the total size of the file system’s metadata. Note that the convergence of the hit rates as cache size increases is predicated on the degree of locality in the workload; a more random distribution of requests will result in a performance similar to smaller cache sizes.

4.5.1.2 Traffic Control

One of the key advantages of a dynamic partitioning strategy is the ability to manage client ignorance to prevent simultaneous access by tens of thousands of users from overwhelming an individual metadata server. Figure 4.7 shows the number of requests processed over time by individual nodes in the simulated MDS cluster when 10,000 clients simultaneously request the same file, a scenario typical of many scientific computing workloads. Requests are directed

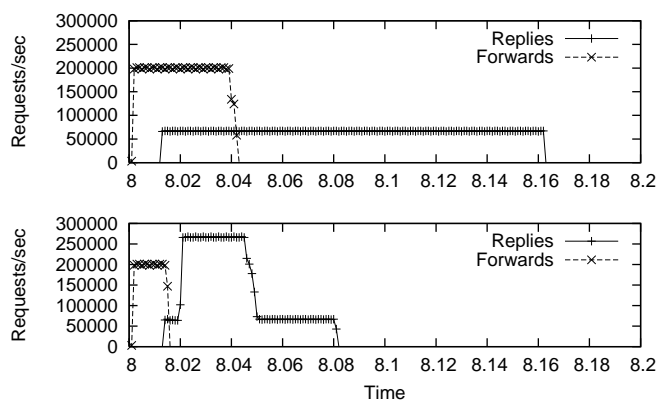


Figure 4.7: No traffic control (top): nodes forward all requests to the authoritative MDS who slowly responds to them in sequence. Traffic control (bottom): the authoritative node quickly replicates the popular item and all nodes respond to requests.

randomly because clients do not already know which MDS node is responsible for the file. Without traffic control (top), MDS nodes simply forward requests to the authoritative node who is quickly saturated and slowly (and, in real situations, inefficiently) responds. When traffic control is enabled (bottom), the authority quickly recognizes the file's sudden popularity and replicates the metadata on other nodes.

The response time from when the flash crowd begins until it is effectively distributed across the cluster is dependent on a number of factors, including the replication threshold, the rate at which client requests can be received and then forwarded by MDS nodes, and the latency of I/O requests that may be required to load the requested metadata into the cache. This response time could be reduced if non-authoritative MDS nodes recognized the sudden flood of requests and preemptively cached the metadata being requested without waiting to be told to do so, or if the authoritative node noticed the flood of requests before waiting for the metadata to be loaded from disk.

4.5.2 Embedded Inodes

I next consider the performance implications of a metadata storage strategy in which inodes are embedded inside directories. I analyze a range of static file system snapshots, including a shared university engineering department file server (*soe.**) [74], an aged (10 year) web and email server file system (*destro*), a software development server (*yakko*), and a collection of file systems of varying vintages (0-8 year) for a commercial web (and email) host (*web.**). I measure the performance of a fully functional implementation with a range of micro-benchmarks and a 2.5 hour web and email workload trace (*destro*).

This set of experiments primarily reflects general purpose—as opposed to scientific computing—file system workloads at scale, mainly because paired file system snapshot and workload data from high performance computing systems was not readily available to me. MDS performance under synthetic workloads based on common scientific computing workload patterns (as described by Wang *et al.* [98]) is considered in Sections 4.5.4 and 4.5.5.

4.5.2.1 Directory Sizes

I begin by looking at typical directory sizes, and show that the increased size of directories due to embedded inodes has little impact on performance compared to the overhead of fetching inodes separately. Figure 4.8 shows the cumulative directory size distributions for a range of file system snapshots and the Linux kernel source (which I use for certain benchmarks). More than 95% of directories contain less than 100 entries, and would occupy less than 20 KB with embedded inodes. This allows for files averaging 20 characters and 160 byte inodes in place of 64-bit inode numbers—a 6-fold increase in storage. However, for small I/Os,




Figure 4.8: Cumulative directory size for a variety of static file system snapshots. Many directories are very small (or even empty), and less than 5% contain more than 100 entries.

positioning delays dominate. For example, a random seek followed by a single 4 KB read takes about 13 ms on a commodity 500 GB 7200rpm SATA disk. Reading 20 KB is not significantly slower, and more than 512 KB (a 2500 file directory) can be read in twice that time. If the seek is less than 1 MB past the previous read, 256 KB (1300 files) can be read in twice the time of 4 KB.²

4.5.2.2 Metadata Prefetching

The effectiveness of loading entire directories of metadata into cache in a single I/O is evaluated by measuring the rate at which inodes are accessed for the first time (*e. g.* due to a *stat* or *open*) relative to the rate directories are loaded. For Linux kernel compilations—the default and a random configuration—I found that 25-70% of loaded inodes were subsequently

²These measurements were made on a Western Digital WD5000YS-01M under Linux 2.6.21 with `O_DIRECT`. The long seek test measures the time to read n 4 KB blocks of data at a random block-aligned disk offset, averaged over 1000 iterations. The short seek test first reads a random 4 KB block and then measures the time to read n blocks of data located a short distance (between zero and 1 MB) past that.




Figure 4.9: Study of hard link prevalence in file system snapshots. For each snapshot, the total file system size and the percentage of files with multiple links are shown. Stacked bars indicate the fraction of multilink files with all links in the same directory, with parallel cohorts linked by multiple directories, or with no obvious locality properties.

used (out of 500 directory loads). Replaying the destro file system trace indicated only a 1.3% hit rate (out of 8000 directory loads), primarily due to large Maildir mailboxes. However, for 44% of all directory loads, at least one inode was used, for 18% at least two were, and for 11% at least five were; as seen above, avoiding each additional I/O is comparable to loading an additional 500-2000 embedded inodes. Most importantly, our strategy eliminates the risk of an I/O storm loading thousands of individual inodes if all files are accessed (*e. g.* as with an email search or `ls -l`).

4.5.2.3 Multilink File Prevalence and Locality

I analyze a range of file system snapshots to determine the prevalence and distribution of hard links. First, in Figure 4.9, we observe that multilinked files are rarely found in typical file systems. When they do occur, all links often appear in the same directory. Moreover, when multilink files do span directories, they usually have cohorts linked in parallel into the same pair(s) of directories (as from `cp -lr`, which creates a duplicate directory tree that links instead of copies the original files). Anecdotally, I found that in each of the analyzed file systems with high rates of hard links, a handful of users were responsible.

4.5.2.4 Anchor Table Performance

When all links for a file occur in the same directory, there is no significant performance impact on reads: the remote dentry and inode are loaded into the cache simultaneously, avoiding any anchor table query. For link creation, an additional transaction occurs against the anchor table and existing inode, increasing latency by 150% to 230% (depending on whether the inode and anchor table are on the same MDS as the new link) in our gigabit ethernet-based environment.

To measure the impact of the parallel link use-case, I consider the Linux 2.6.22 kernel source and a duplicate tree with links (created with `cp -lr`). Starting from a cold cache, a recursive walk (and `stat`) of the original tree took 13.4 ± 1.1 seconds, while the copy took $17.9 \pm .9$ seconds—only 34% longer. The limited impact is due to the fact that only one anchor table query is needed for each directory: once its contents have been loaded into the cache for the first hard link, any subsequent parallel links can also be resolved. Consequently, as seen in

Figure 4.9, a relatively small fraction of all multilink files require significant interaction with the anchor table, limiting its impact on performance.

4.5.2.5 Hard Link Lifetime

Multilink files are rarely encountered in part because they are primarily used for manipulating temporary files. The average time interval during which an inode has a link count greater than one in the *destro* trace is only a few milliseconds. Ceph’s current implementation naïvely keeps the entire anchor table in RAM on a designated MDS, such that 1 GB would store roughly 20 million records, enough for 1–25 billion files (40 TB–1 PB of 40 KB files), given the rates in Figure 4.9. However, because measured workloads rarely require table queries, and most updates are short-lived and never read, a more sophisticated table management approach should scale well.

4.5.3 Journaling

I next consider the performance implications of large journals by analyzing the effect of journaling on metadata I/O under a range of workloads.

4.5.3.1 Journal Interval

Large journals optimize metadata I/O by masking multiple metadata updates to the same objects and combining all dirty data for a directory into a single commit I/O. For example, for a Linux kernel *untar*, approximately 30% of journaled entries in each segment are obsolete by the time they are trimmed (due to the *mknod*, *utime*, and *size/mtime* flush sequence for each

file).

More importantly, each directory's dirty metadata is committed in a single I/O, resulting in significantly fewer directory commits than metadata updates. The extent to which we are able to group updates into a single commit is related to the time interval covered by the journal, and the type of directory locality in the workload. For a kernel compilation workload, the commit rate (relative to the total number of metadata updates), drops from 10% to 3% as the journal interval grows from 1 to 10 seconds, while for the *destro* workload it drops to 5% only after 5 minutes, showing a surprisingly long locality interval. Journal entries in our workloads averaged 300-800 bytes each³; a 200 MB journal would contain on the order of 250,000 to 600,000 entries, a two to five minute interval for an MDS doing 2,000 updates per second.

4.5.3.2 Ancestor Metadata

I measure the overhead of journaling contextual ancestor metadata, which are needed to place updates in the hierarchy during journal replay. Although naïvely including contextual metadata in each entry bloats the journal by 100-200% (for *untar* and compilation workloads), including any given ancestor just once per 1 MB segment increased the journal size by only .05%-.2%. That said, journal size is not terribly significant, as sequential I/O is easy to scale. A typical MDS would only stream to the journal at a few megabytes per second, an order of magnitude slower than a single disk, producing a significantly lighter I/O workload than thousands of updates to hundreds of individual objects.

³Journal entries typically contain multiple inodes, dentries, and other metadata associated with the operation (such as inode number allocation). For example, a file creation journal event includes the inode and dentry for both the new file and its parent directory (whose *mtime* changed). No particular attempt is made to efficiently represent this data, as journal bandwidth—in terms of bytes per second—is relatively low.




Figure 4.10: An MDS cluster adapts to a varying workload. At times 200 and 350 the workload shifts, as 400 clients begin creating files in private directories and then in a shared directory. In both cases load is redistributed, in the latter case after the large (and busy) directory is fragmented. At time 280 two additional MDSs join the cluster.

4.5.4 Adaptive Distribution

Figure 4.10 demonstrates the MDS cluster’s ability to adapt to varying file and directory creation workloads. At time 200, the cluster workload shifts as 400 clients begin creating files in private directories off of the root, all initially managed by *mds0*. After a few seconds the workload is redistributed, and at time 270 two additional servers join the cluster. At time 350, the workload shifts again when all clients begin creating files in the *same* directory. A few seconds later the (now large) directory is fragmented and load is again redistributed. The directory is fragmented a second time around time 380, resulting in a dip in throughput as the new resulting fragments are flushed to disk. Other cluster-wide dips and peaks in throughput are primarily due to contention in the shared storage pool, underscoring the importance of efficient




Figure 4.11: Per-MDS throughput under a variety of workloads and cluster sizes. As the cluster grows to 128 nodes, efficiency drops no more than 50% below perfect linear (horizontal) scaling for most workloads, allowing vastly improved performance over existing systems.

I/O.

4.5.5 Metadata Scaling

I measure the scalability of the MDS cluster using a 430 node partition of the a1c Linux cluster at Lawrence Livermore National Laboratory (LLNL). Figure 4.11 shows per-MDS throughput (y) as a function of MDS cluster size (x), such that a horizontal line represents perfect linear scaling. In the *makedirs* workload, each client creates a tree of nested directories four levels deep, with ten files and subdirectories in each directory. Average MDS throughput drops from 2000 ops per MDS per second with a small cluster, to about 1000 ops per MDS per second (50% efficiency) with 128 MDSs (over 100,000 ops/sec total). In the *makefiles* workload, each client creates thousands of files in the same directory. When the high write levels are detected, fragments the shared directory and relaxes the directory's *mtime* coherence

to distribute the workload across all MDS nodes. The *openshared* workload demonstrates read sharing by having each client repeatedly open and close ten shared files. In the *openssh* workloads, each client replays a captured file system trace of a compilation in a private directory. One variant uses a shared `/lib` for moderate sharing, while the other shares `/usr/include`, which is very heavily read. The *openshared* and *openssh+include* workloads have the heaviest read sharing and show the worst scaling behavior, I believe due to poor replica selection by clients.⁴ *openssh+lib* scales better than the trivially separable *makedirs* because it contains relatively few metadata modifications and little sharing. Although I believe that contention in the network or threading in the messaging layer further lowered performance for larger MDS clusters, my limited time with dedicated access to the large cluster prevented a more thorough investigation.

Figure 4.12 plots latency (y) versus per-MDS throughput (x) for a 4-, 16-, and 64-node MDS cluster under the *makedirs* workload. Larger clusters have imperfect load distributions, resulting in lower average per-MDS throughput (but, of course, much higher total throughput) and slightly higher latencies.

Despite imperfect linear scaling, a 128-node MDS cluster running the Ceph prototype can service more than a quarter million metadata operations per second (128 nodes at 2000 ops/sec). Because metadata transactions are independent of data I/O and metadata size is independent of file size, this corresponds to installations with potentially many hundreds of petabytes of storage or more, depending on average file size. For example, scientific applications creating checkpoints on LLNL's BlueGene/L might involve 64 thousand nodes with two

⁴Due to limited time with access to the large cluster, this hypothesis was not tested.




Figure 4.12: Average latency versus per-MDS throughput for different cluster sizes (*makedirs* workload).

processors each writing to separate files in the same directory (as in the *makefiles* workload). While the current storage system peaks at 6,000 metadata ops/sec and would take minutes to complete each checkpoint, a 128-node MDS cluster could finish in two seconds. If each file were only 10 MB (quite small by HPC standards) and OSDs sustain 50 MB/sec, such a cluster could write 1.25 TB/sec, saturating at least 25,000 OSDs (50,000 with replication). 250 GB OSDs would put such a system at more than six petabytes. More importantly, dynamic metadata distribution allows an MDS cluster (of any size) to reallocate resources based on the current workload, even when all clients access metadata previously assigned to a single MDS, making it significantly more versatile and adaptable than any static partitioning strategy.

4.5.6 Failure Recovery

Recovery from multiple MDS failures is demonstrated in Figure 4.13 in a four node cluster under a file and directory creation workload. At time 400 *mds3* fails (the process is

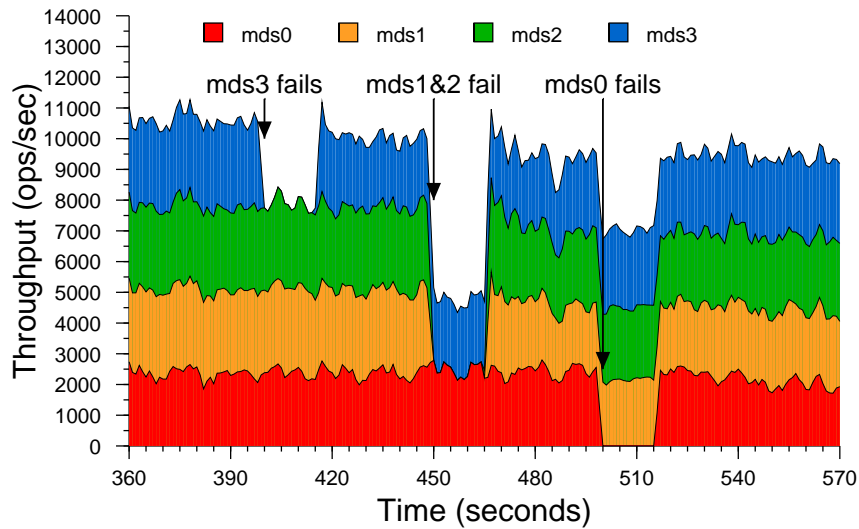


Figure 4.13: Throughput in a small cluster before, during, and after MDS failures at time 400, 500, and two at 450. Unresponsive nodes are declared dead after 15 seconds, and in each case recovery for a 100 MB journal takes 4-6 seconds.

killed) and is soon replaced. Similarly, two nodes fail at time 450 and another at 500. In each case, there is a 15 second delay before the non-responsive node is declared dead, after which point the total recovery time is 4 to 6 seconds, 3 to 4 seconds of which is spent replaying a 100 MB journal. Individual node failures do not disrupt throughput for the rest of the cluster due to a hierarchical partition of localized client workloads.

In this experiment, each recovered MDS immediately functions at full efficiency because there are no caching effects under the create-only workload. In general, nodes will recover warm metadata from the journal, including any recently modified metadata or recently opened files. Under the kernel compilation workload, I measured metadata recovery at a rate of .45 inodes per journal entry (1200 inodes per megabyte).

4.5.6.1 Recovery Time

Recovery time is governed by a number of factors, although in most cases it is dominated by failure detection and journal replay. I arbitrarily chose 15 seconds for the configurable delay before non-responsive nodes are declared dead. Although this period dominates down time in the experiment above, it should be sufficiently long to avoid failing due to transient network disruptions.

Replay is limited by the rate that the journal is sequentially read from the object store—only a few seconds for the 100 MB journal above (partly because it is still cached by the object storage layer), or 20 seconds for a 500 MB journal at 25 MB/second. Large journals prolong downtime but prime the recovering node’s cache with warm metadata, resulting in improved performance once they do become available.

The resolve phase involves exchanging small messages with every other MDS, and is thus dominated by cluster size. For example, a 10-node cluster with a degenerate partition of thousands of subtrees still resolves in less than 300 ms. In extremely large clusters, such an exchange may become significant, although such clusters are also typically equipped with high-performance networks.

The duration of the client reconnect phase is governed by the number of clients in the system, but is bounded by a timeout to avoid waiting for unresponsive clients. This timeout should balance timely recovery versus the risk of stale file handles on laggy clients. I found that a reconnect over 1.7 million open files among 1000 client processes (spread across 36 hosts) took ten seconds.




Figure 4.14: Per-MDS throughput in a 4-node cluster under a compilation workload with heavy read sharing of `/lib`. Replication of popular metadata allows progress to continue even when the MDS managing the heavily shared directory suffers a failure.

Finally, the rejoin phase is related to the number of MDS nodes and the amount of metadata replicated by surviving nodes. As with `resolve`, the exchange of small messages is fast for all but the largest clusters: for highly separable workloads, rejoin consumes only a few hundred milliseconds. In a 10-node metadata cluster with extensive metadata replication, a recovering node spent 6 seconds rejoining: strong rejoin messages, each 5 MB, were processed from 9 surviving peers, each replicating 30,000 inodes. A final message fetched metadata for inodes not present in the journal from a survivor.

4.5.7 Availability

Ancestor metadata replication allows nested subtrees to remain available by facilitating path resolution, as demonstrated in Figure 4.13 by the failure of `mds0` (who manages the

root directory) at time 500. Adaptive replication of popular directories further limits the impact of a node failure: Figure 4.14 shows per-MDS throughput for a four node cluster under a workload of 40 clients, each replaying an openssh compilation trace with heavy read sharing of `/lib`.⁵ By replicating that directory’s contents across the cluster (in response to its popularity and read sharing), the failure of its authoritative MDS does not affect other nodes’ workloads. Clients direct read-only metadata requests (*e. g. stat, open* for read) at replicas when a directory is flagged for replication, improving availability (as in this experiment) and load distribution for read/write workloads in which metadata updates must be handled by the authoritative node.

4.6 Future Work

Although the MDS implementation utilizes an object-based storage layer, the approach is enabled primarily by small inode structures. The extent to which it is applicable to other storage architectures—including block-based storage—has not been considered in detail.

A number of auxiliary metadata structures are managed in a naïve fashion. The anchor table, for instance, is kept entirely in RAM, as is a table of recently client requests (for avoiding repetition when recovering from communication or node failure) and a handful of similar structures. I am investigating alternative data structures and update strategies to serve the same purpose.

I plan to apply simple heuristics to the management of multilink files such that inodes are embedded in the directories through which they are most frequently referenced.

⁵MDS throughputs are higher here than in previous experiments due to a workload consisting primarily of non-update operations.

4.7 Experiences

The MDS design was vastly simplified by separating it from the object storage layer, avoiding concerns about low-level replication, consistency, and data placement. Similarly, despite my original expectation, the use of a large, lazily trimmed journal in combination with versioning had a simplifying effect on the implementation of the MDS. Although “dirty” state has to be associated with the journal segment it was last written to, lazy journal trimming eliminates most timing constraints on when in-memory state must be committed to the primary metadata store.

In my initial design, individual directories could be flagged to distribute metadata using a hash function across all nodes in the cluster, much like some experimental file systems [13, 20, 67, 76]. This led to a large number of corner cases, particularly relating to the subtree partition, and did not mesh well with an architecture in which MDS nodes could be dynamically added or removed from the cluster. By introducing directory fragments and defining the subtree partition in terms of that abstraction, Ceph simultaneously achieves both fine-grained load balancing and addresses storage for very large directories, while adding minimal complexity to the subtree migration infrastructure.

As my distributed metadata cache design and implementation evolved to accommodate arbitrary failure, the importance of setting straightforward invariants became clear—particularly with respect to the subtree migration infrastructure. Simple rules like “a replica must know the identity of the authoritative copy” made it clear how to directly address—or avoid—corner cases as they arose.

One of the largest lessons in Ceph was the importance of the MDS load balancer to overall scalability, and the complexity of choosing what metadata to migrate where and when. Although in principle the design and goals seem quite simple, the reality of distributing an evolving workload over a hundred MDSs highlighted additional subtleties. Most notably, MDS performance has a wide range of performance bounds, including CPU, memory (and cache efficiency), and network or I/O limitations, any of which may limit performance at any point in time. Although I experimented with a variety of MDS load functions (including combinations of metadata popularity, request rates, and queue lengths), in many cases simply using the CPU load average works as well as anything else. Furthermore, it is difficult to quantitatively capture the balance between total throughput and fairness; under certain circumstances unbalanced metadata distributions can increase overall throughput.

4.8 Conclusions

I describe a clustered metadata server that optimizes I/O to the underlying storage system, adapts its metadata distribution to the current workload, and tolerates arbitrary node failure. The MDS embeds inodes inside directories for efficient metadata storage and prefetching, facilitating an efficient and adaptive hierarchical partition of workload. Metadata updates are first written to large per-MDS journals, which aggregate many updates to the same directory into a single commit to the primary metadata structures. This reduces the I/O load on the underlying storage system and improves MDS performance after recovery from a failure by priming its cache. We adaptively partition workload based a hierarchy defined in terms of directory

fragments, facilitating fine-grained load balancing and simple, efficient storage for large directories. Finally, ancestor and popular metadata is replicated across multiple nodes for improved availability—both when the cluster is under load, and when a subset of nodes fail.

I consider the merits of a dynamic subtree-based partitioning strategy relative to alternative approaches in a simulation environment. The metadata storage strategy is evaluated by analyzing a range of static file system snapshots and workloads. I demonstrate the vastly reduced I/O workload resulting from large journals that exploit workload locality, and consider the prevalence of hard links and their impact on the performance of our embedded inode strategy. Finally, I demonstrate the adaptive workload distribution and failure recovery in the implementation in terms of performance and metadata availability.

Chapter 5

Data Distribution

Object-based storage is an emerging architecture that promises improved manageability, scalability, and performance [7]. Unlike conventional block-based hard drives, object-based storage devices (OSDs) manage disk block allocation internally, exposing an interface that allows others to read and write to variably-sized, named objects. In such a system, each file's data is typically striped across a relatively small number of named objects distributed throughout the storage cluster. Objects are replicated across multiple devices (or employ some other data redundancy scheme) in order to protect against data loss in the presence of failures. Object-based storage systems simplify data layout by replacing large block lists with small object lists and distributing the low-level block allocation problem. Although this vastly improves scalability by reducing file allocation metadata and complexity, the fundamental task of distributing data among thousands of storage devices—typically with varying capacities and performance characteristics—remains.

Most systems simply write new data to underutilized devices. The fundamental prob-

lem with this approach is that data is rarely, if ever, moved once it is written. Even a perfect distribution will become imbalanced when the storage system is expanded, because new disks either sit empty or contain only new data. Either old or new disks may be busy, depending on the system workload, but only the rarest of conditions will utilize both equally to take full advantage of available resources.

A robust solution is to distribute all data in a system randomly among available storage devices. This leads to a probabilistically balanced distribution and uniformly mixes old and new data together. When new storage is added, a random sample of existing data is migrated onto new storage devices to restore balance. This approach has the critical advantage that, on average, all devices will be similarly loaded, allowing the system to perform well under any potential workload [84]. Furthermore, in a large storage system, a single large file will be randomly distributed across a large set of available devices, providing a high level of parallelism and aggregate bandwidth. However, simple hash-based distribution fails to cope with changes in the number of devices, incurring a massive reshuffling of data. Further, existing randomized distribution schemes that decluster replication by spreading each disk's replicas across many other devices suffer from a high probability of data loss from coincident device failures.

I have developed CRUSH (Controlled Replication Under Scalable Hashing), a pseudo-random data distribution algorithm that efficiently and robustly distributes object replicas across a heterogeneous, structured storage cluster. CRUSH is implemented as a deterministic function that maps an input value—typically an object or object group identifier—to a list of devices on which to store object replicas. This differs from conventional approaches in that data placement does not rely on any sort of per-file or per-object directory—CRUSH needs only a compact, hi-

erarchical description of the devices comprising the storage cluster and knowledge of the replica placement policy. This approach has two key advantages: first, it is completely distributed such that any party in a large system can independently calculate the location of any object; and second, what little metadata is required is mostly static, changing only when devices are added or removed.

CRUSH is designed to optimally distribute data to utilize available resources, efficiently reorganize data when storage devices are added or removed, and enforce flexible constraints on object replica placement that maximize data safety in the presence of coincident or correlated hardware failures. A wide variety of data safety mechanisms are supported, including n -way replication (mirroring), RAID parity schemes or other forms of erasure coding, and hybrid approaches (*e. g.* RAID-10). These features make CRUSH ideally suited for managing object distribution in extremely large (multi-petabyte) storage systems where scalability, performance, and reliability are critically important.

5.1 Related Work

Object-based storage has recently garnered significant interest as a mechanism for improving the scalability of storage systems. A number of research and production file systems have adopted an object-based approach, including the seminal NASD file system [33], the Panasas file system [68], Lustre [14], and others [76, 30]. Other block-based distributed file systems like GPFS [86] and Federated Array of Bricks (FAB) [82] face a similar data distribution challenge. In these systems a semi-random or heuristic-based approach is used to allocate

new data to storage devices with available capacity, but data is rarely relocated to maintain a balanced distribution over time. More importantly, all of these systems locate data via some sort of metadata directory, while CRUSH relies instead on a compact cluster description and deterministic mapping function. This distinction is most significant when writing data, as systems utilizing CRUSH can calculate any new data’s storage target without consulting a central allocator. The Sorrento [93] storage system’s use of consistent hashing [47] most closely resembles CRUSH, but lacks support for controlled weighting of devices, a well-balanced distribution of data, and failure domains for improving data safety.

Although the data migration problem has been studied extensively in the context of systems with explicit allocation maps [5, 6], such approaches have heavy metadata requirements that functional approaches like CRUSH avoid. Choy, *et al.* [18] describe algorithms for distributing data over disks which move an optimal number of objects as disks are added, but do not support weighting, replication, or disk removal. Brinkmann, *et al.* [17] use hash functions to distribute data to a heterogeneous but static cluster. Brinkmann later describes an improved algorithm for placing replicas among a weighted set of disks in $O(n)$ time [16]. SCADDAR [34] addresses the addition and removal of storage, but only supports a constrained subset of replication strategies. None of these approaches include CRUSH’s flexibility or failure domains for improved reliability.

Brinkmann *et al.* also describe a problem with many placement algorithms—CRUSH included—in which data is imperfectly distributed in small clusters with heterogeneous device weights [16]. In such small clusters, the key properties that CRUSH provides are neither significant nor required (namely, scalability and support for replica separation).

ANU (adaptive, nonuniform randomization) [108] maps data objects and servers onto a sparse subset of the unit interval, and dynamically adjusts server allocations in response to observed load. This bears a strong resemblance to CRUSH’s overload mechanism, although CRUSH requires an initial weight estimate, and Ceph does not implement ANU’s heuristics for limiting load thrashing. However, like consistent hashing, ANU maps objects to individual servers with no support for replication, and further lacks CRUSH’s flexible placement rules, support for failure domains, or a tunable balance between performance and stability.

CRUSH most closely resembles the RUSH [41] family of algorithms upon which it is based. RUSH remains the only existing set of algorithms in the literature that utilizes a mapping function in place of explicit metadata and supports the efficient addition and removal of weighted devices. Despite these basic properties, a number of issues make RUSH an insufficient solution in practice. For example, like Brinkmann’s algorithm [16], $RUSH_R$ biases placement of certain replicas to certain devices, effectively providing a placement *set* instead of an ordered list, while $RUSH_P$ does not support the efficient removal of devices. CRUSH fully generalizes the useful elements of $RUSH_P$ and $RUSH_T$ while resolving previously unaddressed reliability and replication issues, and offering improved performance and flexibility.

5.2 The CRUSH algorithm

The CRUSH algorithm distributes data objects among storage devices according to a per-device weight value, approximating a uniform probability distribution. The distribution is controlled by a hierarchical *cluster map* representing the available storage resources and

composed of the logical elements from which it is built. For example, one might describe a large installation in terms of rows of server cabinets, cabinets filled with disk shelves, and shelves filled with storage devices. The data distribution policy is defined in terms of *placement rules* that specify how many replica targets are chosen from the cluster and what restrictions are imposed on replica placement. For example, one might specify that three mirrored replicas are to be placed on devices in different physical cabinets so that they do not share the same electrical circuit.

Given a single integer input value x , CRUSH will output an ordered list \vec{R} of n distinct storage targets. CRUSH utilizes a strong multi-input integer hash function whose inputs include x , making the mapping completely deterministic and independently calculable using only the cluster map, placement rules, and x . The distribution is pseudo-random in that there is no apparent correlation between the resulting output from similar inputs or in the items stored on any storage device. I say that CRUSH generates a *declustered* distribution of replicas in that the set of devices sharing replicas for one item also appears to be independent of all other items.

5.2.1 Hierarchical Cluster Map

The cluster map is composed of *devices* and *buckets*, both of which have numerical identifiers and weight values associated with them. Buckets can contain any number of devices or other buckets, allowing them to form interior nodes in a storage hierarchy in which devices are always at the leaves. Storage devices are assigned weights by the administrator to control the relative amount of data they are responsible for storing. Although a large system will likely contain devices with a variety of capacity and performance characteristics, randomized data

distributions statistically correlate device utilization with workload, such that device load is on average proportional to the amount of data stored. This results in a one-dimensional placement metric, weight, which should be derived from the device’s capabilities. Bucket weights are defined as the sum of the weights of the items they contain.

Buckets can be composed arbitrarily to construct a hierarchy representing available storage. For example, one might create a cluster map with “shelf” buckets at the lowest level to represent sets of identical devices as they are installed, and then combine shelves into “cabinet” buckets to group together shelves that are installed in the same rack. Cabinets might be further grouped into “row” or “room” buckets for a large system. Data is placed in the hierarchy by recursively selecting nested bucket items via a pseudo-random hash-like function. In contrast to conventional hashing techniques, in which any change in the number of target bins (devices) results in a massive reshuffling of bin contents, CRUSH is based on four different bucket types, each with a different selection algorithm to address data movement resulting from the addition or removal of devices and overall computational complexity.

5.2.2 Replica Placement

CRUSH is designed to distribute data uniformly among weighted devices to maintain a statistically balanced utilization of storage and device bandwidth resources. The placement of replicas on storage devices in the hierarchy can also have a critical effect on data safety. By reflecting the underlying physical organization of the installation, CRUSH can model—and thereby address—potential sources of correlated device failures. Typical sources include physical proximity, a shared power source, and a shared network. By encoding this information

into the cluster map, CRUSH placement policies can separate object replicas across different failure domains while still maintaining the desired distribution. For example, to address the possibility of concurrent failures, it may be desirable to ensure that data replicas are on devices in different shelves, racks, power supplies, controllers, and/or physical locations.

In order to accommodate the wide variety of scenarios in which CRUSH might be used, both in terms of data replication strategies and underlying hardware configurations, CRUSH defines *placement rules* for each replication strategy or distribution policy employed that allow the storage system or administrator to specify exactly how object replicas are placed. For example, one might have a rule selecting a pair of targets for 2-way mirroring, one for selecting three targets in two different data centers for 3-way mirroring, one for RAID-4 over six storage devices, and so on¹.

Each rule consists of a sequence of operations applied to the hierarchy in a simple execution environment, presented as pseudocode in Algorithm 1. The integer input to the CRUSH function, x , is typically an object name or other identifier, such as an identifier for a group of objects whose replicas will be placed on the same devices. The $take(a)$ operation selects an item (typically a bucket) within the storage hierarchy and assigns it to the vector \vec{i} , which serves as an input to subsequent operations. The $select(n, t)$ operation iterates over each element $i \in \vec{i}$, and chooses n distinct items of type t in the subtree rooted at that point. Storage devices have a known, fixed type, and each bucket in the system has a type field that is used to distinguish between classes of buckets (*e. g.* those representing “rows” and those representing “cabinets”). For each $i \in \vec{i}$, the $select(n, t)$ call iterates over the $r \in 1, \dots, n$ items requested and recursively

¹Although a wide variety of data redundancy mechanisms are possible, for simplicity I will refer to the data objects being stored as *replicas*, without any loss of generality.

Algorithm 1 CRUSH placement for object x

```
1: procedure TAKE( $a$ )                                ▷ Put item  $a$  in working vector  $\vec{i}$ 
2:    $\vec{i} \leftarrow [a]$ 
3: end procedure

4: procedure SELECT( $n, t$ )                             ▷ Select  $n$  items of type  $t$ 
5:    $\vec{o} \leftarrow \emptyset$                              ▷ Our output, initially empty
6:   for  $i \in \vec{i}$  do                                   ▷ Loop over input  $\vec{i}$ 
7:      $f \leftarrow 0$                                    ▷ No failures yet
8:     for  $r \leftarrow 1, n$  do                             ▷ Loop over  $n$  replicas
9:        $f_r \leftarrow 0$                                    ▷ No failures on this replica
10:       $retry\_descent \leftarrow false$ 
11:      repeat
12:         $b \leftarrow bucket(i)$                              ▷ Start descent at bucket  $i$ 
13:         $retry\_bucket \leftarrow false$ 
14:        repeat
15:          if “first  $n$ ” then                               ▷ See Section 5.2.2.2
16:             $r' \leftarrow r + f$ 
17:          else
18:             $r' \leftarrow r + f_r n$ 
19:          end if
20:           $o \leftarrow b.c(r', x)$                              ▷ See Section 5.2.4
21:          if  $type(o) \neq t$  then
22:             $b \leftarrow bucket(o)$                              ▷ Continue descent
23:             $retry\_bucket \leftarrow true$ 
24:          else if  $o \in \vec{o}$  or  $failed(o)$  or  $overload(o, x)$  then
25:             $f_r \leftarrow f_r + 1, f \leftarrow f + 1$ 
26:            if  $o \in \vec{o}$  and  $f_r < 3$  then
27:               $retry\_bucket \leftarrow true$                  ▷ Retry collisions locally (see Section 5.2.2.1)
28:            else
29:               $retry\_descent \leftarrow true$                  ▷ Otherwise retry descent from  $i$ 
30:            end if
31:          end if
32:        until  $\neg retry\_bucket$ 
33:      until  $\neg retry\_descent$ 
34:       $\vec{o} \leftarrow [\vec{o}, o]$                              ▷ Add  $o$  to output  $\vec{o}$ 
35:    end for
36:  end for
37:   $\vec{i} \leftarrow \vec{o}$                                        ▷ Copy output back into  $\vec{i}$ 
38: end procedure

39: procedure EMIT                                     ▷ Append working vector  $\vec{i}$  to result
40:    $\vec{R} \leftarrow [\vec{R}, \vec{i}]$ 
41: end procedure
```




Figure 5.1: A partial view of a four-level cluster map hierarchy consisting of rows, cabinets, and shelves of disks. Bold lines illustrate items selected by each *select* operation in the placement rule and fictitious mapping described by Table 5.1.

Action	Resulting \vec{i}
take(root)	root
select(1,row)	row2
select(3,cabinet)	cab21 cab23 cab24
select(1,disk)	disk2107 disk2313 disk2437
emit	

Table 5.1: A simple rule that distributes three replicas across three cabinets in the same row.

descends through any intermediate buckets, pseudo-randomly selecting a nested item in each bucket using the function $c(r, x)$ (defined for each kind of bucket in Section 5.2.4), until it finds an item of the requested type t . The resulting $n|\vec{i}|$ distinct items are placed back into the input \vec{i} and either form the input for a subsequent $select(n, t)$ or are moved into the result vector with an *emit* operation.

As an example, the rule defined in Table 5.1 begins at the root of the hierarchy in Figure 5.1 and with the first $select(1, row)$ chooses a single bucket of type “row” (it selects *row2*).

The subsequent *select(3,cabinet)* chooses three distinct cabinets nested beneath the previously selected *row2* (*cab21,cab23,cab24*), while the final *select(1,disk)* iterates over the three cabinet buckets in the input vector and chooses a single disk nested beneath each of them. The final result is three disks spread over three cabinets, but all in the same row. This approach thus allows replicas to be simultaneously separated across and constrained within container types (*e. g.* rows, cabinets, shelves), a useful property for both reliability and performance considerations. Rules consisting of multiple *take*, *emit* blocks allow storage targets to be explicitly drawn from different pools of storage, as might be expected in remote replication scenarios (in which one replica is stored at a remote site) or tiered installations (*e. g.* fast, near-line storage and slower, higher-capacity arrays).

5.2.2.1 Collisions, Failure, and Overload

The *select(n,t)* operation may traverse many levels of the storage hierarchy in order to locate n distinct items of the specified type t nested beneath its starting point, a recursive process partially parameterized by $r = 1, \dots, n$, the replica number being chosen. During this process, CRUSH may reject and reselect items using a modified input r' for three different reasons: if an item has already been selected in the current set (a collision—the *select(n,t)* result must be distinct), if a device is *failed*, or if a device is *overloaded*. Failed or overloaded devices are marked as such in the cluster map, but left in the hierarchy to avoid unnecessary shifting of data. CRUSH selectively diverts a fraction of an overloaded device's data by pseudo-randomly rejecting with the probability specified in the cluster map—typically related to its reported over-utilization. For failed or overloaded devices, CRUSH uniformly redistributes items across the

storage cluster by restarting the recursion at the beginning of the $select(n, t)$ (see Algorithm 1 line 11). In the case of collisions, an alternate r' is used first at inner levels of the recursion to attempt a local search (see Algorithm 1 line 14) and avoid skewing the overall data distribution away from subtrees where collisions are more probable (*e. g.* where buckets are smaller than n).

5.2.2.2 Replica Ranks

Parity and erasure coding schemes have slightly different placement requirements than replication. In primary copy replication schemes, it is often desirable after a failure for a previous replica target (that already has a copy of the data) to become the new primary. In such situations, CRUSH can use the “first n ” suitable targets by reselecting using $r' = r + f$, where f is the number of failed placement attempts by the current $select(n, t)$ (see Algorithm 1 line 16). With parity and erasure coding schemes, however, the rank or position of a storage device in the CRUSH output is critical because each target stores different bits of the data object. In particular, if a storage device fails, it should be replaced in CRUSH’s output list \vec{R} *in place*, such that other devices in the list retain the same rank (*i. e.* position in \vec{R} , see Figure 5.2). In such cases, CRUSH reselects using $r' = r + f_r n$, where f_r is the number of failed attempts on r , thus defining a sequence of candidates for each replica rank that are probabilistically independent of others’ failures. In contrast, RUSH has no special handling of failed devices; like other existing hashing distribution functions, it implicitly assumes the use of a “first n ” approach to skip over failed devices in the result, making it unwieldy for parity schemes.




Figure 5.2: Reselection behavior of $\text{select}(6, \text{disk})$ when device $r = 2$ (b) is rejected, where the boxes contain the CRUSH output \vec{R} of $n = 6$ devices numbered by rank. The left shows the “first n ” approach in which device ranks of existing devices (c, d, e, f) may shift. On the right, each rank has a probabilistically independent sequence of potential targets; here $f_r = 1$, and $r' = r + f_r n = 8$ (device h).

5.2.2.3 Force-feeding

In a variety of environments it is desirable to specify the placement of the first replica. For example, in distributed computing environments in which each server performs both computation and storage tasks, network utilization can be significantly lowered if the first replica is stored locally. To enable such behavior, CRUSH allows you to “force-feed” a specific first device to the placement algorithm, while still maintaining all other placement constraints imposed by the placement rule. (Note that although force feeding is relatively straightforward, it is not included in Algorithm 1.)

This is accomplished by consulting the device hierarchy that the placement rule is utilizing (*i. e.* that nested beneath the initial *take*), and inferring the choices that would be made by each *select* to ultimately choose the force fed item. Each *select* initially chooses the inferred or force fed item, and then proceeds pseudo-randomly for any additional results, producing a

distinct result list as before. This ensures that constraints imposed by the placement rule (*e. g.* separation of replicas within the hierarchy) are maintained, despite an explicit choice of the initial result.

Force feeding requires that each device only occurs once in the given subtree of the hierarchy in order to unambiguously determine the parent for each node, although the rest of the CRUSH algorithm has no such restriction. This would only be significant for unusual hierarchies not considered here.

5.2.3 Map Changes and Data Movement

A critical element of data distribution in a large file system is the response to the addition or removal of storage resources. CRUSH maintains a uniform distribution of data and workload at all times in order to avoid load asymmetries and the related underutilization of available resources. When an individual device fails, CRUSH flags the device but leaves it in the hierarchy, where it will be rejected and its contents uniformly redistributed by the placement algorithm (see Section 5.2.2.1). Such cluster map changes result in an optimal (minimum) fraction, w_{failed}/W (where W is the total weight of all devices), of total data to be remapped to new storage targets because only data on the failed device is moved.

The situation is more complex when the cluster hierarchy is modified, as with the addition or removal of storage resources. The CRUSH mapping process, which uses the cluster map as a weighted hierarchical decision tree, can result in additional data movement beyond the theoretical optimum of $\frac{\Delta w}{W}$. At each level of the hierarchy, when a shift in relative subtree weights alters the distribution, some data objects must move from subtrees with decreased




Figure 5.3: Data movement in a binary hierarchy due to a node addition and the subsequent weight changes.

weight to those with increased weight. Because the pseudo-random placement decision at each node in the hierarchy is statistically independent, data moving into a subtree is uniformly redistributed beneath that point, and does not necessarily get remapped to the leaf item ultimately responsible for the weight change. Only at subsequent (deeper) levels of the placement process does (often different) data get shifted to maintain the correct overall relative distributions. This general effect is illustrated in the case of a binary hierarchy in Figure 5.3.

The amount of data movement in a hierarchy has a lower bound of $\frac{\Delta w}{W}$, the fraction of data that would reside on a newly added device with weight Δw . Data movement increases with the height h of the hierarchy, with a conservative asymptotic upper bound of $h \frac{\Delta w}{W}$. The amount of movement approaches this upper bound when Δw is small relative to W , because data objects moving into a subtree at each step of the recursion have a very low probability of being mapped to an item with a small relative weight.

5.2.4 Bucket Types

Generally speaking, CRUSH is designed to reconcile two competing goals: efficiency and scalability of the mapping algorithm, and minimal data migration to restore a balanced dis-

Action	Uniform	List	Tree	Straw
Speed	$O(1)$	$O(n)$	$O(\log n)$	$O(n)$
Additions	poor	optimal	good	optimal
Removals	poor	poor	good	optimal

Table 5.2: Summary of mapping speed and data reorganization efficiency of different bucket types when items are added to or removed from a bucket.

tribution when the cluster changes due to the addition or removal of devices. To this end, CRUSH defines four different kinds of buckets to represent internal (non-leaf) nodes in the cluster hierarchy: *uniform buckets*, *list buckets*, *tree buckets*, and *straw buckets*. Each bucket type is based on a different internal data structure and utilizes a different function $c(r, x)$ for pseudo-randomly choosing nested items during the replica placement process, representing a different tradeoff between computation and reorganization efficiency. Uniform buckets are restricted in that they must contain items that are all of the same weight (much like a conventional hash-based distribution function), while the other bucket types can contain a mix of items with any combination of weights. These differences are summarized in Table 5.2.

5.2.4.1 Uniform Buckets

Devices are rarely added individually in a large system. Instead, new storage is typically deployed in blocks of identical devices, often as an additional shelf in a server rack or perhaps an entire cabinet. Devices reaching their end of life are often similarly decommissioned as a set (individual failures aside), making it natural to treat them as a unit. CRUSH uniform buckets are used to represent an identical set of devices in such circumstances. The key advantage in doing so is performance related: CRUSH can map replicas into uniform buckets in

constant time. In cases where the uniformity restrictions are not appropriate, other bucket types can be used.

Given a CRUSH input value of x and a replica number r , we choose an item from a uniform bucket of size m using the function $c(r, x) = (\text{hash}(x) + rp) \bmod m$, where p is a randomly (but deterministically) chosen prime number greater than m . For any $r \leq m$ we can show that we will always select a distinct item using a few simple number theory lemmas.² For $r > m$ this guarantee no longer holds, meaning two different replicas r with the same input x may resolve to the same item. In practice, this means nothing more than a non-zero probability of collisions and subsequent backtracking by the placement algorithm (see Section 5.2.2.1).

If the size of a uniform bucket changes, there is a complete reshuffling of data between devices, much like conventional hash-based distribution strategies.

5.2.4.2 List Buckets

List buckets structure their contents as a linked list, and can contain items with arbitrary weights. To place a replica, CRUSH begins at the head of the list with the most recently added item and compares its weight to the sum of all remaining items' weights. Depending on the value of $\text{hash}(x, r, \text{item})$, either the current item is chosen with the appropriate probability, or the process continues recursively down the list. This approach, derived from RUSH_P , recasts the placement question into that of “most recently added item, or older items?” This is a natural and intuitive choice for an expanding cluster: either an object is relocated to the newest device

²The Prime Number Theorem for Arithmetic Progressions [36] can be used to further show that this function will distribute replicas of object x in $m\phi(m)$ different arrangements, and that each arrangement is equally likely. $\phi(\cdot)$ is the Euler Totient function.

with some appropriate probability, or it remains on the older devices as before. The result is optimal data migration when items are added to the bucket. Items removed from the middle or tail of the list, however, can result in a significant amount of unnecessary movement, making list buckets most suitable for circumstances in which they never (or very rarely) shrink.

The $RUSH_P$ algorithm is approximately equivalent to a two-level CRUSH hierarchy consisting of a single list bucket containing many uniform buckets. Its fixed cluster representation precludes the use for placement rules or CRUSH failure domains for controlling data placement for enhanced reliability.

5.2.4.3 Tree Buckets

Like any linked list data structure, list buckets are efficient for small sets of items but may not be appropriate for large sets, where their $O(n)$ running time may be excessive. Tree buckets, derived from $RUSH_T$, address this problem by storing their items in a binary tree. This reduces the placement time to $O(\log n)$, making them suitable for managing much larger sets of devices or nested buckets. $RUSH_T$ is equivalent to a two-level CRUSH hierarchy consisting of a single tree bucket containing many uniform buckets.

Tree buckets are structured as a weighted binary search tree with items at the leaves. Each interior node knows the total weight of its left and right subtrees and is labeled according to a fixed strategy (described below). In order to select an item within a bucket, CRUSH starts at the root of the tree and calculates the hash of the input key x , replica number r , the bucket identifier, and the label at the current tree node (initially the root). The result is compared to the weight ratio of the left and right subtrees to decide which child node to visit next. This




Figure 5.4: Node labeling strategy used for the binary tree comprising each tree bucket.

process is repeated until a leaf node is reached, at which point the associated item in the bucket is chosen. Only $\log n$ hashes and node comparisons are needed to locate an item.

The bucket's binary tree nodes are labeled with binary values using a simple, fixed strategy designed to avoid label changes when the tree grows or shrinks. The leftmost leaf in the tree is always labeled "1." Each time the tree is expanded, the old root becomes the new root's left child, and the new root node is labeled with the old root's label shifted one bit to the left (1, 10, 100, etc.). The labels for the right side of the tree mirror those on the left side except with a "1" prepended to each value. A labeled binary tree with six leaves is shown in Figure 5.4. This strategy ensures that as new items are added to (or removed from) the bucket and the tree grows (or shrinks), the path taken through the binary tree for any existing leaf item only changes by adding (or removing) additional nodes at the root, at the beginning of the placement decision tree. Once an object is placed in a particular subtree, its final mapping will depend only on the weights and node labels within that subtree and will not change as long as that subtree's items remain fixed. Although the hierarchical decision tree introduces some additional data migration between nested items, this strategy keeps movement to a reasonable level, while offering efficient mapping even for very large buckets.

5.2.4.4 Straw Buckets

List and tree buckets are structured such that a limited number of hash values need to be calculated and compared to weights in order to select a bucket item. In doing so, they divide and conquer in a way that either gives certain items precedence (*e. g.* those at the beginning of a list) or obviates the need to consider entire subtrees of items at all. That improves the performance of the replica placement process, but can also introduce suboptimal reorganization behavior when the contents of a bucket change due an addition, removal, or re-weighting of an item.

The straw bucket type allows all items to fairly “compete” against each other for replica placement through a process analogous to a draw of straws. To place a replica, a straw of random length is drawn for each item in the bucket. The item with the longest straw wins. The length of each straw is initially a value in a fixed range, based on a hash of the CRUSH input x , replica number r , and bucket item i . Each straw length is scaled by a factor $f(w_i)$ based on the item’s weight so that heavily weighted items are more likely to win the draw, *i. e.* $c(r, x) = \max_i (f(w_i) \text{hash}(x, r, i))$. Although this process is almost twice as slow (on average) than a list bucket and even slower than a tree bucket (which scales logarithmically), straw buckets result in optimal data movement between nested items when modified.

The scaling factor $f(w_i)$ for each bucket item is precalculated when the bucket is first created or modified, using the iterative procedure in Algorithm 2³. Although the pseudocode is included here for completeness, its derivation was based on a combination of geometric analysis

³Note that the first portion of the pseudocode is simply creating a sorted mapping *reverse* by weight using an insertion sort.

and trial and error, and is neither concise nor elegant (I would not be surprised to learn that there is an equivalent and trivial closed form). The basic intuition is that items with identical weights will have the same straw scaling value. We sort by weight, and start by giving the least-weighted items a straw multiplier $f(w_i)$ of one. As each successively larger item weight group is “added,” its straw multiplier $f(w_i)$ is chosen based on the desired probability p_{below} of choosing a shorter straw versus the next longer straw, where we consider that items with larger straw multipliers may still result in shorter straw lengths. That is, w_{below} and w_{next} are a product item weight and the number of items with equal or larger weights, and $p_{below} = \frac{w_{below}}{w_{below} + w_{next}}$.

5.2.4.5 Bucket Discussion

The choice of bucket type can be guided based on expected cluster growth patterns to trade mapping function computation for data movement efficiency where it is appropriate to do so. When buckets are expected to be fixed (*e. g.* a shelf of identical disks), uniform buckets are fastest. If a bucket is only expected to expand, list buckets provide optimal data movement when new items are added at the head of the list. This allows CRUSH to divert exactly as much data to the new device as is appropriate, without any shuffle between other bucket items. The downside is $O(n)$ mapping speed and extra data movement when older items are removed or reweighted. In circumstances where removal is expected and reorganization efficiency is critical (*e. g.* near the root of the storage hierarchy), straw buckets provide optimal migration behavior between subtrees. Tree buckets are an all around compromise, providing excellent performance and decent reorganization efficiency.

Algorithm 2 Function to calculate straw scaling factor $f(w_i)$ from the bucket item weights w_i .

```

procedure CALC_STRAW_WEIGHTS( $\vec{w}, \vec{f}$ ) ▷ Calculate  $\vec{f}$  from  $\vec{w}$ 
   $size \leftarrow \text{length}(\vec{w})$ 
   $reverse[0] = 0$  ▷ Determine sort order of  $\vec{w}$  with an insertion sort.
  for  $i = 1$  to  $size - 1$  do
    for  $j = 0$  to  $i - 1$  do
      if  $w[i] < w[reverse[j]]$  then
        for  $k = i$  down to  $j + 1$  do ▷ Insert  $i$  here
           $reverse[k] = reverse[k - 1]$ 
        end for
         $reverse[j] = i$ 
        break
      end if
    end for
    if  $j = i$  then
       $reverse[i] = i$  ▷ Add  $i$  at end
    end if
  end for
   $straw \leftarrow 1$  ▷ Initial straw length is 1
   $numleft \leftarrow size$ 
   $wbelow \leftarrow 0$ 
   $lastw \leftarrow 0$ 
   $i \leftarrow 0$ 
  while  $i < size$  do
     $f[reverse[i]] \leftarrow straw$  ▷ Set this item's straw multiplier
     $i \leftarrow i + 1$ 
    if  $i = size$  then
      break
    end if
    if  $w[reverse[i]] \neq w[reverse[i - 1]]$  then ▷ Different weight than previous item?
       $wbelow \leftarrow (wbelow + w[reverse[i - 1]] - lastw) \times numleft$ 
      for  $j = i$  to  $size - 1$  do ▷ Adjust count of items with greater weight
        if  $w[reverse[j]] == w[reverse[i]]$  then
           $numleft \leftarrow numleft - 1$ 
        else
          break
        end if
      end for
       $wnext \leftarrow numleft \times (w[reverse[i]] - w[reverse[i - 1]])$ 
       $pbelow \leftarrow \frac{wbelow}{wbelow + wnext}$ 
       $straw \leftarrow straw \times \frac{1}{pbelow} \frac{1}{numleft}$ 
       $lastw = w[reverse[i - 1]]$ 
    end if
  end while
end procedure

```

5.3 Evaluation

CRUSH is based on a wide variety of design goals including a balanced, weighted distribution among heterogeneous storage devices, minimal data movement due to the addition or removal of storage (including individual disk failures), improved system reliability through the separation of replicas across failure domains, and a flexible cluster description and rule system for describing available storage and distributing data. I evaluate each of these behaviors under expected CRUSH configurations relative to $RUSH_P$ - and $RUSH_T$ -style clusters by simulating the allocation of objects to devices and examining the resulting distribution. $RUSH_P$ and $RUSH_T$ are generalized by a two-level CRUSH hierarchy with a single list or tree bucket (respectively) containing many uniform buckets. Although RUSH's fixed cluster representation precludes the use of placement rules or the separation of replicas across failure domains (which CRUSH uses to improve data safety), I consider its performance and data migration behavior.

5.3.1 Data Distribution

CRUSH's data distribution should appear random—uncorrelated to object identifiers x or storage targets—and result in a balanced distribution across devices with equal weight. I empirically measured the distribution of objects across devices contained in a variety of bucket types and compared the variance in device utilization to the binomial probability distribution, the theoretical behavior I would expect from a perfectly uniform random process. When distributing n objects with probability $p_i = \frac{w_i}{W}$ of placing each object on a given device i , the expected device utilization predicted by the corresponding binomial $b(n, p)$ is $\mu = np$ with a

standard deviation of $\sigma = \sqrt{np(1-p)}$. In a large system with many devices, we can approximate $1-p \simeq 1$ such that the standard deviation is $\sigma \simeq \sqrt{\mu}$ —that is, utilizations are most even when the number of data objects is large.⁴ As expected, I found that the CRUSH distribution consistently matched the mean and variance of a binomial for both homogeneous clusters and clusters with mixed device weights.

5.3.1.1 Overload Protection

Although CRUSH achieves good balancing (a low variance in device utilization) for large numbers of objects, as in any stochastic process this translates into a non-zero probability that the allocation on any particular device will be significantly larger than the mean. Unlike existing probabilistic mapping algorithms (including RUSH), CRUSH includes a per-device overload correction mechanism that can redistribute any fraction of a device’s data. This can be used to scale back a device’s allocation proportional to its over-utilization when it is in danger of overfilling, selectively “leveling off” overfilled devices. When distributing data over a 1000-device cluster at 99% capacity, I found that CRUSH mapping execution times increase by less than 20% despite overload adjustments on 47% of the devices, and that the variance decreased by a factor of four (as expected).

5.3.1.2 Variance and Partial Failure

Prior research [84] has shown that randomized data distribution offers real-world system performance comparable to (but slightly slower than) that of careful data striping. In my

⁴The binomial distribution is approximately Gaussian when there are many objects (*i. e.* when n is large).

Usage	Devices Overfilled	Devices Adjusted	Time	σ
50%	0	0	1.000	.100
70%	.7%	1.9%	1.003	.098
75%	1.6%	3.8%	1.016	.095
80%	2.9%	6.4%	1.022	.093
85%	7.7%	12.9%	1.035	.081
90%	15.2%	22.1%	1.040	.066
95%	28.6%	35.4%	1.092	.041
97%	38.3%	41.4%	1.161	.029
99%	44.5%	46.3%	1.163	.025

Table 5.3: As the total utilization of available storage approaches capacity, the number of devices that would otherwise overfill and that require adjustment increases. CRUSH computation increases slightly while decreasing variance.

own performance tests of CRUSH in the context of Ceph [100], I found that randomizing object placement resulted in an approximately 5% penalty in write performance due to variance in the OSD workloads, related in turn to the level of variation in OSD utilizations. In practice, however, such variance is primarily only relevant for homogeneous workloads (usually writes) where a careful striping strategy is effective. More often, workloads are mixed and already appear random when they reach the disk (or at least uncorrelated to on-disk layout), resulting in a similar variance in device workloads and performance (despite careful layout), and similarly reduced aggregate throughput. I find that CRUSH’s lack of metadata and robust distribution in the face of any potential workload far outweigh the small performance penalty under a small set of workloads.

This analysis assumes that device capabilities are more or less static over time. Experience with real systems suggests, however, that performance in distributed storage systems is often dragged down by a small number of slow, overloaded, fragmented, or otherwise poorly

performing devices. Traditional, explicit allocation schemes can manually avoid such problem devices, while hash-like distribution functions typically cannot. CRUSH allows degenerate devices to be treated as a “partial failure” using the existing overload correction mechanism, diverting an appropriate amount of data and workload to avoiding such performance bottlenecks and correct workload imbalance over time.

Fine-grained load balancing by the storage system can further mitigate device workload variance by distributing the read workload over data replicas, as demonstrated by the D-SPTF algorithm [62]; such approaches, although complementary, fall outside the scope of the CRUSH mapping function and this paper.

5.3.2 Reorganization and Data Movement

I evaluate the data movement caused by the addition or removal of storage when using both CRUSH and RUSH on a cluster of 7290 devices. The CRUSH clusters are four levels deep: nine rows of nine cabinets of nine shelves of ten storage devices, for a total of 7290 devices. $RUSH_T$ and $RUSH_P$ are equivalent to a two-level CRUSH map consisting of a single tree or list bucket (respectively) containing 729 uniform buckets with 10 devices each. The results are compared to the theoretically optimal amount of movement $m_{optimal} = \frac{\Delta w}{W}$, where Δw is the combined weight of the storage devices added or removed and W is the total weight of the system. Doubling system capacity, for instance, would require exactly half of the existing data to move to new devices under an optimal reorganization.

Figure 5.5 shows the relative reorganization efficiency in terms of the *movement factor* $m_{actual}/m_{optimal}$, where 1 represents an optimal number of objects moved and larger values mean

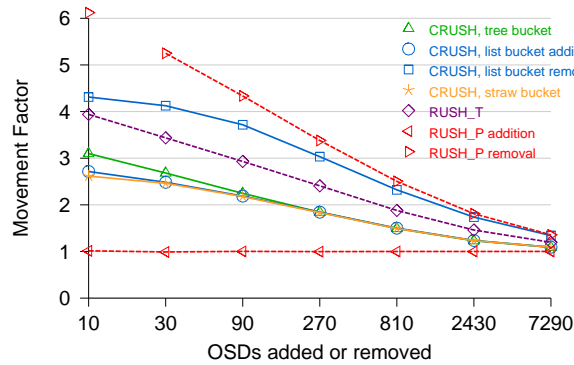


Figure 5.5: Efficiency of reorganization after adding or removing storage devices two levels deep into a four level, 7290 device CRUSH cluster hierarchy, versus $RUSH_P$ and $RUSH_T$. 1 is optimal.

additional movement. The X axis is the number of OSDs added or removed and the Y axis is the movement factor plotted on a log scale. In all cases, larger weight changes (relative to the total system) result in a more efficient reorganization. $RUSH_P$ (a single, large list bucket) dominated the extremes, with the least movement (optimal) for additions and most movement for removals (at a heavy performance penalty, see Section 5.3.3 below). A CRUSH multi-level hierarchy of list (for additions only) or straw buckets had the next least movement. CRUSH with tree buckets was slightly less efficient, but did almost 25% better than plain $RUSH_T$ (due to the slightly imbalanced 9-item binary trees in each tree bucket). Removals from a CRUSH hierarchy built with list buckets did poorly, as expected (see Section 5.2.3).

Figure 5.6 shows the reorganization efficiency of different bucket types (in isolation) when nested items are added or removed. The movement factor in a modified tree bucket is bounded by $\log n$, the depth of its binary tree. Adding items to straw and list buckets is approximately optimal. Uniform bucket modifications result in a total reshuffle of data. Modifications




Figure 5.6: Efficiency of reorganization after adding items to different bucket types. 1 is optimal. Straw and list buckets are normally optimal, although removing items from the tail of a list bucket induces worst case behavior. Tree bucket changes are bounded by the logarithm of the bucket size.

to the tail of a list (*e. g.* removal of the oldest storage) similarly induce data movement proportional to the bucket size. Despite certain limitations, list buckets may be appropriate in places within an overall storage hierarchy where removals are rare and at a scale where the performance impact will be minimal. A hybrid approach combining uniform, list, tree, and straw buckets can minimize data movement under the most common reorganization scenarios while still maintaining good mapping performance.

5.3.3 Algorithm Performance

Calculating a CRUSH mapping is designed to be fast— $O(\log n)$ for a cluster with n OSDs—so that devices can quickly locate any object or reevaluate the proper storage targets for the objects that they already store after a cluster map change. I examine CRUSH’s performance relative to $RUSH_P$ and $RUSH_T$ over a million mappings into clusters of different sizes. Fig-




Figure 5.7: CRUSH and $RUSH_T$ computation times scale logarithmically relative to hierarchy size, while $RUSH_P$ scales linearly.

Figure 5.7 shows the average time (in microseconds) to map a set of replicas into a CRUSH cluster composed entirely of 8-item tree and uniform buckets (the depth of the hierarchy is varied) versus $RUSH$'s fixed two-level hierarchy. The X axis is the number of devices in the system, and is plotted on a log scale such that it corresponds to the depth of the storage hierarchy. CRUSH performance is logarithmic with respect to the number of devices. $RUSH_T$ edges out CRUSH with tree buckets due to slightly simpler code complexity, followed closely by list and straw buckets. $RUSH_P$ scales linearly in this test (taking more than 25 times longer than CRUSH for 32768 devices), although in practical situations where the size of newly deployed disks increases exponentially over time one can expect slightly improved sub-linear scaling [41]. These tests were conducted with a 2.8 GHz Pentium 4, with overall mapping times in the tens of microseconds.

The efficiency of CRUSH depends upon the depth of the storage hierarchy and on the types of buckets from which it is built. Figure 5.8 compares the time (Y) required for $c(r, x)$ to select a single replica from each bucket type as a function of the size of the bucket (X). At a




Figure 5.8: Low-level speed of mapping replicas into individual CRUSH buckets versus bucket size. Uniform buckets take constant time, tree buckets take logarithmic time, and list and straw buckets take linear time.

high level, CRUSH scales as $O(\log n)$ —linearly with the hierarchy depth—provided individual buckets that may be $O(n)$ (list and straw buckets scale linearly) do not exceed a fixed maximum size. When and where individual bucket types should be used depends on the expected number of additions, removals, or re-weightings. List buckets offer a slight performance advantage over straw buckets, although when removals are possible one can expect excessive data shuffling. Tree buckets are a good choice for very large or commonly modified buckets, with decent computation and reorganization costs.

Central to CRUSH’s performance—both the execution time and the quality of the results—is the integer hash function used. Pseudo-random values are calculated using a multiple input integer hash function based on Jenkin’s 32-bit hash *mix* [45]. In its present form, approximately 45% of the time spent in the CRUSH mapping function is spent hashing values, making the hash key to both overall speed and distribution quality and a ripe target for optimization.

5.3.3.1 Negligent Aging

CRUSH leaves failed devices in place in the storage hierarchy both because failure is typically a temporary condition (failed disks are usually replaced) and because it avoids inefficient data reorganization. If a storage system ages in neglect, the number of devices that are failed but not replaced may become significant. Although CRUSH will redistribute data to non-failed devices, it does so at a small performance penalty due to a higher probability of backtracking in the placement algorithm. I evaluated the mapping speed for a 1,000 device cluster while varying the percentage of devices marked as failed. For the relatively extreme failure scenario in which half of all devices are dead, the mapping calculation time increases by 71%. (Such a situation would likely be overshadowed by heavily degraded I/O performance as each devices' workload doubles.) Variance in device utilization increases under such circumstances, with a 17% higher standard deviation at 50% of devices failed and a 50% higher standard deviation at 80% failed.

5.3.4 Reliability

Data safety is of critical importance in large storage systems, where the large number of devices makes hardware failure the rule rather than the exception. Randomized distribution strategies like CRUSH that decluster replication are of particular interest because they expand the number of peers with which any given device shares data. This has two competing and (generally speaking) opposing effects. First, recovery after a failure can proceed in parallel because smaller bits of replicated data are spread across a larger set of peers, reducing recovery times and shrinking the window of vulnerability to additional failures. Second, a larger peer

group means an increased probability of a coincident second failure losing shared data. With 2-way mirroring these two factors cancel each other out, while overall data safety with more than two replicas increases with declustering [111].

However, a critical issue with multiple failures is that, in general, one cannot expect them to be independent—in many cases a single event like a power failure or a physical disturbance will affect multiple devices, and the larger peer groups associated with declustered replication greatly increase the risk of data loss. CRUSH’s separation of replicas across user-defined failure domains (which does not exist with RUSH or existing hash-based schemes) is specifically designed to prevent concurrent, correlated failures from causing data loss. Although it is clear that the risk is reduced, it is difficult to quantify the magnitude of the improvement in overall system reliability in the absence of a specific storage cluster configuration and associated historical failure data to study. Although I hope to perform such a study in the future, it is beyond the scope of this thesis.

5.4 Future Work

Because the CRUSH placement algorithm is designed to preserve the basic properties of the mapping regardless of individual bucket properties, existing placement algorithms for non-replicated data can easily be supported as new bucket types. In particular, we plan to implement a consistent hashing bucket type, as it offers attractive run-time for large buckets (usually $O(1)$) with reasonable stability properties.

The primitive rule structure currently used by CRUSH is just complex enough to

support the data distribution policies I currently envision for Ceph. Other systems may have specific needs that can be met with a more powerful rule structure.

Although data safety concerns related to coincident failures were the primary motivation for designing CRUSH, study of real system failures is needed to determine their character and frequency before Markov or other quantitative models can be used to evaluate their precise effect on a system's mean time to data loss (MTTDL).

CRUSH's performance is highly dependent on a suitably strong multi-input integer hash function. Because it simultaneously affects both algorithmic correctness—the quality of the resulting distribution—and speed, investigation into faster hashing techniques that are sufficiently strong for CRUSH is warranted.

Randomized data distributions statistically correlate device utilization with workload, reducing device performance and capacity characteristics to a one-dimensional weight metric. I have conducted some preliminary investigation into overlaying multiple CRUSH mappings onto the same set of devices to facilitate distribution of data in different “tiers,” each with different bandwidth versus storage requirements. Further investigation of the approach is necessary to determine its feasibility.

5.5 Conclusions

Distributed storage systems present a distinct set of scalability challenges for data placement. CRUSH meets these challenges by casting data placement as a pseudo-random mapping function, eliminating the conventional need for allocation metadata and instead dis-

tributing data based on a weighted hierarchy describing available storage. The structure of the cluster map hierarchy can reflect the underlying physical organization and infrastructure of an installation, such as the composition of storage devices into shelves, cabinets, and rows in a data center, enabling custom placement rules that define a broad class of policies to separate object replicas into different user-defined failure domains (with, say, independent power and network infrastructure). In doing so, CRUSH can mitigate the vulnerability to correlated device failures typical of existing pseudo-random systems with declustered replication. CRUSH also addresses the risk of device overfilling inherent in stochastic approaches by selectively diverting data from overfilled devices, with minimal computational cost.

CRUSH accomplishes all of this in an exceedingly efficient fashion, both in terms of the computational efficiency and the required metadata. Mapping calculations have $O(\log n)$ running time, requiring only tens of microseconds to execute with thousands of devices. This robust combination of efficiency, reliability and flexibility makes CRUSH an appealing choice for large-scale distributed storage systems.

Chapter 6

Distributed Object Storage

At the petabyte scale, storage systems are necessarily dynamic: they are built incrementally, they grow and contract with the deployment of new storage and decommissioning of old devices, devices fail and recover on a continuous basis, and large amounts of data are created and destroyed. Effectively maintaining proper levels of replication and a balanced distribution of data at scale challenges conventional approaches to storage management that rely on centralized controllers and allocation tables.

RADOS is a Reliable, Autonomic Distributed Object Store that provides excellent performance and reliability while scaling to many thousands of storage devices. RADOS facilitates an evolving, balanced distribution of data and workload across a dynamic and heterogeneous storage cluster while providing applications with the illusion of a single logical object store with well-defined safety semantics and strong consistency guarantees. Metadata bottlenecks associated with data layout and storage allocation are eliminated through the use of a compact *cluster map* that describes cluster state and data layout in terms of *placement groups*.

In Ceph, we use a data distribution *function* (see Chapter 5) to specify this layout, although other methods (including an explicit map) are possible. RADOS facilitates scalability, efficient and consistent data access, and seamless adaptation to cluster changes through a protocol that efficiently and safely distributes map changes to OSDs and clients utilizing the storage cluster.

Replication, failure detection, and failure recovery operations are managed by intelligent OSDs, allowing the system to function with minimal oversight while scaling to many petabytes. A special-purpose object file system called EBOFS (described in Chapter 7) provides the necessary interface and safety semantics for low-level object storage with excellent performance.

Although a broad range of existing systems implement various forms of data replication, it is the careful use of the cluster map that allows RADOS to provide its unique combination of scalability, performance, and consistent data access in a cluster environment. Most significantly, clients holding a copy of the map can access data without consulting a centralized object directory or metadata server, and failure detection and recovery are performed in parallel by OSDs with minimal oversight.

6.1 Overview

A RADOS cluster consists of a large collection of OSDs and a small group of monitors responsible for managing OSD cluster membership (Figure 6.1). Each OSD includes a CPU, some volatile RAM, a network interface, and a locally attached disk drive or RAID. Monitors require only a small amount of disk space.




Figure 6.1: A cluster of many thousands of OSDs store all objects in the system. A small, tightly coupled cluster of monitors collectively manages the cluster map that describes the current cluster composition and the distribution of data. Each client instance exposes a simple storage interface to applications.

A *client* running on each host computer exposes an asynchronous I/O interface to applications utilizing the system, while hiding all of the complexity related to the dynamic and distributed nature of the storage cluster. This is exemplified by the read interface, which looks like

read(*oid*, *offset*, *length*, *buffer*, *onfinish*),

where *oid* is an object identifier (currently 128 bits), *offset* and *length* specify a byte range within the object, and *onfinish* identifies a callback to notify the application of completion. The client conceals the details of the physical location of the object, messages exchanged, and any failure scenarios: as long as the cluster is sufficiently available, the operation will eventually complete.

A critical challenge for large-scale distributed systems is reliability: as the number of storage devices scales to petabytes and beyond, the likelihood of device failure and data loss using conventional reliability mechanisms increases to unacceptable levels [112]. Our data

safety model is guided by the observation that, beyond the basic limitations imposed by the finite amounts of RAM in most systems, there are two primary reasons why applications store data. First, to make the data available to other parties; ordinarily this should to happen as quickly as possible to facilitate efficient data sharing and synchronization between cooperating clients. Second, for safety or durability: to know that the data will survive device, power, or other infrastructure failures. However, strong safety typically comes at a high cost: synchronous disk writes and mechanisms like file system journals incur additional latencies that degrade system performance.

For this reason, RADOS disassociates *serialization* from *safety* at all levels of the architecture, while providing *strong consistency* semantics to applications. That is, read and write operations logically occur in some sequential order, and completed write operations are reflected by subsequent read operations. Considering serialization (ordering) and safety independently allows RADOS to provide excellent performance without compromising consistency and safety. As a result, the client write interface differs somewhat from convention:

write(*oid*, *offset*, *length*, *buffer*, *onack*, *oncommit*),

where the *onack* callback indicates that the update is visible to subsequent reads (*e. g.* by other clients), and *oncommit* indicates that it is safely committed to disk. More specifically, although updates are atomic in all circumstances, *ack* is a promise that strong consistency semantics (ordering, in particular) will be preserved, provided the client does not crash before a *commit* is received (clients may need to be able to replay any updates for which they did not receive a commit).

Although commit notification is provided to the application via a callback, it can

safely be (and is usually) ignored. Similarly, although the client process must survive long enough to see a commit to ensure strong consistency is maintained, a client failure typically means a failure of the application as well, in which case update ordering semantics are usually moot. Of course, clients requiring strong safety can simply ignore the *ack* and receive performance comparable to systems based on synchronous writes.

RADOS provides excellent performance, reliability, and scalability via three key design features.

- **The distribution of data and cluster state are managed by manipulating a compact *cluster map*.** The map includes a compact hierarchical description of the devices participating in the cluster that is used by CRUSH, a globally known mapping function that maintains a balanced pseudo-random distribution of objects while taking special care to maintain data safety. This provides all parties—clients and storage devices alike—with complete knowledge of the distribution of data: object locations are calculated when needed, without any need to consult a centralized object directory. Instead, a small, tightly coupled cluster of *monitors* are collectively responsible for managing the map, and through it, the cluster.
- **Synchronization is disassociated from safety** in the update protocol. This facilitates efficient concurrent access to the same objects, while still providing strong safety semantics when applications require it.
- **RADOS leverages device intelligence to distribute data replication, failure detection, failure recovery, and data migration.** OSDs accomplish this distributed management

by observing cluster map differences among peers and following a peering algorithm to maintain consistency and the proper distribution and replication of data. Updates are applied using a combination of primary-copy replication, chain replication [96], and a hybrid scheme in order to minimize update latency while providing strong consistency and data safety semantics. Clients are spared most of the complexity surrounding data replication in a dynamic environment with an evolving data distribution.

I begin by describing the distribution of data and management of the cluster map in Section 6.2. Section 6.3 describes the replication of objects, the reliable processing of object reads and updates, and distributed failure recovery. I also discuss client participation in recovery, as well as the ObjectCacher module, which facilitates both efficient client operation and locking for atomic multi-object operations. I evaluate performance and scalability in Section 6.4.

6.2 Distributed Object Storage

RADOS achieves excellent scalability by eliminating the controllers and gateway servers present in most storage architectures. Instead, clients are given direct access to storage devices. This is enabled by CRUSH, which provides clients and OSDs with complete knowledge of the current data distribution. When device failures or cluster expansion require a change in the distribution of data, OSDs communicate amongst themselves to realize that distribution, without any need for controller oversight.

The RADOS cluster is managed exclusively through manipulation of the *cluster map*, a small data structure that describes what OSDs are participating in the storage cluster and




Figure 6.2: Objects are grouped into *placement groups* (PGs), and distributed to OSDs via CRUSH, a specialized replica placement function. Failed OSDs (e. g. *osd1*) are filtered out of the final mapping.

how data is mapped onto those devices. A small cluster of highly-reliable *monitors* are jointly responsible for maintaining the map and seeing that OSDs learn about cluster changes. Because the cluster map is small, well-known, and completely specifies the data distribution, clients are able to treat the entire storage cluster (potentially tens of thousands of nodes) as a single logical object store.

6.2.1 Data Placement

RADOS employs a data distribution policy in which objects are pseudo-randomly assigned to devices. When new storage is added, a random subsample of existing data is migrated to new devices to restore balance. This strategy is robust in that it maintains a probabilistically balanced distribution that, on average, keeps all devices similarly loaded, allowing the system to perform well under any potential workload [84]. Most importantly, data placement takes the form of a pseudo-random function that *calculates* the proper location of objects; no large or cumbersome centralized allocation table is needed.

Each object stored by the system is first mapped into a *placement group* (PG), a logical collection of objects that are replicated by the same set of devices, as with FaRM [112]. Each object’s PG is determined by a hash of the object name o , the desired level of replication r , and a bit mask m that controls the total number of placement groups in the system. That is, $pgid = (r, \text{hash}(o) \& m)$, where $\&$ is a bit-wise AND and the mask $m = 2^k - 1$, constraining the number of PGs by a power of two.

As the cluster scales, it is periodically necessary to adjust the total number of placement groups. During such adjustments, PG contents can be split in two by adding a bit to m . However, to avoid a massive split from simultaneously affecting all placement groups in the system—resulting in a massive reshuffling of half of all data—in practice we replace the m mask with the $\text{stablemod}(x, n, m)$ function, where $n \& m = n$ and $n \& \bar{m} = 0$ (where the bar indicates a bit-wise NOT). That is, $pgid = (r, \text{stablemod}(\text{hash}(o), n, m))$. This similarly constrains the range of $pgid$ while allowing n to be *any* number of PGs—not just a power of two. If $x \& m$ is less than n , we proceed as before. Otherwise, $\text{stablemod}(x, n, m)$ returns $x \& (m \gg 1)$ (see Algorithm 3). This provides a “smooth” transition between bit masks m , such that PG splits can be spread over time.

Algorithm 3 Function to constrain the number of PGs. Note that $m = 2^k - 1$, $n \& m = n$, and $n \& \bar{m} = 0$.

```

1: procedure STABLEMOD( $x, n, m$ )                                ▷ Choose between mask  $m$  and  $m \gg 1$ 
2:   if  $x \& m < n$  then
3:     return  $x \& m$                                               ▷ Use larger mask.
4:   else
5:     return  $x \& (m \gg 1)$                                        ▷ Use smaller mask.
6:   end if
7: end procedure

```

Placement groups are assigned to OSDs using CRUSH (see Chapter 5), a pseudo-

random data distribution function that efficiently maps each PG to an ordered list of r OSDs upon which to store object replicas. From a high level, CRUSH behaves similarly to a hash function: placement groups are deterministically but pseudo-randomly distributed. Unlike a hash function, however, CRUSH is stable: when one (or many) devices join or leave the cluster, most PGs remain where they are; CRUSH shifts just enough data to maintain a balanced distribution. In contrast, hashing approaches typically force a reshuffle of all prior mappings. CRUSH also uses weights to control the relative amount of data assigned to each device based on its capacity or performance.

Placement groups provide a means of controlling the level of replication declustering. That is, instead of an OSD sharing all of its replicas with one or more devices (mirroring), or sharing each object with different device(s) (complete declustering), the number of replication peers is related to the number of PGs μ it stores—typically on the order of 100 in the current system. As a cluster grows, the PG mask m can be periodically adjusted to “split” each PG in two. Because distribution is stochastic, μ also affects the variance in device utilizations: more PGs result in a more balanced distribution. More importantly, declustering facilitates distributed, parallel failure recovery by allowing each PG to be independently re-replicated from and to different OSDs. At the same time, the system can limit its exposure to coincident device failures by restricting the number of OSDs with which each device shares common data.

6.2.2 Cluster Maps

The cluster map provides a globally known specification of which OSDs are responsible for storing which data, and (more significantly) which devices are allowed to process object

epoch:	map revision
m:	number of placement groups -1
up:	$\text{OSD} \mapsto \{ \text{network address, } \textit{down} \}$
in:	$\text{OSD} \mapsto \{ \textit{in}, \textit{out} \}$
crush:	CRUSH hierarchy and placement rules

Table 6.1: Data elements present in the OSD cluster map, which describes both cluster state and the distribution of data.

reads or updates. Each time the cluster map changes due to an OSD status change, the map *epoch* is incremented. Map epochs allow all parties to agree on what the current distribution of data is, and to determine when their information is (relatively) out of data. Because cluster map changes may be frequent, as in a very large system where OSDs failures and recoveries are the norm, updates are distributed as *incremental maps*: small messages describing the differences between two successive epochs. In most cases, such updates simply state that one or more OSDs have failed or recovered, although in general they may include status changes for many devices, and multiple updates may be bundled together to describe the difference between distant map epochs.

6.2.2.1 Down and Out

The cluster map's hierarchical specification of storage devices is complemented by the current network address of all OSDs that are currently online and reachable (*up*), and indication of which devices are currently *down*. RADOS considers an additional dimension of OSD liveness: *in* devices are included in the mapping and assigned placement groups, while *out* devices are not. For each PG, CRUSH produces a list of exactly r OSDs that are *in* the mapping. RADOS then filters out devices that are *down* to produce the list of active OSDs for

the PG. If the active list is currently empty, PG data is temporarily unavailable, and pending I/O is blocked.

OSDs are normally both *up* and *in* the mapping to actively service I/O, or both *down* and *out* if they have failed, producing an active list of exactly r OSDs. OSDs may also be *down* but still *in* the mapping, meaning that they are currently unreachable but PG data has not yet been remapped to another OSD (similar to the “degraded mode” in RAID systems). Likewise, they may be *up* and *out*, meaning they are online but idle. This facilitates a variety of scenarios, including tolerance of intermittent periods of unavailability (*e. g.* an OSD reboot or network hiccup) without initiating any data migration, the ability to bring newly deployed storage online without using it immediately (*e. g.* to allow the network to be tested), and the ability to safely migrate data off old devices before they are decommissioned.

6.2.3 Communication and Failure Model

RADOS employs an asynchronous, ordered point to point message passing library for communication. For simplicity, the prototype considers a failure on the TCP socket to imply a device failure, and immediately reports it. OSDs exchange periodic heartbeat messages with their peers to ensure that failures are detected. This is somewhat conservative in that an extended ethernet disconnect or a disruption in routing at the IP level will cause an immediate connection drop and failure report. However, it is safe in that any failure of the process, host, or host’s network interface will eventually cause a dropped connection. This strategy can be made somewhat more robust by introducing one or more reconnection attempts to better tolerate network intermittency before reporting a failure. OSDs that discover that they have been marked

down simply sync to disk and kill themselves to ensure consistent behavior.

6.2.4 Monitors

All OSD failures are reported to a small cluster of *monitors*, which are jointly responsible for maintaining the master copy of the cluster map. OSDs can request the latest cluster map from or report failures to any monitor. When an OSD submits a failure report, it expects to receive an acknowledgement in the form of a map update that marks the failed OSD *down* (or back *up* at a new address). If it does not get a response after a few seconds, it simply tries contacting a different monitor.

In order to ensure that responses from all monitors are consistent, the monitor cluster is based on the Paxos part-time parliament algorithm to preserve strong consistency between replicas [54]. Unlike primary-copy or similar replication schemes, the Paxos algorithm emphasizes the durability of updates over availability by requiring that a majority of monitors be available before updates are possible. The monitor cluster simplifies pure Paxos somewhat by allowing only a single update to be proposed at a time (much like Boxwood [63]), simplifying the implementation, while also coordinating updates with a *lease* mechanism to provide a consistent ordering of cluster map read and update operations.

The cluster initially elects a *leader* to serialize map updates and manage consistency. Once elected, the leader begins by requesting the map epochs stored by each monitor. Monitors have a fixed amount of time T (currently two seconds) to respond to the probe and join the *quorum*. The leader ensures that a majority of the monitors are active and that it has the most recent map epoch (requesting incremental updates from other monitors as necessary), and then

begins distributing short-term leases to active monitors.

Each lease grants active monitors permission to distribute copies of the cluster map to OSDs or clients who request it. If the lease term T expires without being renewed, it is assumed the leader has died and a new election is called. Each lease is acknowledged to the leader upon receipt. If the leader does not receive timely acknowledgements when a new lease is distributed, it assumes an active monitor has died and a new election is called. When a monitor first starts up, or finds that a previously called election does not complete after a reasonable interval, an election is called.

When an active monitor receives an update request (*e. g.* a failure report), it first checks to see if it is a new. If, for example, the OSD in question was already marked *down*, the monitor simply responds with the necessary incremental map updates to bring the reporting OSD up to date. New failures are forwarded to the leader, who serializes updates, increments the map epoch, and uses the Paxos update protocol to distribute the update to other monitors, simultaneously revoking leases. Once the update is acknowledged by a majority of monitors a final commit message issues a new lease.¹

The combination of a synchronous two-phase commit and the probe interval T ensures that if the active set of monitors changes, it is guaranteed that all prior leases (which have a matching term T) will have expired before any subsequent map updates take place. Consequently, any sequence of map queries and updates will result in a consistent progression of map versions—significantly, map versions will never “go backwards”—regardless of which monitor messages are sent to and despite any intervening monitor failures, provided a majority of

¹The integration of Paxos with a leasing mechanism is implemented as a generic service and is used to manage other critical state in Ceph, including the MDS cluster map and state for coordination client access to the file system.

monitors is available.

6.2.5 Map Propagation

Because the RADOS cluster may include many thousands of devices or more, it is not practical to simply broadcast map updates to all parties without unduly burdening the central monitors. Fortunately, differences in map epochs are significant only when they vary between two communicating OSDs (or between a client and OSD), which must agree on their proper roles with respect to a particular PG. This property allows RADOS to distribute map updates lazily by combining them with existing inter-OSD messages, shifting the distribution burden to OSDs. Each OSD maintains a history of past map incrementals, tags all messages with its latest epoch, and makes note of its peers' epochs. If an OSD receives a message from a peer with an older map, it shares the necessary incremental(s) to bring that peer in sync. Similarly, when contacting a peer thought to have an older epoch, incremental updates are preemptively shared. The heartbeat messages periodically exchanged for failure detection ensure that updates spread quickly—in $O(\log n)$ time for a cluster of n OSDs.

For example, when an OSD first boots, it begins by informing a monitor that it has come online, and sends its current map epoch. The monitor cluster changes the OSD's status to *up*, and replies with the incremental updates necessary to bring the OSD fully up to date. When the new OSD begins contacting OSDs with whom it shares data (see Section 6.3.5.1), the exact set of devices who are affected by its status change learn about the appropriate map updates. Because a booting OSD does not yet know which epochs its peers have, it shares a safe recent history (*e. g.* at least 30 seconds) of incremental updates.

This preemptive map sharing strategy is conservative: an OSD will always share an update when contacting a peer unless it is certain the peer has already seen it, resulting in OSDs receiving duplicates of the same update. However, the number of duplicates an OSD receives is bounded by the number of peers it has, which is in turn determined by the number of PGs μ it manages. In practice, I find that the actual level of update duplication is much lower than this (see Section 6.4.1.3).

6.3 Reliable Autonomic Storage

RADOS replicates each data object on two or more devices for reliability and availability. Replication and failure recovery are managed entirely by OSDs through a version-based consistency scheme utilizing short-term update logs. A peer to peer recovery protocol avoids any need for controller-driven recovery, facilitating a flat cluster architecture with excellent scalability.

6.3.1 Replication

Storage devices are responsible for update serialization and write replication, shifting the network overhead associated with replication from the client network or WAN to the OSD cluster's internal network, where greater bandwidth and lower latencies are expected. RADOS implements primary-copy replication [3], chain replication [96], and a hybrid I call *splay* replication that combines elements of the two. All three strategies provide strong consistency guarantees, such that read and write operations occur in some sequential order, and completed




Figure 6.3: Replication strategies implemented by RADOS. Primary-copy processes both reads and writes on the first OSD and updates replicas in parallel, while chain forwards writes sequentially and processes reads at the tail. Splay replication combines parallel updates with reads at the tail to minimize update latency.

writes are reflected by subsequent reads.

With primary-copy replication, the first OSD in a PG’s list of active devices is the *primary*, while additional OSDs are called *replicas*. Clients submit both reads and writes to the primary, which serializes updates within each PG. The write is forwarded to the replicas, which apply the update to their local object store and reply to the primary. Once all replicas are updated, the primary applies the update and replies to the client, as shown in Figure 6.3.

Chain replication separates update serialization from read processing. Writes are directed at the first OSD (the *head*), which applies the update locally and forwards it to the next OSD in the list. The last OSD (the *tail*) responds to the client. Reads are directed at the tail, whose responses will always reflect fully replicated (thus, safely applied) updates. For 2× replication, this offers a clear advantage: only three messages and network hops are necessary, versus four for primary-copy replication. However, latency is dependent on the length of the




Figure 6.4: RADOS responds with an *ack* after the write has been applied to the buffer caches on all OSDs replicating the object (shown here with splay replication). Only after it has been safely committed to disk is a second *commit* notification sent to the client.

chain, making the strategy problematic for high levels of replication.

Splay replication combines elements of the two. As with chain replication, updates are directed at the head and reads at the tail. For high levels of replication, however, updates to the middle OSDs occur in parallel, lowering the latency seen by the client. Both primary-copy and splay replication delay the local write in order to maintain strong consistency in the presence of an OSD failure, although splay must do so for less time, lowering OSD memory requirements.

6.3.2 Serialization versus Safety

RADOS disassociates write acknowledgement from safety at all levels in order to provide both efficient update serialization and strong data safety. This approach is illustrated in Figure 6.4, which corresponds to the splay scheme shown in Figure 6.3 with an additional set of messages. During a replicated write, each replica OSD sends an *ack* to the tail immediately after applying the update to the in-memory cache of the local EBOFS object store, and the tail responds to the client with an *ack* only after all replicas have applied the update (as before).

Later, when EBOFS provides each OSD with asynchronous notification that the update is safely committed to disk, they send a second message to the tail, and only after all replicas have done so is the client sent a final *commit*. The strategy is similar for the other schemes: with primary-copy replication, *acks* and *commits* go to the primary instead of the tail, and with chain replication, only *commits* go to the tail (the replicated update itself is an implicit *ack*).

Once clients receive an *ack*, they can be sure their writes are visible to others, and synchronous application level calls can typically unblock. Clients buffer all updates until a *commit* is received, however, allowing clients to participate in recovery if all OSDs replicating the update fail, losing their in-memory (uncommitted) state.

6.3.3 Maps and Consistency

Tagging all RADOS messages—both those originating from clients and from other OSDs—with the map epoch ensures that all update operations are applied in a fully consistent fashion. Because all replicas are involved in any given update operation, any relevant map updates (*i. e.* any update that changes PG membership) will be discovered. Even if the master copy of the cluster map has been updated to change a particular PGs membership, updates may still be processed by the old members, provided they have not yet heard of the change. Because a given set of OSDs who are newly responsible for a PG cannot become active (*i. e.* recover or service I/O) without consulting prior members or determining they are failed (see Section 6.3.5), no updates can be lost, and consistency is maintained.

Achieving similar consistency for read operations is slightly less natural than for updates. In the event of a partial network failure that results in an OSD becoming only partially




Figure 6.5: Each PG has a log summarizing recent object updates and deletions. The most recently applied operation is indicated by *last_update*. All updates above *last_complete* are known to have been applied, while any missing objects in the interval between *last_complete* and *last_update* are summarized in the missing list.

unreachable, the OSD servicing reads for a PG could be declared “failed” but still be reachable by clients with an old map. Meanwhile, the updated map may specify a new OSD in its place. In order to prevent any read operations from being processed by the old OSD after new updates are be processed by the new one, we require timely heartbeat messages between OSDs in each PG in order for the PG to remain available (readable). That is, if the OSD servicing reads hasn’t heard from other replicas in H seconds, reads will block. Then, in order for a new OSD to take over that role from another OSD, it must either obtain positive acknowledgement from the old OSD (ensuring they are aware of their role change), or delay for the same time interval. In my implementation, I choose a relatively short heartbeat interval of two seconds. This ensures both timely failure detection and a short interval of PG data unavailability in the event of a primary OSD failure.

6.3.4 Versions and Logs

RADOS uses versioning to identify individual updates and to serialize them within each placement group. Each version consists of an $(epoch, v)$ pair, where *epoch* reflects the map epoch at the time of the update, and *v* increases monotonically. Each PG has a *last_update* attribute that reflects the most recently applied modification to one its objects, and each object has a similar version attribute to reflect the last time it was modified.

OSDs maintain a short-term log of recent updates (illustrated in Figure 6.5) for each PG, stored both on disk and in RAM or NVRAM (if it is available). Each log entry includes the object name, the type of operation (update, delete), a version number identifying the update, and a unique request identifier consisting of the client name and a client-assigned identifier (not shown). Unique identifiers allow OSDs to detect and ignore duplicate requests, rendering all operations idempotent.

The first OSD in the PG serializes writes by assigning a new version number and appending a new entry to its log. Because by definition only one OSD fills this role during a single map epoch, versions are unique within each PG. The request is then forwarded along with the version stamp to all other replica OSDs (or just to the next replica for chain replication). An OSD processing an update always writes to the log immediately, even if it delays the write for consistency. For this reason, the log may extend below the *last_update* pointer (*i. e.* write-ahead).

Log appends or pointer changes are written to disk wrapped in atomic EBOFS transactions with the updates they describe, such that the log provides a perfect record of which

updates were (and were not) committed before any crash (see Section 7.1). The log also forms the basis for recovery when a PG is being brought up to date or replicated to an entirely new OSD. All updates below the *last_complete* are known to be applied locally, while a *missing* list summarizes the latest versions of modified objects above it. OSDs periodically trim their on-disk logs when requests have been fully flushed to disk on all replicas and clients have been notified.

6.3.5 Failure Recovery

RADOS failure recovery is driven entirely by cluster map updates and subsequent changes in each PG's list of active devices. Such changes may be due to device failures, recoveries, cluster expansion or contraction, or even complete data reshuffling from a totally new CRUSH replica distribution policy—RADOS makes very few assumptions about what kind of map changes are possible. For example, a PG might go from three OSDs to none (a complete power failure) to one OSD (after partial power is restored) and then to two entirely different OSDs (a reorganization), and these changes may happen very quickly, without a sufficient interval between each transition to allow recovery to complete. Moreover, when an OSD crashes and recovers, EBOFS object store will be warped back in time to the most recent snapshot committed to disk.

In all cases, RADOS employs a robust *peering* algorithm to establish a consistent view of PG contents and to restore the proper distribution and replication of data. This strategy relies on the basic design premise that OSDs aggressively replicate the PG log and its record of what the current state of a PG *should* be (*i. e.* what object versions it contains), even when

some object replicas may be missing locally. Thus, even if recovery is slow and object safety is degraded for some time, PG metadata is carefully guarded, simplifying the recovery algorithm and allowing the system to reliably detect data loss.

6.3.5.1 Peering

When an OSD receives a cluster map update, it walks through all new map increments up through the most recent to examine and possibly adjust PG state values. Any locally stored PGs whose active list of OSDs changes are marked *inactive*, indicating that they must re-peer. Considering all map epochs (not just the most recent) ensures that intermediate data distributions are taken into consideration: if an OSD is removed from a PG and then added again, it is important to realize that intervening updates to PG contents may have occurred. As with replication, peering (and any subsequent recovery) proceeds independently for every PG in the system.

The process is driven by the first OSD in the PG (the *primary*). For each PG an OSD stores for which it is not the current primary (*i. e.* it is a *replica*, or a *stray* which is longer in the active set), a *Notify* message is sent to the current primary. This message includes basic state information about the locally stored PG, including *last_update*, *last_complete*, the bounds of the PG log, and *last_epoch_started*, which indicates the most recent known epoch during which the PG successfully peered.

Notify messages ensure that an OSD that is the new primary for a PG discovers its new role without having to consider all possible PGs (of which there may be millions) for every map change. Once aware, the primary generates a *prior set*, which includes all OSDs that may have

participated in the PG since *last_epoch_started*. Because this is a lower bound, as additional notifies are received, its value may be adjusted forward in time (and the prior set reduced). The prior set is explicitly queried to solicit a notify to avoid waiting indefinitely for a prior OSD that does not actually store the PG (*e. g.* if peering never completed for an intermediate PG mapping).

Armed with PG metadata for the entire prior set, the primary can determine the most recent update applied on any replica, and request whatever log fragments are necessary from prior OSDs in order to bring the PG logs up to date on active replicas. That is, the primary must assemble a log that stretches from the oldest log bottom on active replicas to the newest log bottom (most recent update) on any prior OSD. Because the log only reflects recent history, this may not be possible (*e. g.* if the primary is new to the PG and does not have any PG contents at all), making it necessary for the primary to generate or request a *backlog*. A backlog is an extended form of the log that includes entries above the *top* pointer (where the log was last trimmed) to reflect any other objects that exist in the PG (*i. e.* on disk) but have not been modified recently. The backlog is generated by simply scanning locally stored PG contents and creating entries for objects with versions prior to the log top. Because it does not reflect prior deletions, the backlog is only a partial record of the PG's modification history.

Once the primary has assembled a sufficient log, it has a complete picture of the most recent PG contents: they are either summarized entirely by the log (if it has a backlog), or the recent log in combination with locally stored objects. From this, the primary updates its *missing list* by scanning the log for objects it does not have (those updated after its previous *last_update*). All OSDs maintain a missing list for active PGs, and include it when logs are requested by the

primary. The primary can infer where objects can be found by looking at which OSDs include the object in their log but don't list it as missing.

Once the log and missing list are complete, the PG is ready to be activated. The primary first sends a message to all OSDs in the prior set (but not in the active set) to update *last_epoch_started*. Once this is acknowledged, the primary sets its own PG to *active*, and sends a log fragment to each OSD in the active set to bring them up to date and mark them *active* as well. Updating *last_epoch_started* on residual OSDs implicitly renders them *obsolete* in that they know the PG became active in an epoch after their *last_update* and their information is likely out of date. In the future, a primary left with only obsolete information from its prior set can opt to either consider itself crashed or, if an administrator is desperate, bring the PG online with potentially stale data.

6.3.5.2 Recovery

A critical advantage of declustered replication is the ability to parallelize failure recovery [4, 112]. Replicas shared with any single failed device are spread across many other OSDs, and each PG will independently choose a replacement, allowing re-replication to just as many more OSDs. On average, in a large system, any OSD involved in recovery for a single failure will be either pushing or pulling content for only a single PG, making recovery very fast.

The recovery strategy in RADOS is motivated by the observation that I/O is most often limited by read (and not write) throughput. A simple recovery strategy is for each OSD to independently walk through its PG log and “pull” any objects on its missing list from other OSDs, updating *last_complete* along the way, until it reaches the bottom of the log. Although

this strategy works (and was used by previous versions of the system), it has two limitations.

First, if multiple OSDs are independently recovering objects in the same PG, their recovery will not be synchronized. That is, they will probably not pull the same objects from the same OSDs at the same time, resulting in duplication of the most expensive aspect of recovery: seeking and reading. Second, the update replication protocols (described in Section 6.3.1) become increasingly complex if replica OSDs are missing the objects being modified. Although the primary OSD can simply delay updates on missing objects until they are recovered (since it is responsible for ordering requests), replicas do not have that flexibility, which significantly complicates consistency logic.

For these reasons, recovery in RADOS is coordinated by the primary. As before, operations on missing objects are delayed until the primary has a local copy. Since the primary already knows which objects all replicas are missing from the peering process, it can preemptively “push” any missing objects that are about to be modified to replica OSDs, simplifying replication logic while also ensuring that the object is only read once. If a replica is handling reads, as in splay replication, requests for missing objects are delayed until the object can be pulled from the primary. If the primary is pushing an object (*e. g.* in response to a pull request), or if it has just pulled an object for itself, it will always push to all replicas that need a copy while it has the object in memory. Thus, in the aggregate, every re-replicated object is read only once.

6.3.5.3 Client Participation

If a RADOS client has an outstanding (*i. e.* *un-acked* or *un-committed*) request submitted for a PG that experiences a failure, it will simply resubmit the request (with the same unique request identifier) to the new PG primary. This ensures that if the request was not completely replicated or otherwise did not survive the failure, it will still be processed. If the OSD discovers the request was already applied by the request's presence in the log, it will consider the operation a *no-op*, but will otherwise process (*i. e.* *replicate*) it normally so that the client still receives an *ack* and *commit* with the same associated promises.

6.3.5.4 Concurrent Failures

If all OSDs participating in a PG simultaneously fail, the PG is said to have *crashed*, and data becomes unavailable until at least one device comes back online (*e. g.* after a temporary power failure). Recovery, however, is hindered by the probable loss of some updates that were applied and serialized but existed only in RAM. This is problematic, because some clients may have already read the applied (but uncommitted, and now lost) updates.

In order to preserve strong consistency in such situations, OSDs include the version associated with each update in the client *ack*, and the RADOS client buffers all updates it submits until a final *commit* is received. If a PG with which it has uncommitted updates crashes, the client includes the previously assigned version with the resubmitted request. When a crashed PG is recovering, OSDs enter a *replay* period for a fixed amount of time (*e. g.* 30 seconds) after peering but before becoming active. The primary OSD buffers requests such that when the replay period ends, it can reorder any requests that include versions to reestablish the original

order of updates.

From the client's perspective, this preserves consistency: the version number in the *ack* allows the client to reassert the correct ordering (if need be), as long as the client doesn't fail before receiving a *commit*. As described, however, the strategy has a critical flaw that compromises strong consistency semantics: if a read operation is allowed after the update is applied and the *ack* is sent, such that it sees the effects, but the OSD subsequently fails such that the *ack* (and update version) never reaches the client, the client will be unable to reassert the previously established (and witnessed) ordering. Although it is improbable, the asynchronous messaging model makes such a scenario possible. In fact, it would be impossible for the OSD to be sure the client received the *ack* without the client confirming it with an additional message (like a two-phase commit) and an associated increase in latency and protocol complexity.

For this reason, RADOS can delay reads to uncommitted data, while taking steps to expedite their commit to disk (see Figure 6.6). This approach maintains a low latency for writes, and only increases latency for reads if two operations are actually dependent (*i. e.* reference overlapping byte ranges in a data object, a relatively rare occurrence in most workloads [8, 77]). Alternatively, a small amount of NVRAM can be employed on the OSD for PG log storage, allowing serialization to be preserved across power failures such that resubmitted client operations can be correctly reordered, similarly preserving fully consistent semantics.

6.3.6 Client Locking and Caching

Although individual object updates are both strongly consistent and atomic, many applications require atomicity across multiple objects [63, 100]. In contrast to systems that




Figure 6.6: Reads in shared read/write workloads are usually unaffected by write operations. However, reads of uncommitted data can be delayed until the update commits. This increases read latency in certain cases, but maintains a fully consistent behavior for concurrent read and write operations in the event that all OSDs in the placement group simultaneously fail and the write *ack* is not delivered.

use a distributed lock manager [51, 86, 94, 107], RADOS locks are issued and enforced by the OSDs that store objects. Read (shared) and write (exclusive) locks are implemented as object attributes, and lock acquisition and release behave like any other object update: they are serialized and replicated across all OSDs in the PG for consistency and safety. Locks can either time out (*i. e.* be treated as leases) or applications can empower a third party to revoke on behalf of failed clients.

An *ObjectCacher* module layers on top of the RADOS client to manage client lock state and provide basic object caching services and multi-object updates. The *ObjectCacher* transparently acquires the necessary locks to achieve proper cache consistency (read locks on cached objects and write locks to allow write-back). Write locks can also be used to mask latency associated with large updates: ordering is established when the lock is acquired, and is released asynchronously as the data is written back to the OSD. Operations on multiple objects practice deadlock avoidance during lock acquisition.




Figure 6.7: Per-OSD write performance. The horizontal line indicates the upper limit imposed by the physical disk. Replication has minimal impact on OSD throughput, although if the number of OSDs is fixed, n -way replication reduces total *effective* throughput by a factor of n because replicated data must be written to n OSDs.

6.4 Performance and Scalability

I examine the performance, scalability, and failure recovery behavior of the RADOS architecture. I begin by examining the performance of individual OSDs, and evaluate write latency in light of the three replication strategies I implement. Scalability is considered independently in terms of the specific performance-limiting factors. Finally, I discuss failure recovery and its effect on system performance.

Performance tests are conducted using OSDs running on dual-processor Pentiums with SCSI disks. In my experiments, RADOS achieves perfect linear scaling up to 24 OSDs, after which throughput is limited by the network switch.

6.4.0.1 OSD Throughput

I first measure the I/O performance of a 14-node cluster of OSDs. Figure 6.7 shows per-OSD throughput (y) with varying write sizes (x) and replication. Workload is generated by




Figure 6.8: Write latency for varying write sizes and primary-copy replication. More than two replicas incurs minimal additional cost for small writes because replicated updates occur concurrently. For large synchronous writes, transmission times dominate. Clients partially mask that latency for writes over 128 KB by acquiring exclusive locks and asynchronously flushing the data.

400 clients on 20 additional nodes. Performance is ultimately limited by the raw disk bandwidth (around 58 MB/sec), shown by the horizontal line. Replication doubles or triples disk I/O, reducing client data rates accordingly when the number of OSDs is fixed.

6.4.0.2 Write Latency

Figure 6.8 shows the synchronous write latency (y) for a single writer with varying write sizes (x) and primary-copy replication. Because the primary OSD simultaneously retransmits updates to all replicas, small writes incur a minimal latency increase for more than two replicas. For larger writes, the cost of retransmission dominates; 1 MB writes (not shown) take 13 ms for one replica, and 2.5 times longer (33 ms) for three. High update latencies are mitigated for writes over 128 KB with the client ObjectCacher module, which uses exclusive locks to establish serialization for large writes, while flushing data to OSDs and releasing locks asynchronously. Under such circumstances, latencies are determined by the exchange of lock




Figure 6.9: Latency for 4 KB writes under different replication strategies and levels of replication. Splay replication augments chain replication with parallel replica updates to achieve performance similar to primary-copy replication for high levels of replication.

requests, providing latency for isolated large writes which approaches that of small writes.

Figure 6.9 shows the latency associated with different replication strategies for small 4 KB writes. Although chain replication offers the lowest latency for $2\times$ replication, where a minimum number of messages are exchanged, it performs poorly with large numbers of replicas. Primary-copy masks the latency for replication above $2\times$ by parallelizing updates. Splay replication combines the best of both approaches—optimal message exchange for $2\times$ and parallel updates with more replicas—with the added bonus of lowering memory utilization on the tail OSD.

Figure 6.10 shows latency incurred per replica (y) for various replication levels (x) with 1 MB writes. For large writes, the benefits of parallel updates are limited by per-device bandwidth limitations. Here, the crossover point shifts from around $2\times$ to $5\times$, giving chain replication a slight edge.




Figure 6.10: Normalized write latency per replica for 1 MB writes under different replication strategies and levels of replication. Parallel replication does not improve latency for replications levels below six.

6.4.1 Scalability

The scalability of RADOS is potentially limited by three primary factors: the quality of the data distribution generated by CRUSH, interaction with the monitor cluster, and the effective distribution of cluster map updates. Other elements of the system are trivially parallelizable. In particular, failure detection, failure recovery, and replication are all bounded on each OSD by the number of peers, regardless of the cluster size. For our purposes, I assert that a sufficiently large and fast network can be constructed to service many thousands of OSDs [42]. CRUSH includes some provisions for segregating replication traffic (by keeping all replicas inside a suitably large domain), but I otherwise consider the network to be outside the scope of this work.




Figure 6.11: OSD write performance scales linearly with the size of the OSD cluster until the switch is saturated at 24 OSDs. CRUSH and hash performance improves when more PGs lower variance in OSD utilization.

6.4.1.1 Data Placement

As seen in Chapter 5, CRUSH produces a distribution of data that closely matches the mean and variance of a binomial or normal distribution [101], meaning it appears random, even though it is a deterministic and constrained mapping. The primary consequence is that the variance σ^2 in the number PGs per OSD—and subsequently, in OSD storage utilizations and workloads—is related to the average number of PGs per OSD (μ), where $\sigma^2 \approx \mu$. With an average of 100 PGs per OSD, the standard deviation σ is 10%; with 1000 per OSD, σ drops to 3%. This behavior holds even for large clusters composed of heterogeneous (*i. e.* non-uniformly weighted) devices.

Figure 6.11 shows per-OSD write throughput as the cluster scales using CRUSH, a simple hash function, and a linear striping strategy to distribute data in 4096 or 32768 PGs among available OSDs. Linear striping balances load perfectly for maximum throughput to provide a benchmark for comparison, but like a simple hash function, it fails to cope with

device failures or other OSD cluster changes. Because data placement with CRUSH or a hash is stochastic, throughputs are lower with fewer PGs: greater variance in OSD utilizations causes request queue lengths to drift apart under this entangled client workload.

Although a probabilistic data distribution means that some devices may become overloaded (*i. e.* handle many more than μ PGs) with small probability, PGs can be explicitly diverted away from specific devices using the overload mechanism in CRUSH. Unlike the hash and linear strategies, CRUSH also minimizes data migration under cluster expansion while maintaining a balanced distribution.

Most importantly, the computational cost of calculating a CRUSH mapping is $O(\log n)$ for a cluster of size n , allowing mappings in the tens of microseconds for even extremely large clusters.

6.4.1.2 Monitor Interaction

The monitor cluster is designed both for extreme reliability and for high availability. In the general case, monitors do very little work—they process small messages in response to failures, but are otherwise idle. A worst case load for the monitor cluster occurs when large numbers of OSDs appear to fail in a short period. If each OSD stores μ PGs and f OSDs fail, then an upper bound on the number of failure reports generated is on the order of μf , which could be very large if a large OSD cluster experiences a network partition. To prevent such a deluge of messages, OSDs throttle and batch failure reports, imposing an upper bound on monitor load proportional to the cluster size.

Although failures may be reported to multiple random monitors multiple times, only

the first few reports of a given failure will be forwarded to the elected lead monitor. Map updates are quickly propagated among monitors such that subsequent reports of the same failures will be reflected by the current map and result in an immediate response to the OSD. For this reason, OSDs send heartbeats to peers at semi-random intervals to stagger detection of failures, dispersing reports for a given failure over time. Furthermore, map updates returned to a reporting OSDs will also reflect all other failures processed to date, preventing some future failure reports from being sent.

6.4.1.3 Map Propagation

The RADOS map distribution algorithm (Section 6.2.5) ensures that updates reach all OSDs after only $\log n$ hops. However, as the size of the storage cluster scales, the frequency of device failures and related cluster updates increases. Because map updates are only exchanged between OSDs who share PGs, the hard upper bound on the number of copies of a single update an OSD can receive is proportional to μ .

In simulations under near-worst case propagation circumstances with regular map updates, I found that update duplicates approach a steady state even with exponential cluster scaling. In this experiment, the monitors share each map update with a single random OSD, who then shares it with its peers. In Figure 6.12 I vary the cluster size x and the number of PGs on each OSD (which corresponds to the number of peers it has) and measure the number of duplicate map updates received for every new one (y). Update duplication approaches a constant level—less than 20% of μ —even as the cluster size scales exponentially, implying a fixed map distribution overhead. I consider a worst-case scenario in which the only OSD




Figure 6.12: Duplication of map updates received by individual OSDs as the size of the cluster grows. The number of placement groups on each OSD effects number of peers it has who may share map updates.

chatter consists of ping messages for failure detection, which means that, generally speaking, OSDs learn about map updates (and the changes known by their peers) as slowly as possible. Limiting map distribution overhead thus relies only on throttling the map update frequency, which the monitor cluster already does as a matter of course.

6.4.2 Failure Recovery

Figure 6.13 shows write throughput over time as a cluster of 20 (real) OSDs recovers from two (simulated) failures at time 30. 30 clients are writing data with $2\times$ replication and saturating the cluster. As the failed OSDs are initially marked down, replication throughput drops because affected PGs are temporarily unreplicated, while effective client write performance correspondingly increases. At time 50 the OSDs are marked out and recovery is initiated. Performance drops while active and then inactive objects are re-replicated to other OSDs,




Figure 6.13: Write throughput over time as a saturated 20 OSD cluster recovers from two OSD failures at time 30. Data re-replication begins at time 50 and completes at time 80.




Figure 6.14: Write throughput over time as an unsaturated 20 OSD cluster recovers from one OSD failure at time 30 and two more at time 80.

but throughput eventually returns to a level slightly below baseline (due to the two missing OSDs). The throughput penalty associated with recovery is exaggerated in this experiment for two reasons. First, in such a small cluster, all OSDs either share data with the failed devices, or are selected as replacements in one or more PGs. In a large system, each OSD failure will initiate recovery for only a single PG (of a hundred or more) on each of its peers. Second, the relatively naive implementation currently makes almost no attempt to balance recovery with regular workload. Nevertheless, recovery proceeds quickly in parallel and the cluster resumes performing at only a slightly degraded level after about 30 seconds.

This is evident in Figure 6.14, which shows first one and then two OSDs failures on an only partially loaded cluster. The first recovery minimally effects throughput both because more disk bandwidth is available and because fewer PGs contain data to be replicated. The second recovery involves more PGs and more data, with a greater impact on performance.

6.5 Future Work

Although RADOS was developed for use in the Ceph distributed file system [100], the reliable and scalable object storage service it provides is well-suited for a variety of other storage abstractions. In particular, the current interface based on reading and writing byte ranges is primarily an artifact of the intended usage for file data storage. Objects might have any query or update interface or resemble any number of fundamental data structures. Potential services include distributed B-link trees that map ordered keys to data values (as in Boxwood [63]), high-performance distributed hash tables [90], or FIFO queues (as in GFS [30]).

Although RADOS manages scalability in terms of total aggregate storage and capacity, this dissertation does not address the issue of many clients accessing a single popular object. I have implemented a *read shedding* mechanism which allows a busy OSD to shed reads to object replicas for servicing, when the replica's OSD has a lower load and when consistency allows (*i. e.* there are no conflicting in-progress updates). Heartbeat messages exchange information about current load in terms of recent average read latency, such that OSDs can determine if a read is likely to be service more quickly by a peer. This facilitates fine-grained balancing in the presence of transient load imbalance, much like D-SPTF [62]. Notably, this read shedding is only possible with primary-copy replication, as the OSD servicing reads must be aware of any in-progress writes in order to preserve consistency. Although preliminary experiments are promising, a comprehensive evaluation has not yet been conducted.

More generally, the distribution of workload in RADOS is currently dependent on the quality of the data distribution generated by object layout into PGs and the mapping of PGs to OSDs by CRUSH. Although I have considered the statistical properties of such a distribution and demonstrated the effect of load variance on performance for certain workloads, the interaction of workload, PG distribution, and replication can be complex. For example, write access to a PG will generally be limited by the slowest device storing replicas, while workloads may be highly skewed toward possibly disjoint sets of heavily read or written objects. I have conducted only minimal analysis of the effects of such workloads on efficiency in a cluster utilizing declustered replication, or the potential for techniques like read shedding to improve performance in such scenarios.

Although total device failure is addressed, OSDs do not currently consider partial

failures, like corrupted disk blocks. Checksums or preemptive mechanisms like disk scrubbing would dramatically improve data safety.

The integration of intelligent disk scheduling, including the prioritization of replication versus workload and quality of service guarantees, is an ongoing area of investigation within the research group [109].

6.6 Related Work

Most distributed storage systems utilize centralized metadata servers [1, 14, 30] or collaborative allocation tables [86] to manage data distribution, ultimately limiting system scalability. For example, like RADOS, Ursa Minor [1] provides a distributed object storage service (and, like Ceph, layers a file system service on top of that abstraction). In contrast to RADOS, however, Ursa Minor relies on an object manager to maintain a directory of object locations and storage strategies (replication, erasure coding, etc.), limiting scalability and placing a lookup in the data path. RADOS does not provide the same versatility as Ursa Minor's dynamic choice of timing and failure models, or support for online changes to object encoding (although encoding changes are planned for the future); instead, RADOS focuses on scalable performance in a relatively controlled (non-Byzantine) environment.

The Sorrento [93] file system's use of collaborative hashing [47] bears the strongest resemblance to RADOS's application of CRUSH. Many distributed hash tables (DHTs) use similar hashing schemes [21, 80, 90], but these systems do not provide the same combination of strong consistency and performance that RADOS does.

For example, DHTs like PAST [80] rely on an overlay network [81, 90, 114] in order for nodes to communicate or to locate data, limiting I/O performance. More significantly, objects in PAST are immutable, facilitating cryptographic protection and simplifying consistency and caching, but limiting the systems usefulness as a general storage service. CFS [21] utilizes the DHash DHT to provide a distributed peer-to-peer file service with cryptographic data protection and good scalability, but performance is limited by the use of the Chord [90] overlay network. In contrast to these systems, RADOS targets a high-performance cluster or data center environment; a compact cluster map describes the data distribution, avoiding the need for an overlay network for object location or message routing.

Most existing object-based storage systems rely on controllers or metadata servers to perform recovery [14, 68], or centrally micro-manage re-replication [30], failing to leverage intelligent storage devices. Other systems have adopted declustered replication strategies to distribute failure recovery, including OceanStore [52], Farsite [2], and Glacier [37]. These systems focus primarily on data safety and secrecy (using erasure codes or, in the case of Farsite, encrypted replicas) and wide-area scalability (like CFS and PAST), but not performance.

FAB (Federated Array of Bricks) [82] provides high performance by utilizing two-phase writes and voting among bricks to ensure linearizability and to respond to failure and recovery. Although this improves tolerance to intermittent failures, multiple bricks are required to ensure consistent read access, while the lack of complete knowledge of the data distribution further requires coordinator bricks to help conduct I/O. FAB can utilize both replication and erasure codes for efficient storage utilization, but relies on the use of NVRAM for good performance. In contrast, RADOS's cluster maps drive consensus and ensure consistent access

despite a simple and direct data access protocol. RADOS's cluster map bears some resemblance to layout epochs in Palladio [35]: epoch numbers version views of the cluster membership and state, per-store layouts resemble RADOS's PGs and tolerance of cluster map differences when they are not significant, and similar policies guide failure handling and updates. Unlike Palladio, however, RADOS manages cluster maps using a Paxos-like protocol among a dedicated cluster of monitors, and distributes updates with a lazy map distribution approach that is tightly integrated with the data access protocol. This simplifies consensus management, while also avoiding the distributed search and manager selection mechanisms required by Palladio to arbitrate access to healthy stores and recovery from failure. In contrast to both FAB and Palladio, RADOS provides high-performance access to data via a comparatively simple read and update protocol, and (of course) exposes an object-based (instead of block-based) interface.

Wiesmann *et al.* [105] compare replication in distributed systems and databases, including those that rely on *atomic broadcast* or *view synchronous broadcast* group communication primitives to simplify consistency management. RADOS does not utilize atomic broadcast or similar primitives because comparable ordering guarantees are provided by the use of a primary copy and ordered message delivery. RADOS's primary copy approach is analogous to eager primary copy database solutions like that in INGRES [91], although in contrast to replicated databases, replicated storage systems benefit from deterministic updates and inexpensive updates, while the processing associated with databases transactions suggests alternative approaches. This distinction is illustrated in a performance comparison of database replication strategies based on total order broadcast later conducted by Wiesmann *et al.* [106]: the performance of active replication schemes (in which each node processes the update) suffers rel-

ative to alternative (*e. g.* certification) schemes in which a single node calculates the result and broadcasts the results to replicas. Significantly, replication in RADOS further differs from primary-copy replication as described by Wiesmann in that load remains balanced despite a fixed (per-PG) primary due to declustering.

Version-driven consistency and logs are used in many primary-copy replication systems. Like RADOS, Harp [61] uses logs to facilitate efficient recovery from transient failures, and utilizes a limited form of decoupled replication (administrators must manually partition data into sub-volumes). However, although Harp relies on write-ahead logs and UPS to preserve both performance and consistency, RADOS simply delays update application on the readable replica, and dissociates synchronization from safety to preserve performance. Renesee *et al.* describe a recovery mechanism for chain replication [96] that utilizes a log-like structure, although their use of synchronous updates vastly simplifies possible failure modes, and does not address arbitrary changes in the data distribution. Other brick-based storage systems rely on a two-phase commit to maintain update consistency [107] that (like FAB's voting) incur a cost in latency and complexity. PRACTI replication [9] generalizes a range of consistency models and replica placement strategies, but like the systems above, does not address the complexity of doing so in a large-scale environment without a centrally managed metadata directory.

Sorrento utilizes version-based *consistency* model that targets an environment with minimal write sharing; when update conflicts do arise, conflicting transactions are rolled back by simply discarding newer versions. This optimistic approach simplifies updates in the general case at the expense of strong consistency, resulting in an interface in which application open-modify-close sequences may fail to commit and require a retry. In contrast, RADOS writes

are serialized and replicated by a primary OSD, avoiding concurrent update conflicts entirely simply by altering the flow of data.

Seneca [46] is an asynchronous remote-mirroring protocol that supports write coalescing and asynchronous propagation; these features are primarily useful when a slow, wide-area network link separates replicas and when replica consistency is allowed to diverge based on some bounded time interval. In contrast, RADOS is designed for a cluster environment with high speed links, and keeps all replicas consistent by applying updates synchronously. (Although write coalescing could be used to improve recovery performance, RADOS does not yet implement this optimization.) Ji *et al.* also present a taxonomy for remote mirroring which can be applied to a range of replication strategies; RADOS can be described as in-order asynchronous propagation, with no inter-LU (placement group) ordering, no divergence, and full write-through (although Ji's taxonomy does not capture RADOS' two-phase acknowledgement).

Previous systems have provided an abstract distributed object store. One of the earliest is Thor [60], which focused on extending the object (as in object-oriented programming) paradigm in a distributed environment with the notion of persistent objects. BuddyCache [10] provides strong consistency semantics and caching for distributed applications, but it targets small cooperating peer groups and wide area networks. Other systems relax consistency constraints [30] and target widely distributed and weakly connected environments [48, 53]. In contrast, RADOS targets a large, dedicated storage cluster.

Kybos [107] provides a distributed storage service using bricks and network RAID, with an emphasis on resource (storage and bandwidth) reservation and online adjustment of

placement to meet those requirements. In contrast, RADOS employs a pseudo-random distribution (corrected to limit variance and avoid overload) to avoid the metadata associated with defining explicit mappings between individual objects and devices, facilitating scalability and avoiding directory lookup. Kenchammana-Hosekote *et al.* evaluate a range of approaches to data path, consistency, and atomicity in a simulated network RAID system [49]. Although parity-based redundancy differs from simple replication, their analysis is partially applicable to RADOS. Notably, they observe that shifting serialization to an OSD avoids network saturation when client bandwidth is limited, and that the choice of data path is orthogonal to multi-object atomicity. In particular, a two-phase commit strategy for atomic multi-object updates—which is planned but not yet supported by RADOS—has little performance penalty in most circumstances.

Xin *et al.* [112] propose the use of distributed recovery as a means of improving data safety, and conduct a quantitative analysis of system reliability with FaRM, a declustered replication model in which—like RADOS—data objects are pseudo-randomly distributed among placement groups and then replicated by multiple OSDs, facilitating fast parallel recovery. They find that distributed recovery improves reliability at scale, particularly in the presence of relatively high failure rates for new disks (“infant mortality”). Lian *et al.* [59] find that reliability further depends on the number of placement groups per device, and that the optimal choice is related to the amount of bandwidth available for data recovery versus device bandwidth. Although both consider only independent failures, RADOS leverages CRUSH to mitigate correlated failure risk with failure domains.

6.7 Conclusions

RADOS provides a scalable and reliable object storage service without compromising performance. By separating serialization from safety, the architecture provides strong consistency semantics to applications by minimally involving clients in failure recovery.

RADOS utilizes a globally replicated cluster map that provides all parties with complete knowledge of the data distribution, typically specified using a function like CRUSH. This avoids the need for object lookup present in conventional architectures, which RADOS leverages to distribute replication, consistency management, and failure recovery among a dynamic cluster of OSDs while still preserving consistent read and update semantics. A scalable failure detection and cluster map distribution strategy enables the creation of extremely large storage clusters, with minimal oversight by the tightly-coupled and highly reliable monitor cluster that manages the master copy of the map.

Because clusters at the petabyte scale are necessarily heterogeneous and dynamic, OSDs employ a robust recovery algorithm that copes with any combination of device failures, recoveries, or data reorganizations. Recovery from transient outages is fast and efficient, and parallel re-replication of data in response to node failures limits the risk of data loss.

Chapter 7

Local Object Storage

Low-level object storage on individual OSDs is managed by EBOFS, an Extent and B-tree based Object File System. Although a variety of distributed storage architectures leverage existing general-purpose file systems [14, 55] such as ext3 [95], both the performance and the standard POSIX file system interface and safety semantics are inappropriate for RADOS. Because EBOFS is implemented entirely in user-space and interacts directly with a raw block device, it is unencumbered by an unwieldy kernel file system interface, and avoids interaction with the (Linux) VFS inode and page caches, which were designed around a different storage abstraction and workload assumptions. This allows EBOFS to optimize specifically for object workloads [97].

7.1 Object Storage Interface

EBOFS exposes a unique low-level storage interface that forms the basis of RADOS's strong consistency and safety model. Objects are accessed via a simple file-like interface that

provide simple read and write access to byte extents ($\langle offset, length \rangle$ pairs). Objects can be members of zero or more named *collections*, which are indexed to allow efficient enumeration or changes in membership (used by RADOS to manage placement groups). Both objects and collections accept named *attributes*, which map an identifier to an variable length (but normally small) piece of data.

In contrast to conventional file systems, EBOFS exposes an interface that supports compound transactions that allow a sequence of operations to be grouped into a single atomic operation.¹ A single transaction may, for instance, write data to multiple objects, adjust collection membership, update object and collection attributes, and be certain that after a failure the operation will be either fully completed or never started. Similarly, transactions can read both object data and attribute values atomically without fear of race conditions. This allows RADOS to apply an object update, modify the object version, update the PG *last_update*, and append an entry to the PG log in a single atomic transaction, keeping PG metadata perfectly synchronized with data.

EBOFS writes are usually non-blocking, applying changes to the in-memory cache and returning immediately. Because critical metadata is kept in memory, operations block only when the buffer cache is full and data is being flushed out to disk, allowing RADOS to quickly apply updates to establish ordering semantics without waiting for disk I/O. EBOFS provides asynchronous notification via a callback when changes are safely committed to stable storage.

¹Although many kernel-based file systems support compound transactions or group commit internally for efficiency, they do not expose a transaction interface to applications.

7.2 Data Layout

Like most modern file systems [40, 92], EBOFS uses *extents* to keep metadata efficient. In most cases, writes are laid out on disk in large, contiguous and singular extents. Free extents are binned by size and sorted, allowing the allocator to utilize a “closest good fit” strategy in which a free extent of approximately the size required is allocated near related data or the current write position on disk. This facilitates efficient streaming writes and future read access while minimizing long-term fragmentation.

A generalized B-tree [19] library is used to manage most storage metadata, including the free extent lists and the object and collection tables which map object identifiers to *onode* and *cnode* locations on disk. Each *onode* contains the metadata for an individual object, including attributes, collection membership, and the list of extents containing object data. Collection *cnodes* store collection attributes. A single large B-tree is currently used to map a collection ID to the objects it contains.

7.3 Data Safety

EBOFS maintains two superblocks, updated in an alternating fashion, to reference current file system metadata. All other data and metadata is always written to unallocated regions on disk, similar to soft updates [64] and copy-on-write B-tree updates in WAFL [40]. At the end of each *commit cycle*, all pending changes are flushed to disk before the next superblock is written with references to the new metadata. Pending changes are tracked such that subsequent writes do not block while waiting for prior epochs to commit. On mount, EBOFS simply

chooses the newer superblock, secure in the knowledge that it references a fully consistent snapshot of the most recently committed file system state.

In contrast to conventional file systems, which typically keep newly written data in cache before flushing it to disk (in the hopes that it will be modified again or deleted), Ceph clients perform this function before submitting I/Os to the object store. For this reason, EBOFS aggressively schedules disk writes, and cancels pending disk I/O when subsequent write or delete operations render them obsolete. This approach maximizes the length of our I/O queues for greater scheduling efficiency, and ensures that data reaches disk as quickly as possible.

7.4 Journaling

EBOFS can optionally employ an auxiliary journal to reduce the commit latency of object updates. Although the separation of serialization and safety in the storage interface largely masks the commit latency associated with the periodic cycle, certain applications (*e. g.* metadata journaling) rely on both the safety and low latency of writes.

The EBOFS journal is stored on an independent storage device—ideally, one backed by NVRAM—than the rest of the file system. The device is structured as a single ring buffer: once the write pointer reaches the end of the device, it starts over again at the beginning. The journal “tail” pointer is adjusted to reclaim journal storage after each commit cycle completes and older journal entries are no longer needed. If the journal device fills up prematurely, journaling is temporarily disabled for the duration of the commit cycle until it can be safely restarted.

Each update transaction processed by EBOFS is applied first to the in-memory cache,




Figure 7.1: Performance of EBOFS compared to general-purpose file systems. Although small writes suffer from coarse locking in the prototype, EBOFS nearly saturates the disk for writes larger than 32 KB. Since EBOFS lays out data in large extents when it is written in large increments, it has significantly better read performance.

then queued for write in the journal. The final commit callback occurs when the journal event flushes to the journal backing device, or when the next commit cycle completes—whichever comes sooner. For small writes in particular, this offers a significant improvement in performance. For large writes that are I/O bound at the disk, the duplication of data in both the journal and the primary storage device limits the journal’s effectiveness.

7.5 Evaluation

Figure 7.1 compares the performance of EBOFS to that of general-purpose file systems (ext3, ReiserFS, XFS) in handling a Ceph workload. Clients synchronously write out large files, striped over 16 MB objects, and read them back again. Although small read and write performance in EBOFS suffers from coarse threading and locking, EBOFS very nearly

saturates the available disk bandwidth for writes sizes larger than 32 KB, and significantly outperforms the others for read workloads because data is laid out in extents on disk that match the write sizes—even when they are very large. Performance was measured using a fresh file system. Experience with an early EBOFS design suggests it will experience significantly lower fragmentation than ext3, but I have not yet evaluated the current implementation on an aged file system.

7.6 Related Work

The design of EBOFS is based largely on the design of existing file systems. Most notably, these include XFS [92] and WAFL [40]. EBOFS, like XFS, utilizes extents and B-trees to manage allocation and metadata. EBOFS commit strategy is similar to that found in WAFL. Likewise, EBOFS’s journal is very similar to that used by WAFL in filers containing NVRAM. Notably, the loss of the journal (*e. g.* failure of an NVRAM battery) compromises the durability of writes during the last commit cycle, but does not affect the consistency of the rest of the file system.

Like FFS and most file systems that followed, EBOFS attempts to keep related data and metadata together. However, instead of employing explicit cylinder groups, EBOFS sorts the free extent maps by position to locate available storage space that is near related data or metadata.

Although the first incarnation of EBOFS was implemented within the Linux kernel, the implementation was moved to user-space largely based on experience with OBFS [97],

another object-based file system that demonstrated excellent performance outside the kernel by bypassing the kernel page cache with the `O_DIRECT` interface.

7.7 Conclusion

EBOFS provides a reliable, high-performance storage service for local object storage on Ceph OSDs. It is based on a copy-on-write strategy that avoids overwriting any data on disk until it has been deallocated, facilitating fast crash recovery by ensuring that the on-disk image is consistent at all times. Its storage interface tailored specifically to the needs of RADOS, notably including asynchronous notification of update commits to disk and support for compound transactions. This facilitates atomic compound updates to both data and metadata and to multiple objects, which are used by RADOS to keep PG metadata (such as the PG log, discussed in Section 6.3.4) consistent with object data. EBOFS includes support for a journal stored on an auxiliary storage device to facilitate fast update commits. It is implemented entirely in user space.

Chapter 8

Conclusion

I conclude this dissertation by first considering additional avenues of research, and then summarizing the work presented here.

8.1 Future Work

Although the design and implementation described in this work is relatively complete—the file system is functional and the implementation can properly tolerate most combinations of failures—there are a number of planned augmentations and opportunities for future research.

8.1.1 MDS Load Balancing

One of the largest lessons in Ceph was the importance of the MDS load balancer to overall scalability, and the complexity of choosing what metadata to migrate where and when. Although in principle the design and goals seem quite simple, the reality of distributing an evolving workload over a hundred MDSs highlights additional subtleties. Most notably, MDS

performance has a wide range of performance bounds, including CPU, memory (and cache efficiency), and network or I/O limitations, any of which may limit performance at any point in time. Furthermore, it is difficult to quantitatively capture the balance between total throughput and fairness; under certain circumstances unbalanced metadata distributions can increase overall throughput [102].

8.1.2 Client Interaction

The basic protocol used by the client to communicate with the MDS cluster takes a few basic steps to minimize interaction. For example, a *readdir* operation returns both dentry names and inode contents, facilitating a *readdirplus* interface or optionally relaxing strict client consistency to improve application performance. However, the operation of the MDS is still fundamentally synchronous: clients are not issued any leases to keep their metadata caches consistent. Experience with other file systems has shown that such techniques can be quite effective; in most workloads there are typically a small number of directories that are read-only and heavily shared (*e. g.* `/usr`).

A mechanism currently exists to delegate and call back client capabilities to facilitate exclusive, read-only, or read-write sharing of file data; it is likely that reusing this mechanism on a coarse per-directory basis could capture a relatively substantial improvement in metadata performance. A more generalized metadata leasing mechanism would likely prove more effective, although it would be more complicated to implement.

Systems like Envoy [79] seek to migrate metadata management to the same hardware node as the client application. Although this is technically feasible given the Ceph MDS ar-

chitecture, the security and failure recovery implications of doing so have not been carefully considered.

8.1.3 Security

The current implementation does not include any kind of distributed security—security enforcement resembles that of NFS, in which servers trust *uid* and *gid* values embedded in each client request. Although the architecture is based on a capability model to facilitate a strong security infrastructure, OSDs do not yet validate capabilities. Two security architecture and protocol variants have been proposed for Ceph [57, 69], and one has been partially implemented (for a single-MDS system).

8.1.4 Quotas

Pollack *et al.* [50] describe a scalable distributed quota enforcement architecture that is designed to work with distributed storage systems. Their architecture is based on cryptographically signed *vouchers*, generated by a quota server, that clients can redeem with OSDs in order to store data. Used vouchers are tracked and later reconciled, while the system further places bounds on clients’ ability to “cheat” through voucher re-use. Work is currently underway to implement this scheme in Ceph.

8.1.5 Quality of Service

Wu *et al.* [110] describe a quality of service framework designed to isolate the performance of different classes of workloads concurrently accessing the system. EBOFS is

augmented with QoS-aware disk scheduling, while the statistically uniform distribution of load provided by CRUSH is leveraged to approximate global class-based isolation. Research is ongoing in this area to further provide end-to-end performance management with stronger bounds on performance [24].

8.1.6 Snapshots

One of my ulterior motives for introducing the anchor table mechanism was to provide a generic means of making inodes globally addressable. Although this was necessary for support of multilink files, I also plan to use it to facilitate a flexible form of “snapshots” that can be applied at any time to arbitrary subtrees of the file system hierarchy. My hope is to introduce time extents to dentry identifiers in order to easily integrate management of namespace snapshots into the existing MDS, while simultaneously introducing an object-granularity versioning abstraction to the distributed object store.

8.2 Summary

As the scalability and performance requirements for storage systems increase, system designers must look to new architectures to meet those demands, in many cases abandoning conventional approaches. I have described Ceph, a distributed file system that provides excellent performance, reliability, and scalability. Ceph’s design eschews a number of conventions in file system design dating back to early Unix in order to push the bounds of reliable and scalable performance.

Metadata in Ceph is managed by a distributed, adaptive, fault-tolerant cluster of metadata servers (MDSs) that collectively forms a high-performance consistent cache of the file system namespace (Chapter 4). Metadata is stored with file metadata (inodes) embedded inside their containing directories, and each MDS maintains an extremely large journal, dramatically improving the I/O profile generated by a busy MDS server, while support for POSIX hard links is preserved through an auxiliary table. Workload is partitioned in terms of the file system namespace—made possible in part by the embedded inode strategy—allowing the cluster to adapt to changing file system workloads to effectively utilize available resources.

Scalability and self-management in RADOS is made possible in part through the use of CRUSH, and specialized data distribution function that logically takes the place of a conventional allocation table. CRUSH (described in Chapter 5) allows the location of a data object to be cheaply calculated when it is needed, eliminating the need to store object locations in a table or query a directory service. CRUSH provides functionality similar to a hash function, while taking additional steps to ensure that object locations are stable when devices fail or join the cluster. Most importantly, the algorithm addresses reliability in the presence of correlated failures through the use of a flexible, hierarchical placement strategy.

RADOS (Chapter 6) enables the management of a large, dynamic cluster of storage devices by distributing data replication, failure detection, and failure recovery to intelligent, semi-autonomous devices. A cluster map ensures a consistent view of cluster state and data layout, facilitating consistent read and update semantics despite a lazy update propagation scheme. This allows replication, failure detection, and recovery to be managed in a decentralized and scalable fashion, with minimal oversight from the tightly-coupled cluster of monitors responsi-

ble for the master copy of the cluster map. The update protocol separates synchronization from safety in order to improve performance while preserving strong consistency and data safety semantics, enabling excellent performance, data safety, and scalability in a dynamic cluster environment.

In addition to describing key enabling components of Ceph's design, I have evaluated a working prototype of the system under a range of benchmarks and real-world workloads, demonstrating excellent performance, reliability, and system scalability.

Appendix A

Communications Layer

A.1 Abstract Interface

Message passing in Ceph is based on a simple abstract messaging interface that defines node addressing and naming, methods for sending and receiving messages, and a notification mechanism for handling delivery failures.

Each node in Ceph has both a logical name and a network address associated with it. The logical name (*e. g.* mds0, osd1) will typically remain the same across restarts, as it refers to the node's role in the system, while the address is (by design) unique for every incarnation of a daemon or process participating in the system. The address consists of the IP address and port used for sending messages to the given node, as well as a (typically random) nonce value to keep the address unique, even across starts. This allows other nodes in the system to easily determine whether a peer has restarted (and potentially lost shared state) by comparing addresses.¹

¹In certain cases, the nonce values are not used. The monitors' address nonce values are fixed at 0 since commu-

A.2 SimpleMessenger

SimpleMessenger is an implementation of the abstract communications interface that utilized TCP for message transport. The use of TCP naturally provides reliable, ordered delivery as well as notification of communication faults (in the form of connection drops). When a message is queued for a new peer, a TCP session is established, and the peers identify themselves to ensure that their nonce values match and the address is not stale. Once established, messages can be passed across the same TCP session in both directions. An orderly connection teardown process is used to shut down the session to ensure that no messages are lost.

The implementation currently makes no attempt to retry in the event of a communications error, which makes it relatively vulnerable to certain types of intermittent network problems. For example, packet loss is tolerated relatively well (eventually causing a timeout), while routing changes can cause an immediate TCP session drop. Although explicit message receipt acknowledgement and buffering of outgoing messages would allow the implementation to transparently attempt reconnect and redelivery without violating the ordering requirements, this has not yet been implemented.

A.3 FakeMessenger

FakeMessenger is an alternate implementation of the messaging interface that is used for debugging purposes. Messages are exchanged between logical entities existing within a single process's address space. Because no inter-process communication mechanism is provided,

nication with the monitors is stateless.

multiple logical system components (OSDs, metadata servers, clients, monitors) are compiled and executed within the same process. Sent messages are simply added to in-memory queues, and later delivered in a round-robin fashion that ensures that only a single message is begin processed at a time (for ease in debugging).

Appendix B

MDS Implementation Details

This appendix describes certain elements of the MDS implementation in greater detail. In particular, I describe the structure of the distributed cache and the mechanisms through which replication consistency is preserved in the face of failure.

B.1 MDS Distributed Cache

Central to the dynamic subtree partitioning approach is the treatment of the file system as a hierarchy. The file system is partitioned by delegating authority for subtrees of the hierarchy to different metadata servers. Delegations may be nested: `/usr` may be assigned to one MDS, for instance, while `/usr/local` is reassigned to another. In the absence of an explicit subtree assignment, however, the entire directory tree nested beneath a point is assumed to reside on the same server.

Implicit in this structure is the process of hierarchy traversal in order for nested inodes to be located and opened for subsequent descent into the file hierarchy. Such path traversal

is also necessary to verify user access permissions for nested items as required by POSIX semantics. Although this process may seem costly for locating a file deep within the directory hierarchy, the locality of reference typical of both scientific and general purpose computing workloads [28, 98] allows those costs to be amortized over subsequent accesses to the same directories. More importantly, unlike LH permission management [15], a hierarchically defined structure allows the system to move or change the effective permissions of arbitrarily sized subtrees of the directory tree by modifying the relevant ancestor directory with fixed cost. Likewise, individual subtrees of the hierarchy are fully independent from their siblings; semantics are dependent only on the prefix (ancestor) directories leading to the root of the file system.

B.1.1 Cache Structure

All metadata that exists in the cache is attached directly or indirectly to the root inode¹. That is, if the `/usr/bin/vi` inode is in the cache, then `/usr/bin`, `/usr`, and `/` are too, including the inodes, directory objects, and dentries. Only leaf items may be expired from the cache; directories may not be removed until items contained within them are expired first. This allows permission verification for all known items to proceed without any additional I/O costs, and for hierarchical consistency to be preserved. It also facilitates the embedding of inodes in directories.

¹Metadata may also be rooted by a *stray* inode; see section B.1.1.3.

B.1.1.1 Directory Fragments

Directory contents are managed in terms of one or more directory fragments (*dir frags*) that are associated with each directory inode. In most cases, there is a single fragment that contains the full directory contents. In certain cases, however, directory contents are broken up to ease load balancing. Each directory fragment is stored in a different object in the object store, allowing them to be managed completely independently of each other.

Fragments are described by a binary split tree (a *frag tree*) stored in the directory inode. Each directory fragment corresponds to a *frag*, which is defined by a bit pattern and mask. Much like network names and masks in IP networking, the mask specifies which bits are significant and compared to the bit pattern. For example, frag 12/4 is a 4 bit mask and a matching bit pattern of 12 (1100 in binary). Directory entries are mapped into fragments by hashing the filename to an integer value and then matching that against all fragments in the fragtree.²

Because the cache structure is defined such that a directory inode has any number of children (depending on its fragmentation level), the subtree partitioning mechanisms can be leveraged to individually delegate directory fragments to other nodes in the cluster. This allows all of the existing infrastructure for migration and failure recovery to be re-used, while providing a simple internal abstraction for breaking large or busy directory into smaller pieces.

²The implementation is considerably more efficient than this sounds; mapping a hash value to a fragment in a fragtree is $O(\log n)$, where n is the number of fragments.

B.1.1.2 Dentries and Inodes

Each directory fragment is a collection of directory entries, or dentries. Because Ceph embeds inodes into the directories that reference them, this gives rise to two types of dentries. The first (or only) dentry to reference an inode is called the *primary*, and is always accompanied by the inode itself. *Remote* dentries reference an inode by number only; an additional lookup in the anchor table may be required to locate the inode itself. A *null* dentry is also possible (at least in-memory) to capture caching or locking state for a name that does not (yet) reference an inode but is still of interest (*e. g.* for replicating knowledge of a name's non-existence), or to capture the *dirty* state for a deleted file.

B.1.1.3 The “Stray” Directory

Because the cache is structured such that all inodes are associated with a primary dentry, management of files that are open for I/O but have been unlinked from the namespace presents a problem. Maintaining the primary dentry/inode relationship is desirable because of the subtree-based approach for partitioning workload and the existing metadata storage infrastructure.

To maintain that arrangement, unlinked files are moved into a hidden *stray* directory if it is not possible for them to be immediately removed. Each MDS maintains a separate stray directory inode (with potentially many directory fragments), allowing operations that unlink files from the namespace to proceed locally on each MDS. Stray directory inodes have predictable numbers (based on the MDS node), allowing discovery and replication of stray metadata on other MDS nodes, and allowing unlinked inodes to be referenced by the anchor table.

B.1.2 Replication and Authority

The authority maintains a list of what nodes cache each inode. Additionally, each replica is assigned a nonce (initial 0) to disambiguate multiple replicas of the same item (see below).

```
map<int, int> replicas; // maps replicating mds# to nonce
```

The replicas set *always* includes all nodes that cache the particular object, but may additionally include nodes that used to cache it but recently trimmed it from their cache. In those cases, an expire message should be in transit. We have two invariants:

1. The authority's replica set will always include all actual replicas, and
2. cache expiration notices will be reliably delivered to the authority.

The second invariant is particularly important because the presence of replicas will pin the metadata object in memory on the authority, preventing it from being trimmed from the cache. Notification of expiration of the replicas is required to allow previously replicated objects to eventually be trimmed from the cache as well.

Each metadata object has a authority bit that indicates whether it is authoritative or a replica. Although this information is also discernible from the subtree partition, the bits are faster to check and provide an additional sanity check for debugging.

Each replicated object maintains a "nonce" value, issued by the authority at the time the replica was created. If the authority has already created a replica for the given MDS, the new replica will be issues a new (incremented) nonce. This nonce is attached to cache expirations,

and allows the authority to disambiguate expirations when multiple replicas of the same object are created and cache expiration is coincident with replication. (In certain cases, the authority may push a new replica to another MDS to ensure that a replica exists there for some operation.) That is, when an old replica is expired from the replicating MDS at the same time that a new replica is issued by the authority and the resulting messages cross paths, the authority can tell that it was the old replica that was expired and effectively ignore the expiration message. A replica is removed from the authority's replicas map only if the nonce matches.

B.1.3 Subtree Partition

Authority of the file system namespace is partitioned using a subtree-based partitioning strategy. This strategy effectively separates directory inodes from directory contents, such that the directory contents are the unit of re-delegation. That is, if `/` is assigned to `mds0` and `/usr` to `mds1`, the inode for `/usr` will be managed by `mds0` (it is part of the `/` directory), while the contents of `/usr` (and everything nested beneath it) will be managed by `mds1`.

The description for this partition exists solely in the collective memory of the MDS cluster and in the individual MDS journals. It is not described in the regular on-disk metadata structures. This is related to the fact that authority delegation is a property of the *directory fragment* and not the directory's *inode*.

Subsequently, if an MDS is authoritative for a directory inode and does not yet have any state associated with the directory in its cache, then it can assume that it is also authoritative for the directory.

Directory state consists of a data object `CDir` that describes any cached dentries con-

tained in the directory, information about the relationship between the cached contents and what appears on disk, and any delegation of authority. Each CDir object has a `dir_auth` element. Normally `dir_auth` has a value of `AUTH_PARENT`, meaning that the authority for the directory is the same as the directory's inode. When `dir_auth` specifies another metadata server, that directory is point of authority delegation and becomes a *subtree root*. A CDir is a subtree root IFF its `dir_auth` specifies an MDS id (and is not `AUTH_PARENT`). That is,

1. A dir is a subtree root IFF `dir_auth != AUTH_PARENT`.
2. If `dir_auth = AUTH_PARENT` then the inode `auth == dir_auth`, but the converse may not be true.

The authority for any metadata object in the cache can be determined by following the parent pointers toward the root until a subtree root CDir object is reached, at which point the authority is specified by its `dir_auth`.

Each MDS cache maintains a subtree data structure that describes the subtree partition for all objects currently in the cache:

```
map< CDir*, set<CDir*> > subtrees;
```

A dirfrag (represented by a CDir) will appear in the subtree map (as a key) IFF it is a subtree root. The map value is a set of all other subtree roots nested immediately beneath that point. Nested subtree roots effectively bound or prune a subtree. For example, if we had the following partition:

MDS	Path
mds0	/
mds1	/usr
mds0	/usr/local
mds0	/home

The subtree map on mds0 would be

Subtree	Bounds
/	(/usr, /home)
/usr/local	()
/home	()

and on mds1:

Subtree	Bounds
/usr	(/usr/local)

B.2 Metadata Storage

Ceph metadata is stored in regular objects in the shared object store. This is advantageous primarily because it keeps all metadata in a shared medium, facilitating migration of metadata authority between nodes and recovery in the event of failure. Metadata resides both in the per-MDS journals and primary per-directory fragment objects (see Chapter 4).

B.2.1 Directory Fragments and Versioning

Metadata updates that exist in the MDS journal but not in the regular metadata objects are called *dirty* and pinned in the MDS cache. In order to keep track of which updates have been committed, each inode, dentry, and dir object in the MDS cache maintains a version value. Versions are generated relative to the directory fragment that contains them, since that is the underlying unit for metadata storage: when a cache item is dirtied, its new version value is generated by incrementing the directory fragment version.

Each directory fragment maintains four version values: *version*, the current version; *projected_version*, the anticipated version pending updates that are still being journaled; and *committing_version* and *committed_version*, the last versions to be queued for or commit to

disk. The projected version values are assigned to updates that are currently being written to the journal, but have not yet been applied to the in-memory metadata cache (and whose side-effects are thus not yet visible). Thus, *committed* \leq *committing* \leq *version* \leq *projected*.

The committed version normally indicates the version of the directory that is written to stable storage on disk. Any directory entry or inode contained in the directory with a version value greater than this should be flagged as dirty in the cache. After journal replay due to an MDS recovery from a failure, however, the MDS does not know the committed version of directory fragments it extracted from the journal (it would be expensive to probe all of these objects just to find out). In this case, the committed version is zero, and many cache items may be marked as dirty (due to journaled updates) even though they were safely stored by the prior incarnation of the MDS. When such a directory is *fetch*ed from disk, it will compare dirty dentry versions with the actual committed version and mark any unnecessarily dirty items clean.

When a directory item is removed, it is replaced by a (dirty and versioned) *null* dentry. This ensures that the deletion is reflected in the cache until the directory is committed. Clean null dentries can be safely trimmed.

Each directory fragment has a *complete* flag that indicates whether all directory contents are currently cached. If any non-null item is expired from the cache, the flag is cleared. Currently, directory fragments can only be committed to disk in their entirety, which means that a fetch must sometimes be conducted (to fill in any missing contents) before a commit (*i. e.* a read-modify-write). Partial commits will be made possible in the future by extending the object interface to provide a key/value interface in addition to the simple byte extent (file-like) model. A higher-resolution complete flag (*e. g.* one with a separate bit for different subsets of the den-

try namespace, or a Bloom filter [11]) would allow create operations to safely proceed without caching the complete directory fragment contents.

B.2.2 Journal Entries

Journal entries reflect metadata updates that have not yet been *committed* to the per-directory metadata storage. Entries employ a common *metablob* structure for describing metadata updates. Each metablob consists of one or more directory fragment IDs, dentries, and inodes. Normally, each updated item in the metablob is accompanied by any ancestor metadata necessary to connect it to the root of the current subtree. Journal trimming is constrained such that at all times any non-obsolete journal entry is preceded by a SubtreeMap entry that includes all subtree roots and bounds, with ancestor metadata up to the filesystem root. Thus, as long as replay begins with a SubtreeMap, all included metadata can be placed properly within the hierarchy.

Metadata objects in the metablob also include a few flags: namely, all items may be marked *dirty*, and directory fragments may be marked *complete*. The *dirty* flag simply sets the object's dirty flag when the entry is replayed after a failure, inducing an eventual commit of the containing directory fragment. The complete flag likewise sets the fragment's complete flag; it is used only when migrating subtrees between MDS nodes when the fragment's entire contents are included in the metablob.

A few other metadata updates can be described by the metablob, allowing them to be committed atomically with other updates. These include anchor table transactions (see below), inode number allocations, inode truncations, and client request identifiers (which are journaled

to keep all operations idempotent from the perspective of the client).

B.2.3 Anchor Table

The anchor table is an auxiliary structure, managed by a single MDS, that allows inodes to be located within the directory hierarchy by their inode number. It is necessary because Ceph lacks a conventional inode table that facilitates inode lookup by inode number. Inode contents are always embedded in directories adjacent to a directory entry that references them. Because they can exist only in one place, that dentry is deemed the *primary dentry*. Any additional *remote dentries* refer simply to the inode number.

Because there are no back pointers from the inode to reference remote dentries, and because directory renames can affect arbitrarily large portions of the hierarchy, it is necessarily to be able to locate the actual inode contents within the hierarchy with only an inode number. This is accomplished by *anchoring* only those inodes who have multiple hard links or directories inodes with children having multiple hard links. The anchor table maintains back pointers with reference counts for all anchored inodes to the directory fragment that contains them.

For example, if `/usr/bin/nano` and `/usr/lib/ld.so` have additional hard links, the table might look like Figure B.1.

This simple structure for the anchor table keeps the overall size of the table relatively small (particularly when considering that in most cases, very few inodes have multiple hard links). It also is easy to update when directory renames effect large portions of the hierarchy: only the backpointer for the renamed item, the reference count for its immediate old and new parents, and any missing ancestors need to be updated in the table.

Path	Inode	Parent	Ref
/usr/bin/nano	123	12#0	1
/usr/bin	12	10#0	1
/usr/lib/ld.so	456	11#0	1
/usr/lib	11	10#0	1
/usr	10	1#0	2
/	1		1

Table B.1: Sample anchor table with two anchored inodes. Note that the full path is shown only for illustration and is not part of the table. The hash mark (#) in the parent column is used to denote a particular fragment of the given directory, where 0 indicates the directory is not fragmented. The /usr entry has a ref count of two because it is referenced by both /usr/bin and /usr/lib.

Anchor table queries return all records necessary to reach the root for the given inode. Updates are conducted with a two-phase commit. The anchor table MDS first journals a *prepare* event. The resulting transaction ID is included in the journaled metablob for the atomic operation effecting the update, and a *commit* event is later journaled by the anchor table MDS to close the transaction.

B.3 Migration

B.3.1 Cache Infrastructure

B.3.1.1 Ambiguous Authority

While metadata for a subtree is being migrated between two MDS nodes, the `dir_auth` for the subtree root is allowed to be ambiguous. That is, it will specify both the old and new MDS IDs to indicate that a migration is in progress.

If a replicated metadata object is expired from the cache from a subtree whose author-

ity is ambiguous, the cache expiration is sent to both potential authorities. This ensures that the message will be reliably delivered, even if either of those nodes fails. A number of alternative strategies were considered. Sending the expiration to the old or new authority and having it forwarded if needed can result in message loss if the forwarding node fails. Pinning ambiguous metadata in cache is computationally expensive for implementation reasons, and delaying the transmission of expiration messages is difficult to implement because the replicating MDS must send the final expiration messages only when the subtree authority is disambiguated, forcing it to keep certain elements of it cached in memory. Although duplicating expirations incurs a small communications overhead, the implementation is much simpler and easier to verify.

B.3.1.2 Auth Pins

Most operations that modify metadata must allow some amount of time to pass in order for the operation to be journaled or for communication to take place between the object's authority and any replicas. For this reason it must not only be pinned in the authority's metadata cache, but also be locked such that the object's authority is not allowed to change until the operation completes. This is accomplished using *auth pins*, which increment a reference counter on the object in question, as well as all parent metadata objects up to the root of the subtree. As long as the pin is in place, it is impossible for that subtree (or any fragment of it that contains one or more auth pins) to be migrated to a different MDS node. Auth pins can be placed on inodes, dentries, and directories.

Auth pins can only exist for authoritative metadata, because they are only created if the object is authoritative, and their presence prevents the migration of authority.

B.3.1.3 Freezing

More specifically, auth pins prevent a subtree from being *frozen*. When a subtree is frozen, all updates to metadata are forbidden. This includes updates to the replicas map that describes which replicas (and nonces) exist for each object.

In order for metadata to be migrated between MDS nodes, it must first be frozen. The root of the subtree is initially marked as *freezing*. This prevents the creation of any new auth pins within the subtree. After all existing auth pins are removed, the subtree is then marked as *frozen*, at which point all updates are forbidden. This allows metadata state to be packaged up in a message and transmitted to the new authority, without worrying about intervening updates.

If the directory at the base of a freezing or frozen subtree is not also a subtree root (that is, it has `dir_auth == AUTH_PARENT`), the directory's parent inode is auth pinned. A frozen tree root dir will auth_pin its inode IFF it is authoritative AND not a subtree root. This prevents a parent directory from being concurrently frozen, and a range of resulting implementation complications relating to determining the bounds of the frozen region.

B.3.1.4 Cache Expiration for Frozen Subtrees

Cache expiration messages that are received for a subtree that is frozen are temporarily set aside instead of being processed. Only when the subtree is unfrozen are the expirations either processed (if the MDS is authoritative) or discarded (if it is not). Because either the exporting or importing metadata can fail during the migration process, the MDS cannot tell whether it will be authoritative or not until the migration completes.

During a migration, the subtree will first be frozen on both the exporter and importer,

and then all other replicas will be informed of a subtree's ambiguous authority. This ensures that all expirations during migration will go to both parties, and nothing will be lost in the event of a failure.

B.3.2 Normal Migration

The exporter begins by doing some checks in `export_dir()` to verify that it is permissible to export the subtree at this time. In particular, the cluster must not be degraded, the subtree root may not be freezing or frozen, and the full path must be read locked (*i. e.* not conflicted with a rename). If these conditions are met, the subtree root directory is temporarily auth pinned, the subtree freeze is initiated, and the exporter is committed to the subtree migration, barring an intervening failure of the importer or itself.

An `MExportDiscover` message sent from the exporter to the importer serves simply to ensure that the inode for the base directory being exported is open on the destination node. It is pinned by the importer to prevent it from being trimmed. This occurs before the exporter completes the freeze of the subtree to ensure that the importer is able to replicate the necessary metadata. When the exporter receives the `MDiscoverAck`, it allows the freeze to proceed by removing its temporary auth pin.

The `MExportPrep` message then follows to populate the importer with a spanning tree that includes all directories, inodes, and dentries necessary to reach any nested subtrees within the exported region. This replicates metadata as well, but it is pushed out by the exporter, avoiding deadlock with the regular discover and replication process. The importer is responsible for opening the bounding directory fragments from any third parties authoritative for those

subtrees before acknowledging. This ensures that the importer has correct `dir_auth` information about where authority is re-delegated for all points nested beneath the subtree being migrated. While processing the `MExportPrep`, the importer freezes the entire subtree region to prevent any new replication or cache expiration.

A *warning* stage occurs only if the base subtree directory is open by nodes other than the importer and exporter. If it is not, then this implies that no metadata within or nested beneath the subtree is replicated by any node other than the importer and exporter. If it is, then a `MExportWarning` message informs any bystanders that the authority for the region is temporarily ambiguous, and lists both the exporter and importer as authoritative MDS nodes. In particular, bystanders who are trimming items from their cache must send `MCacheExpire` messages to both the old and new authorities (see above). Lock-related messages are also delayed until the authority is no longer ambiguous.

The exporter walks the subtree hierarchy and packages up an `MExport` message containing all metadata and important state (*e. g.* information about metadata replicas). At the same time, the exporter's metadata objects are flagged as non-authoritative. The `MExport` message sends the actual subtree metadata to the importer. Upon receipt, the importer inserts the data into its cache, marks all objects as authoritative, and logs a copy of all metadata in an `EImportStart` journal message. Once that has safely flushed to the journal, it replies with an `MExportAck`. The exporter can now log an `EExport` journal entry, which ultimately specifies that the export was a success. In the presence of failures, it is the existence of the `EExport` entry only that disambiguates authority during recovery.

Once logged, the exporter will send an `MExportNotify` to any bystanders, informing

them that the authority is no longer ambiguous and cache expirations should be sent only to the new authority (the importer). Once these are acknowledged back to the exporter, implicitly flushing the bystander to exporter message streams of any stray expiration notices, the exporter unfreezes the subtree, cleans up its migration-related state, and sends a final MExportFinish to the importer. Upon receipt, the importer logs an EImportFinish(true) (noting locally that the export was indeed a success), unfreezes its subtree, processes any queued cache expirations, and cleans up its state.

B.4 Failure Recovery

MDS recovery after a failure (*e. g.* host crash, process segfault, sufficiently long network outage) is made possible by the journal kept by each MDS. Journal replay is complicated, however, because not all state is written to the journal. Although all metadata updates are logged, the state of the distributed cache (*i. e.* which nodes have in-memory replicas of which metadata) is not. Furthermore, some metadata operations involve multiple metadata objects on different MDS nodes, requiring a two-phase commit and a resolution stage during recovery.

The full recovery process is broken down into four stages. During the *replay* stage, the MDS simply re-reads the contents of the journal, accumulating state in memory. During the *resolve* stage, the fate of any two-phase updates are determined. During *reconnect*, client sessions are reestablished. Finally, during the *rejoin* stage, MDS nodes exchange information about what metadata objects are replicated with peers in the cluster, reestablishing distributed cache and lock state.

Certain stages of the recovery occur with respect to the *recovery set*, the set of MDS nodes that were participating in the cluster at the time of the first failure. Once one failure has occurred, the cluster is considered *degraded*, and no new MDS nodes are allowed to join (unless they are taking over for a failed node) or leave until all nodes have fully recovered.

B.4.1 Journal Replay

Journal replay is relatively straightforward. The Journaler interface is used to read and write to a journal constructed with objects in the distributed object store, and probes the journal size on its own. Replay begins at the last known expire point (the journal's size metadata is only written periodically). Each journal event is read from the journal in sequence, and its `replay()` method is called to recover its state. The only caveat to replay is that events prior to the first SubtreeMap event are ignored. Once the first subtree map is replayed, the remainder of the journal is processed in its entirety.

When replay completes, the MDS moves to the *resolve* stage, unless it is the only MDS in the cluster, in which case it moves directly to the *reconnect* stage.

B.4.2 Resolve Stage

The *resolve* stage serves to disambiguate the fate any operations that span multiple metadata servers. These include imports of subtrees of metadata (migrations from other MDS nodes to the current node), updates to the anchor table, and client operations like *rename* or *unlink* that sometimes affect metadata managed by different MDS nodes.

An MDS entering the resolve stage begins by sending a Resolve message to all other

resolving nodes in the recovery set. Surviving MDS nodes also send a Resolve message to any MDS that enters the resolve stage. The Resolve message includes a summary of all subtrees currently managed by that node (specified as the root dirfrag and its list of bounds, if any), a list of ambiguous imports (subtrees that were partially imported when the failure occurred), and a list of uncommitted slave request IDs (for updates initiated by the target MDS, but affecting metadata on the recovering node). Ambiguous imports and uncommitted slave requests are defined by the presence of a *prepare* event in the journal but no matching *committed*, or (in the case of a surviving MDS sending a Resolve to a recovering MDS) by the corresponding in-memory state.

The recovering MDS waits for Resolve messages from all other MDS nodes, assimilating information about subtree authority for metadata it has recovered from the journal by setting its CDir dir_auth values appropriately. If the Resolve message lists uncommitted slave requests, a ResolveAck message is sent in response that specifies which requests committed and aborted. Ambiguous imports are noted, but not processed until all Resolve messages have been processed. At that point, subtree authority ambiguities are resolved by simply checking if authority for an ambiguous subtree is claimed by another node: if it is, the import clearly didn't complete.

Once subtree authority has been resolved, the recovering node trims all non-authoritative metadata from its cache, with the exception of that necessary to connect authoritative subtrees to the root. On recovery, the MDS has no way of knowing whether non-authoritative metadata was subsequently updated, requiring all such metadata to be revalidated during the rejoin stage; trimming non-authoritative metadata reduces the amount of metadata that must be exchanged,

allowing that process to complete more quickly, and significantly simplifies the process.

B.4.3 Reconnect Stage

The recovering MDS next reestablishes connectivity with any clients with whom it had open sessions at the time of the crash. It announces itself by sending each client a copy of the current MDSMap (which specifies that the sending MDS is in the reconnect state). In response, each client sends an MClientReconnect, which lists any open files, their full pathname, and issued capabilities. This is necessary primarily because the MDS does not (synchronously) journal file opens. Although the file inode number is often sufficient to restore state, the full path name (noted by the client at the time the file was successfully opened) is included in case the inode was not replayed from the journal (*e. g.* , it was not recently modified). The MDS updates the inode *file lock* state such that it is compatible with the currently issued capabilities (*e. g.* if multiple clients have the read and write capability bits, the lock is put in the *mixed* state).

B.4.4 Rejoin Stage

During the *rejoin* stage, state about locks and the replication of metadata in memory is reestablished between MDS nodes. This replication state is necessary for the distributed MDS cache and locking to function, but is too expensive to journal. Rejoin involves both recovering MDS nodes (those that crashed and have replayed their journal) and those that have survived (not failed).

A recovering MDS begins the rejoin process only after all other recovering nodes also

reach the rejoin stage. At that point, each MDS sends a message summarizing what metadata it replicates. Such messages come in two varieties: *weak* or *strong*, depending on whether it is a recovering or surviving MDS, respectively. A recovering node expects a rejoin (either strong or weak) from every other node in the cluster. Surviving MDS nodes send strong rejoin messages to each recovering node only (*i. e.* not to other survivors).

B.4.4.1 Weak and Strong Rejoin

The initial rejoin message is a declaration of replication, sent from an MDS replicating metadata to the metadata's authoritative MDS. A *weak rejoin* is sent by nodes recovering from a failure. Rejoin messages are generated by walking all subtrees in the node's subtree map for which it is not authoritative, and declaring any replicated metadata object in the message bound for its authority. The declarations consist of dirfrag identifiers (inode number and frag), dentry names and types, and inode numbers. The weak rejoin also includes a list of all locally open files and capabilities that clients have declared during the reconnect stage that do not fall within subtrees the recovering node is authoritative for.

A *strong rejoin* is sent by surviving MDS nodes, and also includes the lock states associated with each replicated object, any capabilities that are wanted by clients with locally opened files, and any slave auth_pins or xlocks held by client requests currently being processed. The strong rejoin does not include an enumeration of open files.

A recovering node may thus receive a combination of weak and strong rejoin messages, depending on whether its peers also suffered a failure. A surviving node will only receive weak rejoins (from recovering nodes); surviving nodes have lost no shared state and do not re-

join with each other.

When processing a weak rejoin, the MDS will check for any capability exports it should claim. (Such client capabilities are migrated to the authoritative MDS much like they are during a subtree migration.) If the inode does not already exist in the cache, it will add the path to a list of inodes to load (see below; surviving nodes will always have the inode in their cache). The MDS will then walk the list of replica declarations, adding the sender to the appropriate metadata object replica maps. If the node is recovering, the dirlock on directories will also be set to the SCATTER state to ensure that any remote mtime updates are captured.

If a surviving MDS node receives a weak rejoin (the sender is therefore a recovering node), the rejoin consists of an exhaustive list of items replicated by the sender. The recipient (survivor) takes the additional step of scouring its own cache for objects that were previously replicated by the sender node but were not recovered from the journal, and adjusting their replica maps accordingly.

Once a recovering node receives a rejoin from all other nodes, it fetches any metadata for previously issued (or imported) capabilities from disk. Capability imports are then processed (by updating the issued state and notifying the client). Finally, rejoin acks are generated for all replicated metadata in the cache.

B.4.4.2 Ack

An *ack* message is generated in response to each weak or strong rejoin. The ack contains a newly issued nonce for each replicated object, as well as the initial states to initialize replica locks.

If the rejoin recipient is a survivor, it can immediately respond with an ack, because metadata object locks are in a known correct state. If the rejoin recipient is a recovering node, however, it must first wait until all rejoins have been received before responding, in case any of its peers includes a strong declaration that forces a lock into a non-default state.

Once a recovering node receives all expected rejoins and acks, it moves from the *rejoin* to the *active* state.

B.4.4.3 Missing and Full

It is possible that a strong rejoin will include declarations for replicated metadata objects that are not in the recovering node's local cache (*i. e.* were not in the journal). In this case, it is less expensive to acquire those metadata objects from the surviving MDS node than from disk. To allow the rejoin to proceed smoothly, the recovering node simply creates any metadata objects that are missing, and flags inodes to indicate their contents are (as yet) undefined. The recovering node generates a *missing* message for the sender that lists any inodes it lacked.

The surviving node responds with a *full* message that includes the full inode contents (inode_t struct, symlink target, and/or dirfrag tree).

B.5 Anchor Table

B.5.1 Table Updates

Anchor table updates are based on a two phase commit. The MDS initiating an update sends a *prepare* request to the anchor table MDS. The prepare is identified by the inode number and operation type; only one operation for each type (create, update, destroy) can be pending per inode at any time. Both parties may actually be the same MDS node, but for simplicity we treat that situation the same. (That is, we act as if they may fail independently, even though they can't.)

The anchor table journals the proposed update, and responds with an *agree* message and an anchor table version number. This uniquely identifies the request.

The initiating MDS can then update the file system metadata however it sees fit (*e. g.* to perform an *unlink* or *rename*). When it is finished and the operation has been journaled, it sends a *commit* message to the anchor table. The table journals the commit, frees any state from the transaction, and sends an *ack*. The initiating MDS should then journal the *ack* to complete the transaction.

B.5.2 Failure Recovery

B.5.2.1 Anchor Table MDS Failure

If the anchor table fails before journaling the *prepare* and sending the *agree*, the initiating MDS will simply retry the request.

If the anchor table fails after journaling *prepare* but before journaling *commit*, it will

resend *agree* to the initiating MDS.

If the anchor table fails after the *commit*, the transaction has been closed, and it takes no action. If it receives a *commit* for which it has no open transaction, it will reply with *ack*.

B.5.2.2 Initiating MDS Failure

If the MDS fails before the metadata update has been journaled, no action is taken, since nothing is known about the previously proposed transaction. If an *agree* message is received and there is no corresponding *prepare* or pending-commit state, and *rollback* is sent to the anchor table.

If the MDS fails after journaling the metadata update but before journaling the *ack*, it resends *commit* to the anchor table. If it receives an *agree* after resending the *commit*, it simply ignores the *agree*. The anchor table will respond with an *ack*, allowing the initiating MDS to journal the final *ack* and close out the transaction locally.

On journal replay, each metadata update (metablob) encountered that includes an anchor transaction is noted in the anchor table client by adding it to the pending-commit list, and each journaled *ack* is removed from that list. Journal replay may encounter *acks* with no prior metadata update; these are ignored. When recovery finishes, a *commit* is sent for all outstanding transactions.

Bibliography

- [1] Michael Abd-El-Malek, William V. Courtright II, Chuck Cranor, Gregory R. Ganger, James Hendricks, Andrew J. Klosterman, Michael Mesnier, Manish Prasad, Brandon Salmon, Raja R. Sambasivan, Shafeeq Sinnamohideen, John D. Strunk, Eno Thereska, Matthew Wachs, and Jay J. Wylie. Ursa minor: versatile cluster-based storage. In *Proceedings of the 4th USENIX Conference on File and Storage Technologies (FAST)*, pages 59–72, San Francisco, CA, December 2005.
- [2] Atul Adya, William J. Bolosky, Miguel Castro, Ronnie Chaiken, Gerald Cermak, John R. Douceur, John Howell, Jacob R. Lorch, Marvin Theimer, and Roger Wattenhofer. FAR-SITE: Federated, available, and reliable storage for an incompletely trusted environment. In *Proceedings of the 5th Symposium on Operating Systems Design and Implementation (OSDI)*, Boston, MA, December 2002. USENIX.
- [3] Peter A. Alsberg and John D. Day. A principle for resilient sharing of distributed resources. In *Proceedings of the 2nd International Conference on Software Engineering*, pages 562–570. IEEE Computer Society Press, 1976.
- [4] Guillermo A. Alvarez, Walter A. Burkhard, and Flaviu Cristian. Tolerating multiple

- failures in RAID architectures with optimal storage and uniform declustering. In *Proceedings of the 24th Int'l Symposium on Computer Architecture*, pages 62–72, Denver, CO, June 1997. ACM.
- [5] Eric Anderson, Joseph Hall, Jason Hartline, Michael Hobbs, Anna R. Karlin, Jared Saia, Ram Swaminathan, and John Wilkes. An experimental study of data migration algorithms. In *Proceedings of the 5th International Workshop on Algorithm Engineering*, pages 145–158, London, UK, 2001. Springer-Verlag.
- [6] Eric Anderson, Michael Hobbs, Kimberly Keeton, Susan Spence, Mustafa Uysal, and Alistair Veitch. Hippodrome: running circles around storage administration. In *Proceedings of the 2002 Conference on File and Storage Technologies (FAST)*, Monterey, CA, January 2002.
- [7] Alain Azagury, Vladimir Dreizin, Michael Factor, Ealan Henis, Dalit Naor, Noam Rinetzky, Ohad Rodeh, Julian Satran, Ami Tavory, and Lena Yerushalmi. Towards an object store. In *Proceedings of the 20th IEEE / 11th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 165–176, April 2003.
- [8] Mary G. Baker, John H. Hartman, Michael D. Kupfer, Ken W. Shirriff, and John K. Ousterhout. Measurements of a distributed file system. In *Proceedings of the 13th ACM Symposium on Operating Systems Principles (SOSP '91)*, pages 198–212, October 1991.
- [9] N. Belaramani, M. Dahlin, L. Gao, A. Nayate, A. Venkataramani, P. Yalagandula, and

- J. Zheng. PRACTI replication. In *Proceedings of the 3rd Symposium on Networked Systems Design and Implementation (NSDI)*, pages 59–72, May 2006.
- [10] Magnus E. Bjornsson and Liuba Shrira. Buddycache: Cache coherence for transactional peer group applications. In *Second IEEE Workshop on Internet Applications (WIAPP '01)*, volume 00, page 57, Los Alamitos, CA, 2001. IEEE Computer Society.
- [11] Burton H Bloom. Space/time trade-offs in hash coding with allowable errors. *Communications of the ACM*, 13(7):422–426, July 1970.
- [12] IEEE Standard Board. *Information technology—Portable Operating System Interface (POSIX) Part 1: System Application Program Interface (API) [C Language]*. Institute of Electrical and Electronics Engineers, Inc., April 1990.
- [13] Peter Braam, Michael Callahan, and Phil Schwan. The intermezzo file system. In *Proceedings of the 3rd of the Perl Conference, O'Reilly Open Source Convention*, Monterey, CA, USA, August 1999.
- [14] Peter J. Braam. The Lustre storage architecture. <http://www.lustre.org/documentation.html>, Cluster File Systems, Inc., August 2004.
- [15] Scott A. Brandt, Lan Xue, Ethan L. Miller, and Darrell D. E. Long. Efficient metadata management in large distributed file systems. In *Proceedings of the 20th IEEE / 11th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 290–298, April 2003.

- [16] A. Brinkmann, S. Effert, F. Meyer auf der Heide, and C. Scheideler. Dynamic and redundant data placement. In *Proceedings of the 27th International Conference on Distributed Computing Systems (ICDCS '07)*, 2007.
- [17] André Brinkmann, Kay Salzwedel, and Christian Scheideler. Efficient, distributed data placement strategies for storage area networks. In *Proceedings of the 12th ACM Symposium on Parallel Algorithms and Architectures (SPAA)*, pages 119–128. ACM Press, 2000. Extended Abstract.
- [18] D. M. Choy, R. Fagin, and L. Stockmeyer. Efficiently extendible mappings for balanced data distribution. *Algorithmica*, 16:215–232, 1996.
- [19] Douglas Comer. The ubiquitous B-tree. *ACM Computing Surveys*, 11(2):121–137, June 1979.
- [20] Peter F. Corbett and Dror G. Feitelson. The Vesta parallel file system. *ACM Transactions on Computer Systems*, 14(3):225–264, 1996.
- [21] Frank Dabek, M. Frans Kaashoek, David Karger, Robert Morris, and Ion Stoica. Wide-area cooperative storage with CFS. In *Proceedings of the 18th ACM Symposium on Operating Systems Principles (SOSP '01)*, pages 202–215, Banff, Canada, October 2001. ACM.
- [22] Michael Dahlin, Clifford Mather, Randolph Wang, Thomas E. Anderson, and David A. Patterson. A quantitative analysis of cache policies for scalable network file systems. In *Measurement and Modeling of Computer Systems*, pages 150–160, 1994.

- [23] Michael Dahlin, Randy Wang, Tom Anderson, and David Patterson. Cooperative caching: Using remote client memory to improve file system performance. In *Proceedings of the 1st Symposium on Operating Systems Design and Implementation (OSDI)*, pages 267–280, November 1994.
- [24] Tim Kaldewey Roberto C. Pineiro Anna Povzner Scott A. Brandt Richard A. Golding Theodore M. Wong Carlos Maltzahn David O. Bigelow, Suresh Iyer. End-to-end performance management for scalable distributed storage. In *Proceedings of the 2007 ACM Petascale Data Storage Workshop (PDSW 07)*, November 2007.
- [25] John R. Douceur and Jon Howell. Distributed directory service in the farsite file system. In *Proceedings of the 7th Symposium on Operating Systems Design and Implementation (OSDI)*, pages 321–334, Seattle, WA, November 2006. Usenix.
- [26] Fred Douglass and John K. Ousterhout. Beating the I/O bottleneck: A case for log-structured files systems. Technical Report UCB/CSD 88/467, University of California, Berkeley, October 1988.
- [27] Daniel Ellard, Jonathan Ledlie, Pia Malkani, and Margo Seltzer. Passive nfs tracing of email and research workloads. In *Proceedings of the Second USENIX Conference on File and Storage Technologies (FAST)*, pages 203–216, San Francisco, CA, March 2003. USENIX.
- [28] Richard A. Floyd and Carla Schlatter Ellis. Directory reference patterns in hierarchical

- file systems. *IEEE Transactions on Knowledge and Data Engineering*, 1(2):238–247, 1989.
- [29] Gregory R. Ganger and M. Frans Kaashoek. Embedded inodes and explicit groupings: Exploiting disk bandwidth for small files. In *Proceedings of the 1997 USENIX Annual Technical Conference*, pages 1–17. USENIX Association, January 1997.
- [30] Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. The Google file system. In *Proceedings of the 19th ACM Symposium on Operating Systems Principles (SOSP '03)*, Bolton Landing, NY, October 2003. ACM.
- [31] Garth A. Gibson, David F. Nagle, Khalil Amiri, Jeff Butler, Fay W. Chang, Howard Gobioff, Charles Hardin, Erik Riedel, David Rochberg, and Jim Zelenka. A cost-effective, high-bandwidth storage architecture. In *Proceedings of the 8th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pages 92–103, San Jose, CA, October 1998.
- [32] David K. Gifford, Roger M. Needham, and Michael D. Schroeder. The Cedar file system. *Communications of the ACM*, 31(3), March 1988.
- [33] Howard Gobioff, Garth Gibson, and Doug Tygar. Security for network attached storage devices. Technical Report TR CMU-CS-97-185, Carnegie Mellon, October 1997.
- [34] Ashish Goel, Cyrus Shahabi, Did Shu-Yuen Yao, and Roger Zimmerman. SCADDAR: An efficient randomized technique to reorganize continuous media blocks. In *Proceed-*

ings of the 18th International Conference on Data Engineering (ICDE '02), pages 473–482, February 2002.

- [35] Richard Golding and Elizabeth Borowsky. Fault-tolerant replication management in large-scale distributed storage systems. In *Proceedings of the 18th Symposium on Reliable Distributed Systems (SRDS '99)*, pages 144–155, October 1999.
- [36] Andrew Granville. On elementary proofs of the Prime Number Theorem for Arithmetic Progressions, without characters. In *Proceedings of the 1993 Amalfi Conference on Analytic Number Theory*, pages 157–194, Salerno, Italy, 1993.
- [37] Andreas Haeberlen, Alan Mislove, and Peter Druschel. Glacier: Highly durable, decentralized storage despite massive correlated failures. In *Proceedings of the 2nd Symposium on Networked Systems Design and Implementation (NSDI)*, Boston, MA, May 2005. USENIX.
- [38] C. R. Hertel. *Implementing CIFS: The Common Internet File System*. Prentice Hall, 2003.
- [39] Dean Hildebrand and Peter Honeyman. Exporting storage systems in a scalable manner with pNFS. Technical Report CITI-05-1, CITI, University of Michigan, February 2005.
- [40] Dave Hitz, James Lau, and Michael Malcom. File system design for an NFS file server appliance. In *Proceedings of the Winter 1994 USENIX Technical Conference*, pages 235–246, San Francisco, CA, January 1994.
- [41] R. J. Honicky and Ethan L. Miller. Replication under scalable hashing: A family of

- algorithms for scalable decentralized data distribution. In *Proceedings of the 18th International Parallel & Distributed Processing Symposium (IPDPS 2004)*, Santa Fe, NM, April 2004. IEEE.
- [42] Andy Hospodor and Ethan L. Miller. Interconnection architectures for petabyte-scale high-performance storage systems. In *Proceedings of the 21st IEEE / 12th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 273–281, College Park, MD, April 2004.
- [43] John H. Howard, Michael L. Kazar, Sherri G. Menees, David A. Nichols, M. Satyanarayanan, Robert N. Sidebotham, and Michael J. Wes. Scale and performance in a distributed file system. *ACM Transactions on Computer Systems*, 6(1):51–81, February 1988.
- [44] IBM Corporation. IBM white paper: IBM storage tank – a distributed storage system, January 2002.
- [45] Robert J. Jenkins. Hash functions for hash table lookup. <http://burtleburtle.net/bob/hash/evahash.html>, 1997.
- [46] Minwen Ji, Alistair Veitch, and John Wilkes. Seneca: remote mirroring done write. In *Proceedings of the 2003 USENIX Annual Technical Conference*, pages 253–268, June 2003.
- [47] David Karger, Eric Lehman, Tom Leighton, Mathew Levine, Daniel Lewin, and Rina Panigrahy. Consistent hashing and random trees: Distributed caching protocols for re-

- lieving hot spots on the World Wide Web. In *ACM Symposium on Theory of Computing*, pages 654–663, May 1997.
- [48] Peter J Keleher. Decentralized replicated object protocols. In *Proceedings of the Eighteenth ACM Symposium on Principles of Distributed Computing (PODC 1999)*, pages 143–151, Atlanta, Georgia, 1999.
- [49] Deepak R. Kenchammana-Hosekote, Richard A. Golding, Claudio Fleiner, and Omer A. Zaki. The design and evaluation of network raid protocols. Technical Report RJ 10316 (A0403-006), IBM Research, Almaden Center, March 2004.
- [50] Richard A. Golding Ralph A. Becker-Szendy Kristal T. Pollack, Darrell D. E. Long. Quota enforcement for high-performance distributed storage systems. In *Proceedings of the 24rd IEEE / 14th NASA Goddard Conference on Mass Storage Systems and Technologies*, September 2007.
- [51] Nancy P. Kronenberg, Henry M. Levy, and William D. Strecker. VAXclusters: A closely-coupled distributed system. *ACM Transactions on Computer Systems*, 4(2):130–146, 1986.
- [52] John Kubiawicz, David Bindel, Yan Chen, Patrick Eaton, Dennis Geels, Ramakrishna Gummadi, Sean Rhea, Hakim Weatherspoon, Westly Weimer, Christopher Wells, and Ben Zhao. OceanStore: An architecture for global-scale persistent storage. In *Proceedings of the 9th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, Cambridge, MA, November 2000. ACM.

- [53] Rivka Ladin, Barbara Liskov, Liuba Shrira, and Sanjay Ghemawat. Providing high availability using lazy replication. *ACM Transactions on Computer Systems*, 10(4):360–391, November 1992.
- [54] Leslie Lamport. The part-time parliament. *ACM Transactions on Computer Systems*, 16(2):133–169, 1998.
- [55] Rob Latham, Neill Miller, Robert Ross, and Phil Carns. A next-generation parallel file system for Linux clusters. *LinuxWorld*, pages 56–59, January 2004.
- [56] Edward K. Lee and Chandramohan A. Thekkath. Petal: Distributed virtual disks. In *Proceedings of the 7th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pages 84–92, Cambridge, MA, 1996.
- [57] Andrew Leung and Ethan L. Miller. Scalable security for large, high performance storage systems. In *Proceedings of the 2006 ACM Workshop on Storage Security and Survivability*. ACM, October 2006.
- [58] Andrew W. Leung, Ethan L. Miller, and Stephanie Jones. Scalable security for petascale parallel file systems. In *Proceedings of the 2007 ACM/IEEE Conference on Supercomputing (SC '07)*, November 2007.
- [59] Qiao Lian, Wei Chen, and Zheng Zhang. On the impact of replica placement to the reliability of distributed brick storage systems. In *Proceedings of the 25th International Conference on Distributed Computing Systems (ICDCS '05)*, pages 187–196, Los Alamitos, CA, 2005. IEEE Computer Society.

- [60] Barbara Liskov, Mark Day, and Liuba Shrira. Distributed object management in thor. In *International Workshop on Distributed Object Management*, pages 79–91, 1992.
- [61] Barbara Liskov, Sanjay Ghemawat, Robert Gruber, Paul Johnson, Liuba Shrira, and Michael Williams. Replication in the Harp file system. In *Proceedings of the 13th ACM Symposium on Operating Systems Principles (SOSP '91)*, pages 226–238. ACM, 1991.
- [62] Christopher R. Lumb, Gregory R. Ganger, and Richard Golding. D-SPTF: Decentralized request distribution in brick-based storage systems. In *Proceedings of the 11th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pages 37–47, Boston, MA, 2004.
- [63] John MacCormick, Nick Murphy, Marc Najork, Chandramohan A. Thekkath, and Li-dong Zhou. Boxwood: Abstractions as the foundation for storage infrastructure. In *Proceedings of the 6th Symposium on Operating Systems Design and Implementation (OSDI)*, San Francisco, CA, December 2004.
- [64] M. Kirk McKusick and Gregory R. Ganger. Soft updates: A technique for eliminating most synchronous writes in the Fast File System. In *Proceedings of the Freenix Track: 1999 USENIX Annual Technical Conference*, pages 1–18, June 1999.
- [65] Marshall Kirk McKusick, William N. Joy, Samuel J. Leffler, and Robert S. Fabry. A fast file system for UNIX. *ACM Transactions on Computer Systems*, 2(3):181–197, August 1984.

- [66] J. Menon, D. A. Pease, R. Rees, L. Duyanovich, and B. Hillsberg. IBM Storage Tank—a heterogeneous scalable SAN file system. *IBM Systems Journal*, 42(2):250–267, 2003.
- [67] Ethan L. Miller and Randy H. Katz. RAMA: An easy-to-use, high-performance parallel file system. *Parallel Computing*, 23(4):419–446, 1997.
- [68] David Nagle, Denis Serenyi, and Abbie Matthews. The Panasas ActiveScale storage cluster—delivering scalable high bandwidth storage. In *Proceedings of the 2004 ACM/IEEE Conference on Supercomputing (SC '04)*, November 2004.
- [69] Christopher A. Olson and Ethan L. Miller. Secure capabilities for a petabyte-scale object-based distributed file system. In *Proceedings of the 2005 ACM Workshop on Storage Security and Survivability*, Fairfax, VA, November 2005.
- [70] Oracle. Rds: Reliable datagram sockets. <http://oss.oracle.com/project/rds/>.
- [71] John K. Ousterhout, Andrew R. Cherenon, Frederick Douglass, Michael N. Nelson, and Brent B. Welch. The Sprite network operating system. *IEEE Computer*, 21(2):23–36, February 1988.
- [72] Brian Pawlowski, Chet Juszczak, Peter Staubach, Carl Smith, Diane Lebel, and Dave Hitz. NFS version 3: Design and implementation. In *Proceedings of the Summer 1994 USENIX Technical Conference*, pages 137–151, 1994.
- [73] Juan Piernas, Toni Cortes, and José M. García. Dualfs: a new journaling file system without meta-data duplication. In *Proceedings of the 16th International Conference on Supercomputing*, pages 84–95, New York, NY, 2002. ACM.

- [74] Kristal T. Pollack and Scott A. Brandt. Efficient access control for distributed hierarchical file systems. In *Proceedings of the 22nd IEEE / 13th NASA Goddard Conference on Mass Storage Systems and Technologies*, April 2005.
- [75] Sean Quinlan and Sean Dorward. Venti: A new approach to archival storage. In *Proceedings of the 2002 Conference on File and Storage Technologies (FAST)*, pages 89–101, Monterey, California, USA, 2002. USENIX.
- [76] Ohad Rodeh and Avi Teperman. zFS—a scalable distributed file system using object disks. In *Proceedings of the 20th IEEE / 11th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 207–218, April 2003.
- [77] Drew Roselli, Jay Lorch, and Tom Anderson. A comparison of file system workloads. In *Proceedings of the 2000 USENIX Annual Technical Conference*, pages 41–54, San Diego, CA, June 2000. USENIX Association.
- [78] Mendel Rosenblum and John K. Ousterhout. The design and implementation of a log-structured file system. *ACM Transactions on Computer Systems*, 10(1):26–52, February 1992.
- [79] Russel Glen Ross. Cluster storage for commodity computation. Technical Report UCAM-CL-TR-690, University of Cambridge, Cambridge, UK, June 2007.
- [80] Antony Rowstron and Peter Druschel. Storage management and caching in PAST, a large-scale, persistent peer-to-peer storage utility. In *Proceedings of the 18th ACM Sym-*

posium on Operating Systems Principles (SOSP '01), pages 188–201, Banff, Canada, October 2001. ACM.

- [81] Antony Rowstrong and Peter Druschel. Pastry: Scalable, decentralized object location, and routing for large-scale peer-to-peer systems. In *Proceedings of the IFIP/ACM International Conference on Distributed Systems Platforms (Middleware)*, pages 329–350, 2001.
- [82] Yashushi Saito, Svend Frølund, Alistair Veitch, Arif Merchant, and Susan Spence. FAB: Building distributed enterprise disk arrays from commodity components. In *Proceedings of the 11th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pages 48–58, 2004.
- [83] Yasushi Saito, Christos Karamanolis, Magnus Karlsson, and Mallick Mahalingam. Taming aggressive replication in the Pangaea wide-area file system. In *Proceedings of the 5th Symposium on Operating Systems Design and Implementation (OSDI)*. USENIX, December 2002.
- [84] Jose Renato Santos, Richard R. Muntz, and Berthier Ribeiro-Neto. Comparing random data allocation and data striping in multimedia servers. In *Proceedings of the 2000 SIGMETRICS Conference on Measurement and Modeling of Computer Systems*, pages 44–55, Santa Clara, CA, June 2000. ACM Press.
- [85] M. Satyanarayanan, James J. Kistler, Puneet Kumar, Maria E. Okasaki, Ellen H. Siegel,

and David C. Steere. Coda: A highly available file system for a distributed workstation environment. *IEEE Transactions on Computers*, 39(4):447–459, 1990.

- [86] Frank Schmuck and Roger Haskin. GPFS: A shared-disk file system for large computing clusters. In *Proceedings of the 2002 Conference on File and Storage Technologies (FAST)*, pages 231–244. USENIX, January 2002.
- [87] Margo Seltzer, Keith Bostic, M. Kirk McKusick, and Carl Staelin. An implementation of a log-structured file system for UNIX. In *Proceedings of the Winter 1993 USENIX Technical Conference*, pages 307–326, January 1993.
- [88] Margo Seltzer, Keith A. Smith, Hari Balakrishnan, Jacqueline Chang, Sara McMains, and Venkata Padmanabhan. File system logging versus clustering: A performance comparison. In *Proceedings of the Winter 1995 USENIX Technical Conference*, pages 249–264, 1995.
- [89] K. W. Shirriff and J. K. Ousterhout. A trace-driven analysis of name and attribute caching in a distributed system. In *Proceedings of the Winter 1992 USENIX Technical Conference*, pages 315–331, San Francisco, CA, January 1992.
- [90] Ion Stoica, Robert Morris, David Karger, M. Frans Kaashoek, and Hari Balakrishnan. Chord: A scalable peer-to-peer lookup service for Internet applications. In *Proceedings of the Conference on Applications, Technologies, Architectures, and Protocols for Computer Communication (SIGCOMM '01)*, pages 149–160, San Diego, CA, 2001.
- [91] Michael Stonebraker. Concurrency control and consistency of multiple copies of data in

distributed INGRES. *IEEE Transactions on Software Engineering*, SE-5:188–194, May 1979.

- [92] A. Sweeney, D. Doucette, W. Hu, C. Anderson, M. Nishimoto, and G. Peck. Scalability in the XFS file system. In *Proceedings of the 1996 USENIX Annual Technical Conference*, pages 1–14, January 1996.
- [93] Hong Tang, Aziz Gulbeden, Jingyu Zhou, William Strathearn, Tao Yang, and Lingkun Chu. A self-organizing storage cluster for parallel data-intensive applications. In *Proceedings of the 2004 ACM/IEEE Conference on Supercomputing (SC '04)*, Pittsburgh, PA, November 2004.
- [94] Chandramohan A. Thekkath, Timothy Mann, and Edward K. Lee. Frangipani: A scalable distributed file system. In *Proceedings of the 16th ACM Symposium on Operating Systems Principles (SOSP '97)*, pages 224–237, 1997.
- [95] Theodore Y. Ts'o and Stephen Tweedie. Planned extensions to the Linux EXT2/EXT3 filesystem. In *Proceedings of the Freenix Track: 2002 USENIX Annual Technical Conference*, pages 235–244, Monterey, CA, June 2002. USENIX.
- [96] Robbert van Renesse and Fred B. Schneider. Chain replication for supporting high throughput and availability. In *Proceedings of the 6th Symposium on Operating Systems Design and Implementation (OSDI)*, pages 91–104, San Francisco, CA, December 2004.
- [97] Feng Wang, Scott A. Brandt, Ethan L. Miller, and Darrell D. E. Long. OBFS: A file

- system for object-based storage devices. In *Proceedings of the 21st IEEE / 12th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 283–300, College Park, MD, April 2004. IEEE.
- [98] Feng Wang, Qin Xin, Bo Hong, Scott A. Brandt, Ethan L. Miller, Darrell D. E. Long, and Tyce T. McLarty. File system workload analysis for large scale scientific computing applications. In *Proceedings of the 21st IEEE / 12th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 139–152, College Park, MD, April 2004.
- [99] Randolph Y. Wang and Thomas E. Anderson. xFS: A wide area mass storage file system. In *Proceedings of the Fourth Workshop on Workstation Operating Systems*, pages 71–78, 1993.
- [100] Sage A. Weil, Scott A. Brandt, Ethan L. Miller, Darrell D. E. Long, and Carlos Maltzahn. Ceph: A scalable, high-performance distributed file system. In *Proceedings of the 7th Symposium on Operating Systems Design and Implementation (OSDI)*, Seattle, WA, November 2006. USENIX.
- [101] Sage A. Weil, Scott A. Brandt, Ethan L. Miller, and Carlos Maltzahn. CRUSH: Controlled, scalable, decentralized placement of replicated data. In *Proceedings of the 2006 ACM/IEEE Conference on Supercomputing (SC '06)*, Tampa, FL, November 2006. ACM.
- [102] Sage A. Weil, Kristal T. Pollack, Scott A. Brandt, and Ethan L. Miller. Dynamic metadata management for petabyte-scale file systems. In *Proceedings of the 2004 ACM/IEEE Conference on Supercomputing (SC '04)*, Pittsburgh, PA, November 2004. ACM.

- [103] Brent Welch. POSIX IO extensions for HPC. In *Proceedings of the 4th USENIX Conference on File and Storage Technologies (FAST)*, December 2005.
- [104] Brent Welch and Garth Gibson. Managing scalability in object storage systems for HPC Linux clusters. In *Proceedings of the 21st IEEE / 12th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 433–445, April 2004.
- [105] M. Wiesmann, F. Pedone, A. Schiper, B. Kemme, and G. Alonso. Understanding replication in databases and distributed systems. In *Proceedings of the 20th International Conference on Distributed Computing Systems (ICDCS '00)*, 2000.
- [106] Matthias Wiesmann and Andre Schiper. Comparison of database replication techniques based on total order broadcast. *IEEE Transactions on Knowledge and Data Engineering*, 17(4):551–566, April 2005.
- [107] Theodore M. Wong, Richard A. Golding, Joseph S. Glider, Elizabeth Borowsky, Ralph A. Becker-Szendy, Claudio Fleiner, Deepak R. Kenchammana-Hosekote, and Omer A. Zaki. Kybos: self-management for distributed brick-based storage. Research Report RJ 10356, IBM Almaden Research Center, August 2005.
- [108] Changxun Wu and Randal Burns. Tunable randomization for load management in shared-disk clusters. *ACM Transactions on Storage*, 1(1):108–131, December 2004.
- [109] Joel C. Wu and Scott A. Brandt. The design and implementation of AQUA: an adaptive quality of service aware object-based storage device. In *Proceedings of the 23rd IEEE*

- / *14th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 209–218, College Park, MD, May 2006.
- [110] Joel C. Wu and Scott A. Brandt. Providing quality of service support in object-based file system. In *Proceedings of the 24rd IEEE / 14th NASA Goddard Conference on Mass Storage Systems and Technologies*, September 2007.
- [111] Qin Xin, Ethan L. Miller, and Thomas J. E. Schwarz. Evaluation of distributed recovery in large-scale storage systems. In *Proceedings of the 13th IEEE International Symposium on High Performance Distributed Computing (HPDC)*, pages 172–181, Honolulu, HI, June 2004.
- [112] Qin Xin, Ethan L. Miller, Thomas J.E. Schwarz, Darrell D. E. Long, Scott A. Brandt, and Witold Litwin. Reliability mechanisms for very large storage systems. In *Proceedings of the 20th IEEE / 11th NASA Goddard Conference on Mass Storage Systems and Technologies*, pages 146–156, April 2003.
- [113] Zhihui Zhang and Kanad Ghose. hfs: A hybrid file system prototype for improving small file and metadata performance. In *Proceedings of EuroSys 2007*, March 2007.
- [114] Ben Y. Zhao, Lin gHuang, Jeremy Stribling, Sean C. Rhea, and Anthony D. Joseph nad John D. Kubiatowicz. Tapestry: A global-scale overlay for rapid service deployment. *IEEE Journal on Selected Areas in Communications*, 22:41–53, January 2003.